

Fundamentos de los Sistemas Operativos (FSO)

Departamento de Informática de Sistemas y Computadoras (DISCA)
Universitat Politècnica de València

Part 3: Memory management

Seminar 7

Memory map of a Linux process

fso

DISCA



UNIVERSITAT
POLITÈCNICA
DE VALÈNCIA

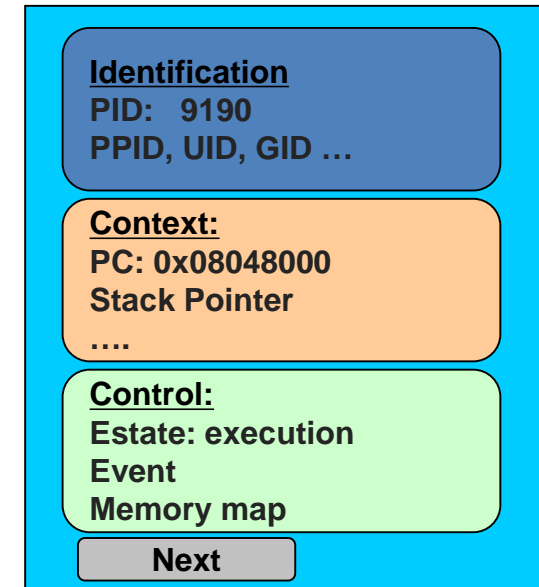
- Goals
 - To understand the **process memory map concept**
 - To know the features of **process memory map in Linux**
 - To know the technique used to **map files into memory**
 - To be aware about the advantages and disadvantages of using **static and dynamic libraries**
- Bibliography
 - Carretero, chapter 5

- **Introduction**
- Memory map of a Linux process
- Memory mapped files
- Dynamic linking libraries

- **Process memory map**

- The OS manages the memory map of every process during its lifetime
- Memory map is a process **attribute** -> it is included in its PCB
- It contains information about a process memory regions:
 - Code, data, stack, etc

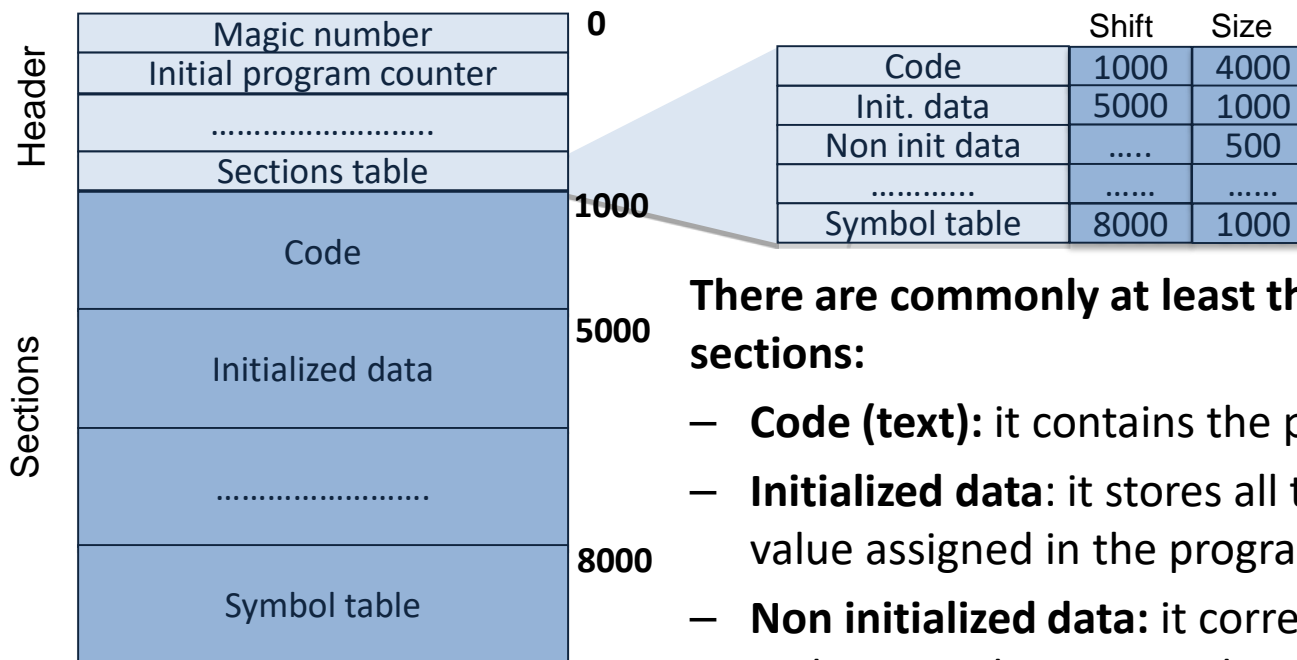
PCB



- The **initial memory map** of a process is strongly linked to its executable file
- Nowadays OSs offer a **dynamic memory model** that gives support to process memory regions allocation:
 - **New regions could be created** to allocate dynamic process entities: stack, mapped files, dynamic memory, etc.
 - **Unused regions could be removed**

• Simplified format of an executable file

- After compiling and linking an executable file is generated that contains the program machine code ready for execution
- An executable file is structured into a header and a sections set
 - **Header:** it contains control information that allows reading into the remaining executable file content
 - **Sections:** Executable file content is organized in sections

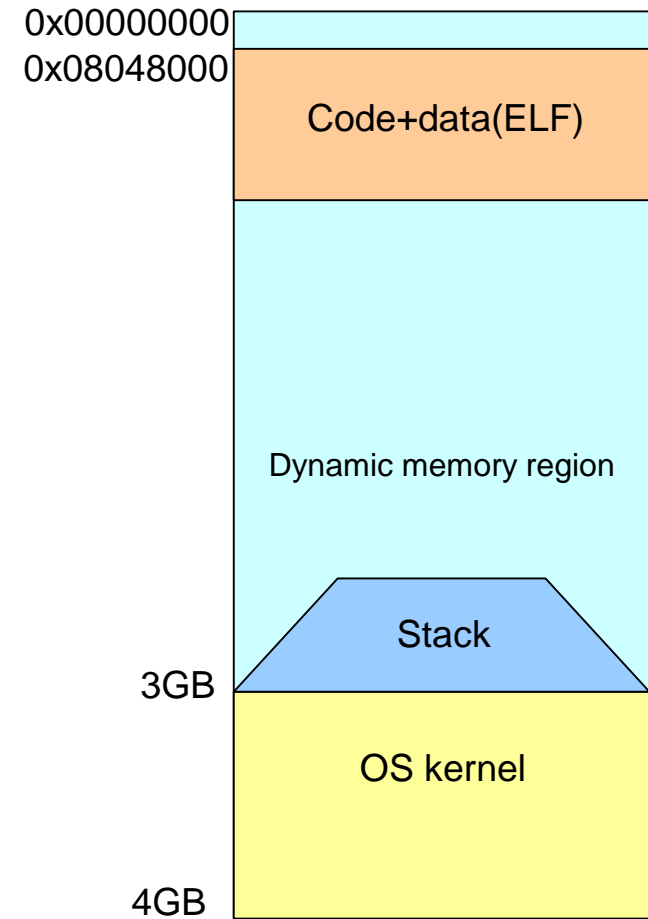


There are commonly at least the following three sections:

- **Code (text):** it contains the program code
- **Initialized data:** it stores all the variables with a value assigned in the program text
- **Non initialized data:** it corresponds to variables without a value assigned in the program text

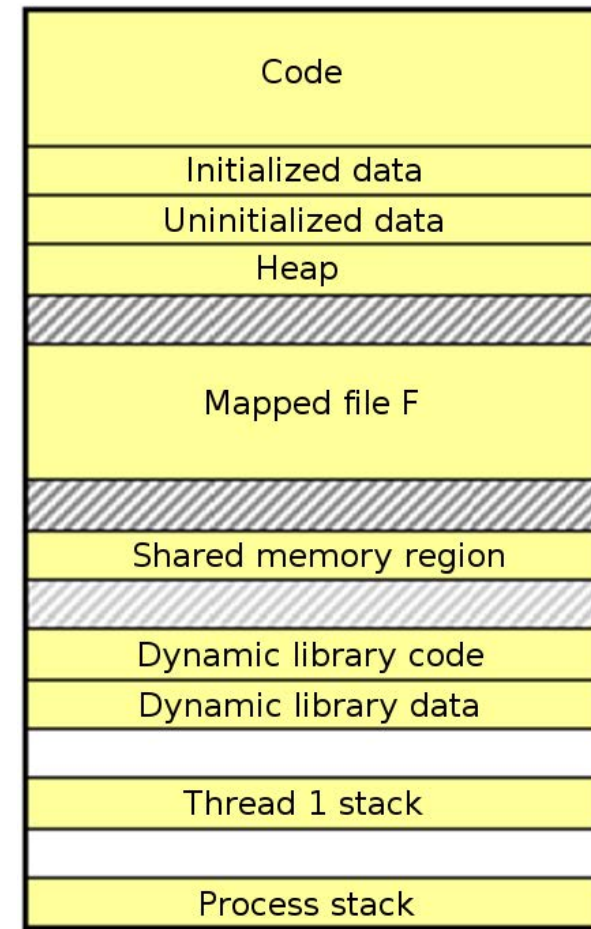
- Introduction
- **Memory map of a Linux process**
- Memory mapped files
- Dynamic linking libraries

- Linux 2.6.x versions on 32 bits architecture:
 - Logical space is 4GB.
 - 1st GByte: Code + initialized data begin at address 0x08048000 (ELF format).
 - 3rd GByte: The stack starts in the 3rd GB upper side and it grows downwards.
 - 4th GByte: The upper GB is reserved to the OS
 - The remaining space is available for data (uninitialized + heap) and dynamic linked libraries

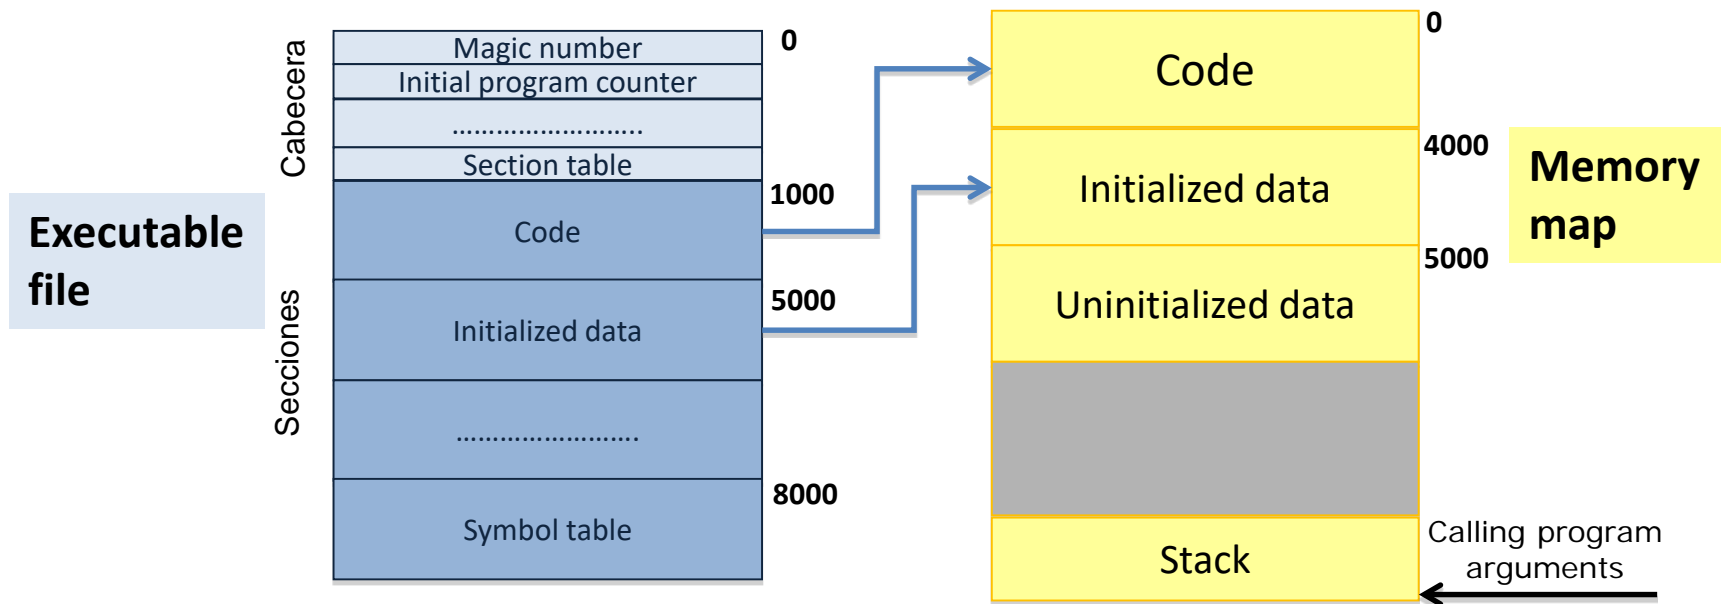


- **Memory maps are region made**

- A region has a certain type of information associated
 - Code, initialized data, uninitialized data, mapped files and dynamic linked libraries
- A region is a **contiguous memory chunk** featured by the address inside the process memory map where it begins and its size
- Region features:
 - **Support:** where the region information is stored
 - File support: region information is stored in a file
 - Without support: region without initial content
 - **Share type:**
 - Private (p) : content is only accessible by the owner process
 - Shared: content is accessible by several processes
 - **Protection:** region access type allowed
 - Read, write and/or execution
 - **Size:** fixed or variable



- Program execution starts building a process memory map from the executable file
- **Every executable file section becomes an initial map region**
 - **Code (text):** shared region, read and execution, fixed size, executable file support
 - **Initialized data:** private region (every process needs a private copy), read and write, fixed size, executable file support
 - **Uninitialized data:** private region, read and write, fixed size, without support (some compilers/languages initialize it to zero)
 - **Stack:** private region, read and write, variable size, without support. It grows towards lower addresses. When execution starts it only contains calling program arguments



- Process memory map is **dynamic** -> along process lifetime some regions can be created like:
 - **Heap**
 - Dynamic memory support (i.e. pointers)
 - Private, read and write, variable size, without support (initialized to zero)
 - It grows towards upper addresses
 - **Mapped files**
 - When a file is mapped into memory a new region is created
 - Variable size, file support
 - Protection and sharing specified in the mapping
 - **Shared memory**
 - Region that supports interprocess communication
 - Shared, variable size, without support (initialized to zero)
 - Protection specified in the program
 - **Thread stacks**
 - Every thread stack has its own region
 - Same features as process stack

- **Two visualization methods of a process memory map:**

- Viewing the process maps file:

\$cat /proc/PID/maps

- » Logic address range
- » Permissions
- » Shift from the beginning of the executable file
- » Device
- » Node-i
- » Mapped file name

- Executing the **shell command**

\$pmap PID

- » Logic base address
- » Size
- » Permissions
- » Mapped file name

Nota: Replacing PID by variable **\$\$** we refer to the **process** in **execution** like:

\$ pmap \$\$

\$ cat /proc/\$\$/maps

32 bit architecture

pblanes@pblanes-desktop:~\$ pmap \$\$

1608: /bin/bash

00140000 32K r-x-- /lib/tls/i686/cmov/libnss_nis-2.11.1.so

00148000 4K r---- /lib/tls/i686/cmov/libnss_nis-2.11.1.so

00149000 4K rw--- /lib/tls/i686/cmov/libnss_nis-2.11.1.so

00266000 8K r-x-- /lib/tls/i686/cmov/libdl-2.11.1.so

00268000 4K r---- /lib/tls/i686/cmov/libdl-2.11.1.so

00269000 4K rw--- /lib/tls/i686/cmov/libdl-2.11.1.so

00319000 1356K r-x-- /lib/tls/i686/cmov/libc-2.11.1.so

0046c000 4K ----- /lib/tls/i686/cmov/libc-2.11.1.so

0046d000 8K r---- /lib/tls/i686/cmov/libc-2.11.1.so

0046f000 4K rw--- /lib/tls/i686/cmov/libc-2.11.1.so

00470000 12K rw--- [anon]

005b0000 4K r-x-- [anon]

007f7000 208K r-x-- /lib/libncurses.so.5.7

0082b000 4K ----- /lib/libncurses.so.5.7

0082c000 8K r---- /lib/libncurses.so.5.7

0082e000 4K rw--- /lib/libncurses.so.5.7

00cbc000 108K r-x-- /lib/ld-2.11.1.so

00cd7000 4K r---- /lib/ld-2.11.1.so

00cd8000 4K rw--- /lib/ld-2.11.1.so

00cf7000 24K r-x-- /lib/tls/i686/cmov/libnss_compat-2.11.1.so

00cfd000 4K r---- /lib/tls/i686/cmov/libnss_compat-2.11.1.so

00cfe000 4K rw--- /lib/tls/i686/cmov/libnss_compat-2.11.1.so

00f56000 40K r-x-- /lib/tls/i686/cmov/libnss_files-2.11.1.so

00f60000 4K r---- /lib/tls/i686/cmov/libnss_files-2.11.1.so

00f61000 4K rw--- /lib/tls/i686/cmov/libnss_files-2.11.1.so

00fc1000 76K r-x-- /lib/tls/i686/cmov/libnsl-2.11.1.so

00fd4000 4K r---- /lib/tls/i686/cmov/libnsl-2.11.1.so

00fd5000 4K rw--- /lib/tls/i686/cmov/libnsl-2.11.1.so

00fd6000 8K rw--- [anon]

08048000 780K r-x-- /bin/bash

0810b000 4K r---- /bin/bash

0810c000 20K rw--- /bin/bash

08111000 20K rw--- [anon]

0876e000 1368K rw--- [anon]

b75ef000 156K r---- /usr/share/locale-langpack/es/LC_MESSAGES/bash.mo

b7616000 252K r---- /usr/lib/locale/es_ES.utf8/LC_CTYPE

b7655000 4K r---- /usr/lib/locale/es_ES.utf8/LC_NUMERIC

b7656000 4K r---- /usr/lib/locale/es_ES.utf8/LC_TIME

b7657000 1144K r---- /usr/lib/locale/es_ES.utf8/LC_COLLATE

b7775000 8K rw--- [anon]

b7777000 4K r---- /usr/lib/locale/es_ES.utf8/LC_MONETARY

b7778000 4K r---- /usr/lib/locale/es_ES.utf8/LC_MESSAGES/SYS_LC_MESSAGES

b7779000 4K r---- /usr/lib/locale/es_ES.utf8/LC_PAPER

b777a000 4K r---- /usr/lib/locale/es_ES.utf8/LC_NAME

b777b000 4K r---- /usr/lib/locale/es_ES.utf8/LC_ADDRESS

b777c000 4K r---- /usr/lib/locale/es_ES.utf8/LC_TELEPHONE

b777d000 4K r---- /usr/lib/locale/es_ES.utf8/LC_MEASUREMENT

b777e000 28K r--s- /usr/lib/gconv/gconv-modules.cache

b7785000 4K r---- /usr/lib/locale/es_ES.utf8/LC_IDENTIFICATION

b7786000 8K rw--- [anon]

bf95a000 84K rw--- [stack]

total 5868K

pblanes@shell-sisop:~\$ pmap -d \$\$

29916: -bash

Address	Kbytes	Mode	Offset	Device	Mapping
0000000000400000	760	r-x--	0000000000000000	008:00002	bash
00000000006bd000	40	rw---	00000000000bd000	008:00002	bash
00000000006c7000	2616	rw---	00000000006c7000	000:00000	[anon]
00007fa6e6728000	40	r-x--	0000000000000000	008:00002	libnss_files-2.7.so
00007fa6e6732000	2048	----	000000000000a000	008:00002	libnss_files-2.7.so
00007fa6e6932000	8	rw---	000000000000a000	008:00002	libnss_files-2.7.so
00007fa6e6934000	40	r-x--	0000000000000000	008:00002	libnss_nis-2.7.so
00007fa6e693e000	2044	----	000000000000a000	008:00002	libnss_nis-2.7.so
00007fa6e6b3d000	8	rw---	0000000000009000	008:00002	libnss_nis-2.7.so
00007fa6e6b3f000	88	r-x--	0000000000000000	008:00002	libnsl-2.7.so
00007fa6e6b55000	2044	----	0000000000016000	008:00002	libnsl-2.7.so
00007fa6e6d54000	8	rw---	0000000000015000	008:00002	libnsl-2.7.so
00007fa6e6d56000	8	rw---	00007fa6e6d56000	000:00000	[anon]
00007fa6e6d58000	32	r-x--	0000000000000000	008:00002	libnss_compat-2.7.so
00007fa6e6d60000	2044	----	0000000000008000	008:00002	libnss_compat-2.7.so
00007fa6e6f5f000	8	rw---	0000000000007000	008:00002	libnss_compat-2.7.so
00007fa6e6f61000	1376	r-x--	0000000000000000	008:00002	libc-2.7.so
00007fa6e70b9000	2048	----	00000000000158000	008:00002	libc-2.7.so
00007fa6e72b9000	12	r----	00000000000158000	008:00002	libc-2.7.so
00007fa6e72bc000	8	rw---	0000000000015b000	008:00002	libc-2.7.so
00007fa6e72be000	20	rw---	00007fa6e72be000	000:00000	[anon]
00007fa6e72c3000	8	r-x--	0000000000000000	008:00002	libdl-2.7.so
00007fa6e72c5000	2048	----	0000000000002000	008:00002	libdl-2.7.so
00007fa6e74c5000	8	rw---	0000000000002000	008:00002	libdl-2.7.so

64 bit architecture

00007fa6e74c7000	220	r-x--	0000000000000000	008:00002	libncurses.so.5.6
00007fa6e74fe000	2044	----	00000000000037000	008:00002	libncurses.so.5.6
00007fa6e76fd000	20	rw---	00000000000036000	008:00002	libncurses.so.5.6
00007fa6e7702000	116	r-x--	0000000000000000	008:00002	ld-2.7.so
00007fa6e77da000	60	r----	0000000000000000	008:00002	bash.mo
00007fa6e77e9000	252	r----	0000000000000000	008:00002	LC_CTYPE
00007fa6e7828000	900	r----	0000000000000000	008:00002	LC_COLLATE
00007fa6e7909000	8	rw---	00007fa6e7909000	000:00000	[anon]
00007fa6e790b000	4	r----	0000000000000000	008:00002	LC_NUMERIC
00007fa6e790c000	4	r----	0000000000000000	008:00002	LC_TIME
00007fa6e790d000	4	r----	0000000000000000	008:00002	LC_MONETARY
00007fa6e790e000	4	r----	0000000000000000	008:00002	SYS_LC_MESSAGES
00007fa6e790f000	4	r----	0000000000000000	008:00002	LC_PAPER
00007fa6e7910000	4	r----	0000000000000000	008:00002	LC_NAME
00007fa6e7911000	4	r----	0000000000000000	008:00002	LC_ADDRESS
00007fa6e7912000	4	r----	0000000000000000	008:00002	LC_TELEPHONE
00007fa6e7913000	4	r----	0000000000000000	008:00002	LC_MEASUREMENT
00007fa6e7914000	28	r--s-	0000000000000000	008:00002	gconv-modules.cache
00007fa6e791b000	4	r----	0000000000000000	008:00002	LC_IDENTIFICATION
00007fa6e791c000	12	rw---	00007fa6e791c000	000:00000	[anon]
00007fa6e791f000	8	rw---	0000000000001d000	008:00002	ld-2.7.so
00007fff8bc0f000	84	rw---	00007fffffe9000	000:00000	[stack]
00007fff8bd1d000	8	r-x--	00007fff8bd1d000	000:00000	[anon]
fffffffff600000	4	r-x--	0000000000000000	000:00000	[anon]

mapped: 21168K writeable/private: 2864K shared: 28K
gandreu@shell-sisop:~\$

- **File /proc/PID/maps**

- It contains actual memory regions associated to process PID and their access permissions
- Maps file format:

Address	perm	shift	device	node-i	path
08048000-08056000	r-xp	00000000	03:0c	64593	/usr/sbin/gpm
08056000-08058000	rw-p	0000d000	03:0c	64593	/usr/sbin/gpm
08058000-0805b000	rwxp	00000000	00:00	0	
40000000-40013000	r-xp	00000000	03:0c	4165	/lib/ld-2.2.4.so
40013000-40015000	rw-p	00012000	03:0c	4165	/lib/ld-2.2.4.so
4001f000-40135000	r-xp	00000000	03:0c	45494	/lib/libc-2.2.4.so
40135000-4013e000	rw-p	00115000	03:0c	45494	/lib/libc-2.2.4.so
4013e000-40142000	rw-p	00000000	00:00	0	
bf f f f 000-c0000000	rwxp	00000000	00:00	0	

Address: Logical address ranges for process regions

Permissions:
r = read
w = write
x = execute
s = shared
p = private (copy on write)

Node-i: device node-i, 0 means no node-i

Device: device id (major number: minor number)

Shift: shift inside the supporting file

32 bit architecture

pblanes\$ cat /proc/\$\$/maps

00140000-00148000	r-xp	00000000	08:01	266235	/lib/tls/i686/cmov/libnss_nis-2.11.1.so
00148000-00149000	r--p	00007000	08:01	266235	/lib/tls/i686/cmov/libnss_nis-2.11.1.so
00149000-0014a000	rw-p	00008000	08:01	266235	/lib/tls/i686/cmov/libnss_nis-2.11.1.so
00266000-00268000	r-xp	00000000	08:01	266220	/lib/tls/i686/cmov/libdl-2.11.1.so
00268000-00269000	r--p	00001000	08:01	266220	/lib/tls/i686/cmov/libdl-2.11.1.so
00269000-0026a000	rw-p	00002000	08:01	266220	/lib/tls/i686/cmov/libdl-2.11.1.so
00319000-0046c000	r-xp	00000000	08:01	266214	/lib/tls/i686/cmov/libc-2.11.1.so
0046c000-0046d000	---p	00153000	08:01	266214	/lib/tls/i686/cmov/libc-2.11.1.so
0046d000-0046f000	r--p	00153000	08:01	266214	/lib/tls/i686/cmov/libc-2.11.1.so
0046f000-00470000	rw-p	00155000	08:01	266214	/lib/tls/i686/cmov/libc-2.11.1.so
00470000-00473000	rw-p	00000000	00:00	0	
005b0000-005b1000	r-xp	00000000	00:00	0	[vdso]
007f7000-0082b000	r-xp	00000000	08:01	261740	/lib/libncurses.so.5.7
0082b000-0082c000	---p	00034000	08:01	261740	/lib/libncurses.so.5.7
0082c000-0082e000	r--p	00034000	08:01	261740	/lib/libncurses.so.5.7
0082e000-0082f000	rw-p	00036000	08:01	261740	/lib/libncurses.so.5.7
00cbc000-00cd7000	r-xp	00000000	08:01	261663	/lib/ld-2.11.1.so
00cd7000-00cd8000	r--p	0001a000	08:01	261663	/lib/ld-2.11.1.so
00cd8000-00cd9000	rw-p	0001b000	08:01	261663	/lib/ld-2.11.1.so

64 bit architecture

gandreu\$ cat /proc/\$\$/maps

```
00400000-004be000 r-xp 00000000 08:02 65607
006bd000-006c7000 rw-p 000bd000 08:02 65607
006c7000-00955000 rw-p 006c7000 00:00 0
7fa6e728000-7fa6e732000 r-xp 00000000 08:02 81942
7fa6e732000-7fa6e6932000 ---p 0000a000 08:02 81942
7fa6e6932000-7fa6e6934000 rw-p 0000a000 08:02 81942
7fa6e6934000-7fa6e693e000 r-xp 00000000 08:02 81944
7fa6e693e000-7fa6e6b3d000 ---p 0000a000 08:02 81944
7fa6e6b3d000-7fa6e6b3f000 rw-p 00009000 08:02 81944
7fa6e6b3f000-7fa6e6b55000 r-xp 00000000 08:02 81939
7fa6e6b55000-7fa6e6d54000 ---p 00016000 08:02 81939
7fa6e6d54000-7fa6e6d56000 rw-p 00015000 08:02 81939
7fa6e6d56000-7fa6e6d58000 rw-p 7fa6e6d56000 00:00 0
7fa6e6d58000-7fa6e6d60000 r-xp 00000000 08:02 81940
7fa6e6d60000-7fa6e6f5f000 ---p 00008000 08:02 81940
7fa6e6f5f000-7fa6e6f61000 rw-p 00007000 08:02 81940
7fa6e6f61000-7fa6e70b9000 r-xp 00000000 08:02 81930
7fa6e70b9000-7fa6e72b9000 ---p 00158000 08:02 81930
7fa6e72b9000-7fa6e72bc000 r--p 00158000 08:02 81930
7fa6e72bc000-7fa6e72be000 rw-p 0015b000 08:02 81930
7fa6e72be000-7fa6e72c3000 rw-p 7fa6e72be000 00:00 0
7fa6e72c3000-7fa6e72c5000 r-xp 00000000 08:02 81936
7fa6e72c5000-7fa6e74c5000 ---p 00002000 08:02 81936
7fa6e74c5000-7fa6e74c7000 rw-p 00002000 08:02 81936
7fa6e74c7000-7fa6e74fe000 r-xp 00000000 08:02 82217
7fa6e74fe000-7fa6e76fd000 ---p 00037000 08:02 82217
7fa6e76fd000-7fa6e7702000 rw-p 00036000 08:02 82217
7fa6e7702000-7fa6e771f000 r-xp 00000000 08:02 81927
7fa6e771f000-7fa6e77e9000 r--p 00000000 08:02 271736
```

```
/bin/bash
/bin/bash
[heap]
/lib/libnss_files-2.7.so
/lib/libnss_files-2.7.so
/lib/libnss_files-2.7.so
/lib/libnss_nis-2.7.so
/lib/libnss_nis-2.7.so
/lib/libnss_nis-2.7.so
/lib/libnsl-2.7.so
/lib/libnsl-2.7.so
/lib/libnsl-2.7.so
/lib/libnss_compat-2.7.so
/lib/libnss_compat-2.7.so
/lib/libnss_compat-2.7.so
/lib/libc-2.7.so
/lib/libc-2.7.so
/lib/libc-2.7.so
/lib/libc-2.7.so
/lib/libdl-2.7.so
/lib/libdl-2.7.so
/lib/libdl-2.7.so
/lib/libcurses.so.5.6
/lib/libcurses.so.5.6
/lib/libcurses.so.5.6
/lib/libd-2.7.so
/usr/share/locale-langpack/es/LC_MESSAGES/bash.mo
```

```
7fa6e77e9000-7fa6e7828000 r--p 00000000 08:02 439402
7fa6e7828000-7fa6e7909000 r--p 00000000 08:02 439411
7fa6e7909000-7fa6e790b000 rw-p 7fa6e7909000 00:00 0
7fa6e790b000-7fa6e790c000 r--p 00000000 08:02 439403
7fa6e790c000-7fa6e790d000 r--p 00000000 08:02 28628
7fa6e790d000-7fa6e790e000 r--p 00000000 08:02 28629
7fa6e790e000-7fa6e790f000 r--p 00000000 08:02 21591
/usr/lib/locale/es_ES.utf8/LC_MESSAGES/SYS_LC_MESSAGES
7fa6e790f000-7fa6e7910000 r--p 00000000 08:02 439406
7fa6e7910000-7fa6e7911000 r--p 00000000 08:02 439410
7fa6e7911000-7fa6e7912000 r--p 00000000 08:02 28631
7fa6e7912000-7fa6e7913000 r--p 00000000 08:02 28633
7fa6e7913000-7fa6e7914000 r--p 00000000 08:02 439407
7fa6e7914000-7fa6e791b000 r--s 00000000 08:02 446759
7fa6e791b000-7fa6e791c000 r--p 00000000 08:02 28635
7fa6e791c000-7fa6e791f000 rw-p 7fa6e791c000 00:00 0
7fa6e791f000-7fa6e7921000 rw-p 0001d000 08:02 81927
7fff8bc0f000-7fff8bc24000 rw-p 7fffffe9000 00:00 0
7fff8bd1d000-7fff8bd1f000 r-xp 7fff8bd1d000 00:00 0
ffffffffff600000-ffffffffff601000 r-xp 00000000 00:00 0
gandreu@shell-sisop:~$
/usr/lib/locale/es_ES.utf8/LC_CTYPE
/usr/lib/locale/es_ES.utf8/LC_COLLATE
/usr/lib/locale/es_ES.utf8/LC_NUMERIC
/usr/lib/locale/es_ES.utf8/LC_TIME
/usr/lib/locale/es_ES.utf8/LC_MONETARY
/usr/lib/locale/es_ES.utf8/LC_PAPER
/usr/lib/locale/es_ES.utf8/LC_NAME
/usr/lib/locale/es_ES.utf8/LC_ADDRESS
/usr/lib/locale/es_ES.utf8/LC_TELEPHONE
/usr/lib/locale/es_ES.utf8/LC_MEASUREMENT
/usr/lib/gconv/gconv-modules.cache
/usr/lib/locale/es_ES.utf8/LC_IDENTIFICATION
/lib/ld-2.7.so
[stack]
[vdso]
[vsyscall]
```

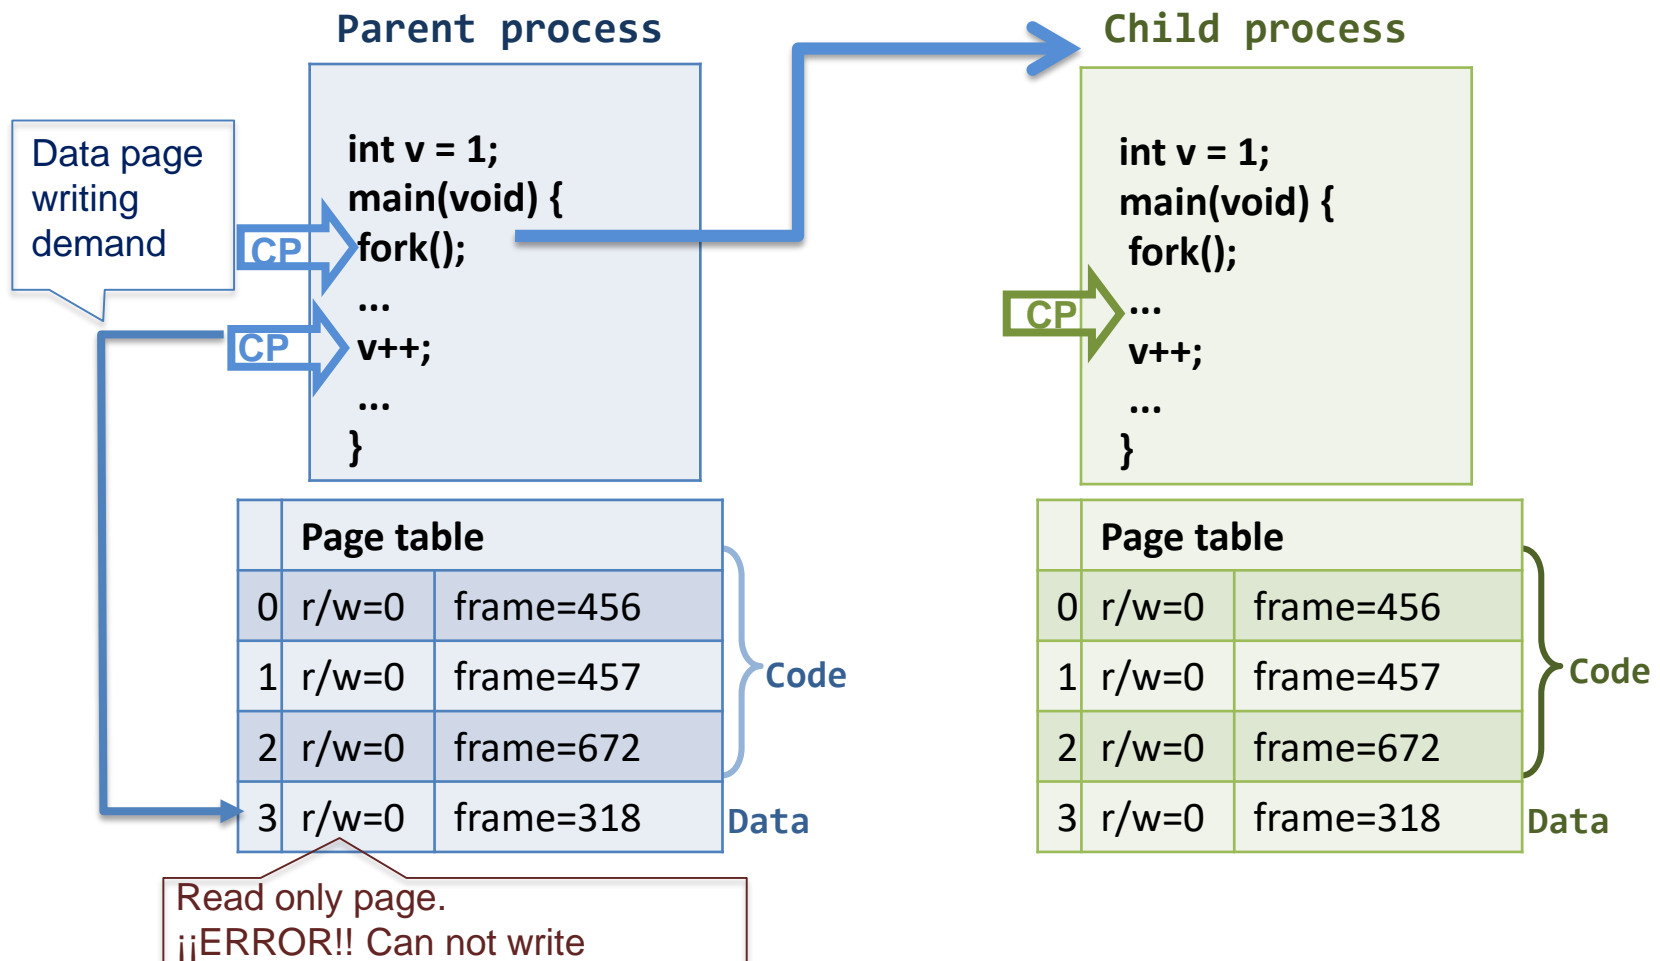

- **Copy-on-Write**

- Linux technique to **efficiently perform copies of memory pages** (it saves memory and time)
- When a process creates a new one, parent and child share data and stack pages in memory
 - All shared pages are marked as “read-only”
 - A write access attempt to these pages makes the MMU to send a page access failure interrupt, then:
 - The kernel does a copy of the troubling page to the process that wants to write it
 - » If there are more than two processes, the remaining processes continue being unable to write
 - » If it remains only one process using the page it will be able to write changing previously its descriptor bit
 - The interrupted instruction is restarted
- **Advantage:** Time and space required to copy unused pages are avoided

Memory map of a Linux process

- ***fork()*** call “Copy-on-Write”

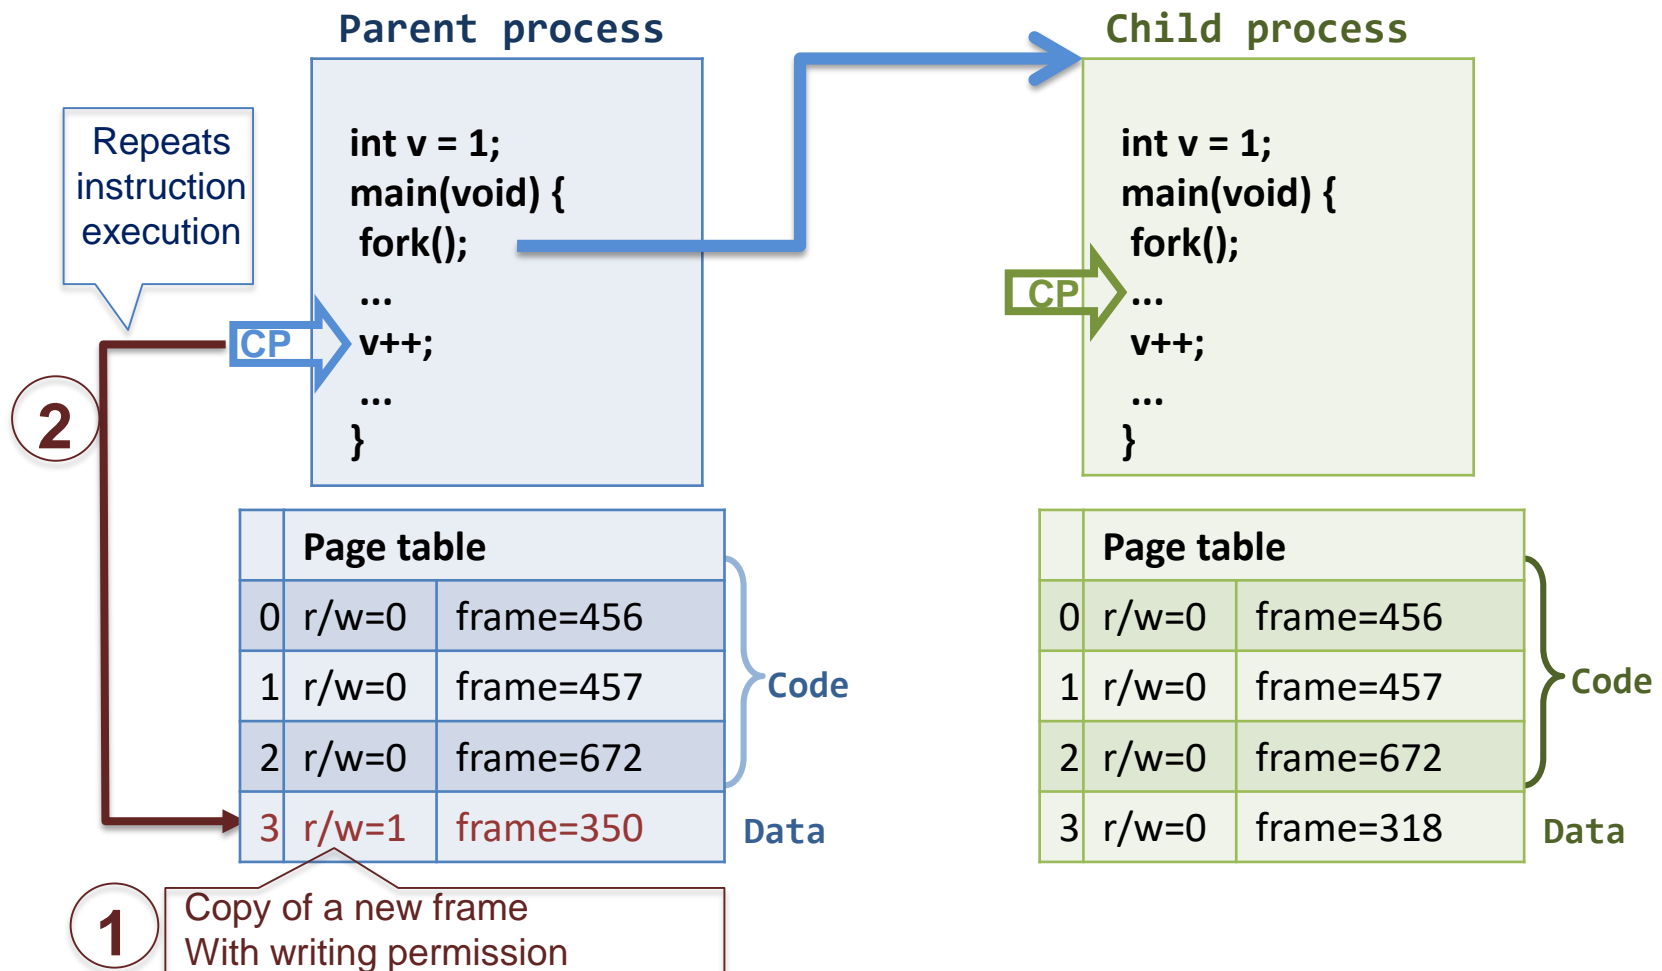
- **Remember:** when a process creates a new one, parent and child share data and stack pages in memory
 - All shared pages are marked as “read-only”
 - An write access attempt to these pages make the MMU to send a page access failure interrupt, then what happens is the following:



Memory map of a Linux process

- ***fork()* call “Copy-on-Write”**

- **Remember:** when a process creates a new one, parent and child share data and stack pages in memory
 - All shared pages are marked as “read-only”
 - An write access attempt to these pages make the MMU to send a page access failure interrupt, then what happens is the following:



- Introduction
- Memory map of a Linux process
- **Memory mapped files**
- Dynamic linking libraries

- **Memory mapped file**

- A file (whole or part) is included inside a process memory map

POSIX: `mmap()`

Win32: `CreateFileMapping()`

- Advantages:
 - **File access time improved**, once mapped into memory access time is set by memory speed instead of hard disk
 - **Intermediate copies avoided** the OS transfers data directly between mapped file memory region and file
- It is an alternative file access method instead of using I/O calls *read* and *write*
 - The file becomes an array of byte

POSIX call **mmap** creates a new region in the process memory map and some of its properties can be set, like sharing and permissions

```
caddr_t mmap (caddr_t addr, size_t length, int protec,  
              int indicator, int fd, off_t shift)
```

- **addr**: memory address for file mapping. If 0 the OS decides. mmap always returns the mapping address used
- **fd**: file descriptor for the file to map (it must be opened)
- **shift** and **length** define the region to map the file, it is from shift to shift+length-1
- **protec**: PROT_READ, PROT_WRITE, PROT_EXEC and combinations (i.e. PROT_WRITE | PROT_EXEC)
- **indicator**: MAP_SHARED, MAP_PRIVATE

munmap: removes a previous whole or part file mapping

```
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/mman.h>
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>
```

count_char.c It maps a file in memory and counts the number of appearances of a given character. The character to count is specified as the first parameter and the file to map as the second parameter.

```
$ gcc count_char.c -o count_char
$ ./count_char c count_char.c
```

```
int main(int argc, char *argv[] ) {
    int i,fd,count;
    char *p,*org;
    struct stat bstat;
    char ch;


    ch = argv[1][0];
    fd = open(argv[2], O_RDONLY); /* File open */
    fstat(fd, &bstat); /* Gets file length*/
    /* File mapping */
    org = mmap((caddr_t) 0, bstat.st_size, PROT_READ,MAP_SHARED, fd, 0);
    close(fd); /* File close */
    /* Access loop */
    p = org;
    count = 0;
    for (i=0; i<bstat.st_size; i++)
        if (*p++ == ch) count++;
    /* Remove mapping */
    munmap(org, bstat.st_size);
    printf("%d\n", count);
}
```

- Memory map before file mapping

```
pplanes$ pmap 21914:
21914   ./count_char c count_char.c
08048000    4K r-x-- /home/naomac/fso/count_char
08049000    4K rw--- /home/naomac/fso/count_char
b7de1000    4K rw--- [ anon ]
b7de2000  1316K r-x-- /lib/tls/i686/cmov/libc-2.7.so
b7f2b000    4K r---- /lib/tls/i686/cmov/libc-2.7.so
b7f2c000    8K rw--- /lib/tls/i686/cmov/libc-2.7.so
b7f2e000   12K rw--- [ anon ]
b7f42000   12K rw--- [ anon ]
b7f45000    4K r-x-- [ anon ]
b7f46000  104K r-x-- /lib/ld-2.7.so
b7f60000    8K rw--- /lib/ld-2.7.so
bf986000   84K rw--- [ stack ]
total   1564K
```

- Memory map after file mapping

```
pplanes$ pmap 21914
21914: ./count_char c count_char.c
08048000    4K r-x-- /home/naomac/fso/count_char
08049000    4K rw--- /home/naomac/fso/count_char
b7de1000    4K rw--- [ anon ]
b7de2000  1316K r-x-- /lib/tls/i686/cmov/libc-2.7.so
b7f2b000    4K r---- /lib/tls/i686/cmov/libc-2.7.so
b7f2c000    8K rw--- /lib/tls/i686/cmov/libc-2.7.so
b7f2e000   12K rw--- [ anon ]
b7f41000    4K r--s- /home/naomac/fso/count_char.c
b7f42000   12K rw--- [ anon ]
b7f45000    4K r-x-- [ anon ]
b7f46000  104K r-x-- /lib/ld-2.7.so
b7f60000    8K rw--- /lib/ld-2.7.so
bf986000   84K rw--- [ stack ]
total   1568K
```



Mapped file count_char.c

Memory mapped files

“**map.c**” maps a file into memory and it shows the process memory map before and after mapping the file which name is specified in the first program parameter

```
$ gcc map.c -o map
```

```
$ map map.c
```

```
int main (int argc, char *argv[])
{
    int fd;
    void *map;
    struct stat statbuf;
    char path_maps[80];

    // Open the file to map
    if (argc!=2) {
        puts("Usage: map FileName \n");
        exit(EXIT_FAILURE) ;
    }
    if ((fd=open(argv[1],O_RDONLY))<0)
        error("Open file failure (open) \n");

    // Get file length
    fstat(fd, &statbuf);
    // fstat dumps its information to statbuf

    // SHOW MAP
    printf(" PROCESS MEMORY MAP /proc/%d/maps \n", getpid());
    build_command(path_maps);
    system(path_maps); // Command execution system call
}
```

```
#include <sys/types.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <stdio.h>
```

```
void error (char * message) {
    perror(message);
    exit(EXIT_FAILURE);
}
```

```
void build_command(char command[80]) {
    // Build command to show memory map
    sprintf(command,"cat /proc/%d/maps",getpid());
}
```

continues...

“map.c” continuing...

```
// Map input file
if ((map = mmap(0, statbuf.st_size, PROT_READ, MAP_SHARED, fd, 0)) == MAP_FAILED)
    error("Mapping failure(mmap)");

close(fd); // Close file

// SHOW MAP
printf ("\n\n MEMORY MAPPED FILE \n");

system(path_maps); // Command execution system call

munmap(map, statbuf.st_size); // Remove mapping

printf ("\n\n MEMORY MAPPING MAP REMOVED \n");

system(path_maps);

exit(EXIT_SUCCESS);

} /* main end */
```

System call to execute the previously defined command:

cat /proc/%d/maps

It shows this process “maps” file

- Introduction
- Memory map of a Linux process
- Memory mapped files
- **Dynamic linking libraries**

- **Programming libraries**

- Binary (no text) files that contain functions code
- There are two way of linking programs with libraries:
 - **Static linking:** the executable file includes all library functions code
 - Program code = Own program code + Library functions code
 - **.lib files on Windows, .a files on UNIX/Linux**
 - **Dynamic linking:** the executable file contains references to library functions that it uses and the memory region required to store them
 - Library functions are loaded in memory on demand mapping them in the corresponding process memory map region. This is done by the library loader program (i.e. **ld** in Linux) by means of **dlopen** call
 - **Windows:** .dll files
 - **UNIX/Linux:** .so files

- **Static linking**

- Disadvantages
 - Generally big executable files
 - Library functions code replicated in many executables in the file system and in memory during execution
 - A library update requires rebuilding programs
- Advantages
 - Executable files are self contained

- **Dynamic linking**

- Advantages
 - Smaller executable size that saves disk and memory space
 - No library functions code replication
 - Processes share library code in memory
 - Library updates don't require rebuilding and several library version can coexists
- Disadvantages
 - Executable files rely in library files
 - Dynamic linking introduces execution time overhead

- `ejemplo1` executable size with static and dynamic linking of math library
 - Static linking

```
pblanes$ gcc ejemplo1.c -static -o ejemplo1 -lm
pblanes$ ls -l
total 660
-rwxr-xr-x 1 pblanes disca-upvnet 670227 2011-10-20 15:54 ejemplo1
-rw-r--r-- 1 pblanes disca-upvnet 905 2011-10-20 13:30 ejemplo1.c
```

- Dynamic linking (gcc default)

```
pblanes$ gcc ejemplo1.c -o ejemplo1 -lm
pblanes$ ls -l
total 16
-rwxr-xr-x 1 pblanes disca-upvnet 10301 2011-10-20 15:56 ejemplo1
-rw-r--r-- 1 pblanes disca-upvnet 905 2011-10-20 13:30 ejemplo1.c
```