

Fundamentos de los Sistemas Operativos (FSO)

Departamento de Informática de Sistemas y Computadoras (DISCA)
Universitat Politècnica de València

Part 3: Memory management

Unit 9

Virtual memory (I)

fSO

DISCA



UNIVERSITAT
POLITÈCNICA
DE VALÈNCIA

- **Goals**

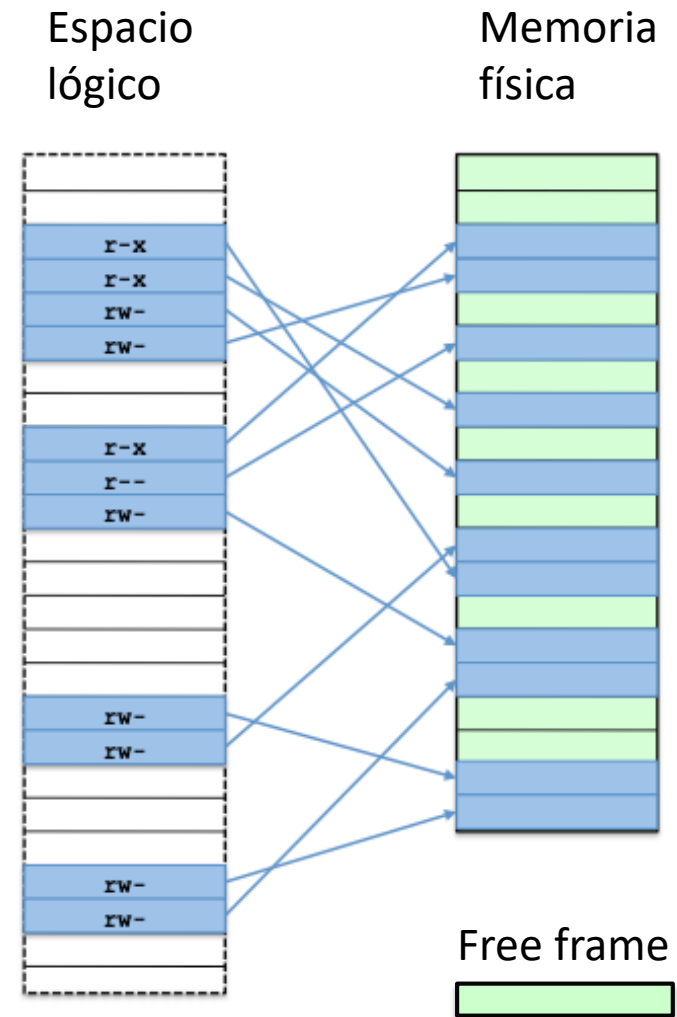
- To know the **advantages** of virtual memory and its effects on **system performance**
- To understand **demand paging concept**
- To know the **page replacement techniques**

Bibliography

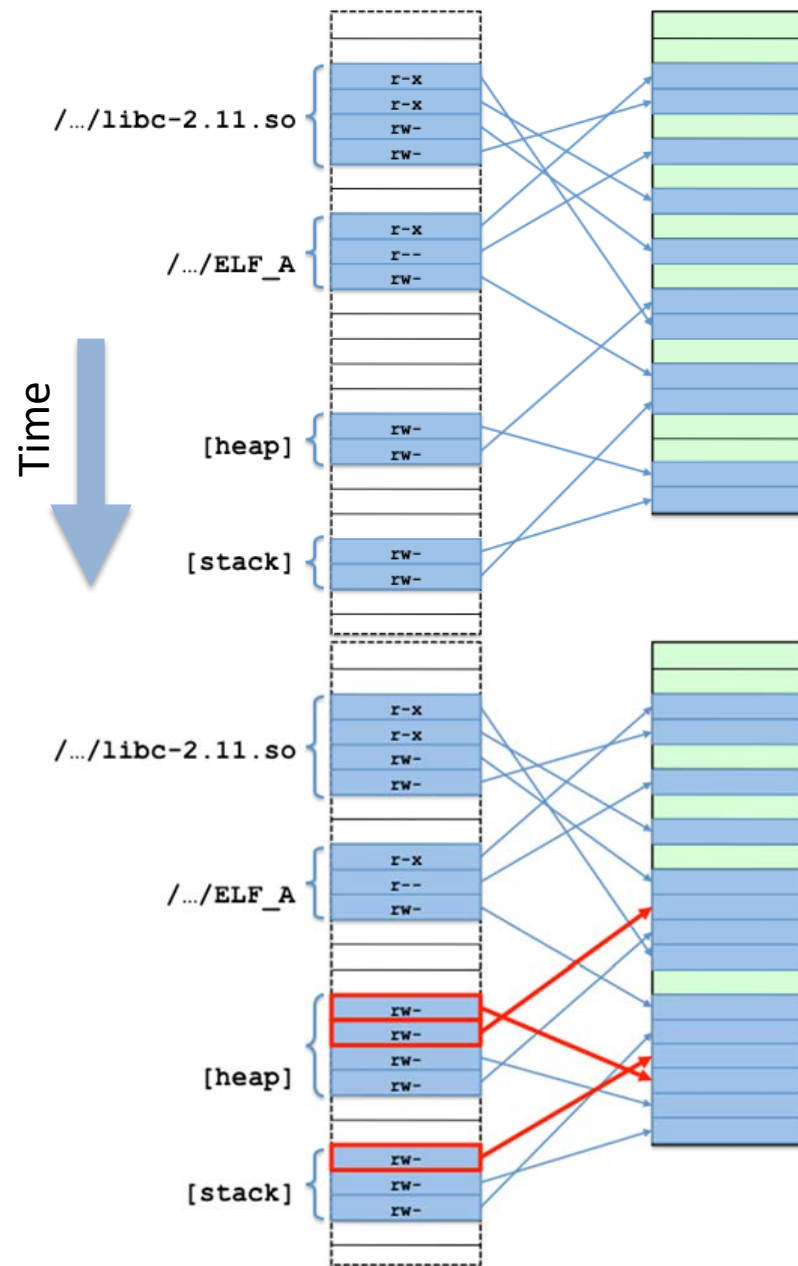
- Silberschatz, chapter 9

- **Virtual memory concept**
- Virtual memory support
- Demand paging
- Virtual Memory and process management
- FIFO replacement algorithm
- Optimal replacement algorithm
- LRU replacement algorithm

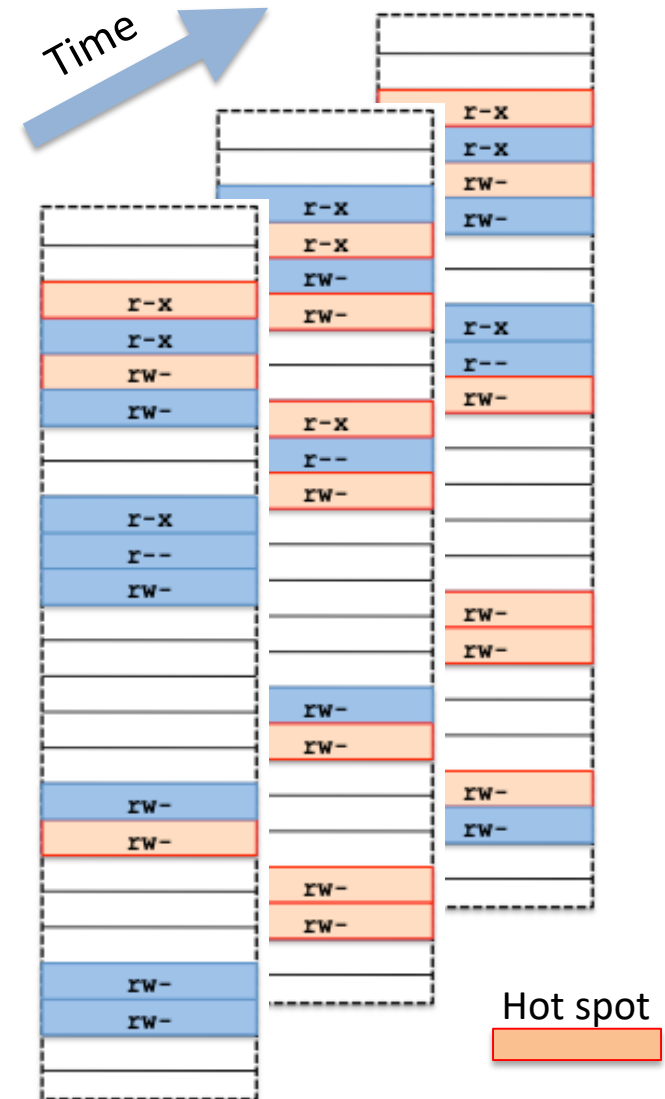
- Paging without virtual memory
 - The OS reserves **all the physical memory** required by a process to start it
 - All memory accesses take the same amount of time
 - When a process ends the OS releases all the frames used by the process



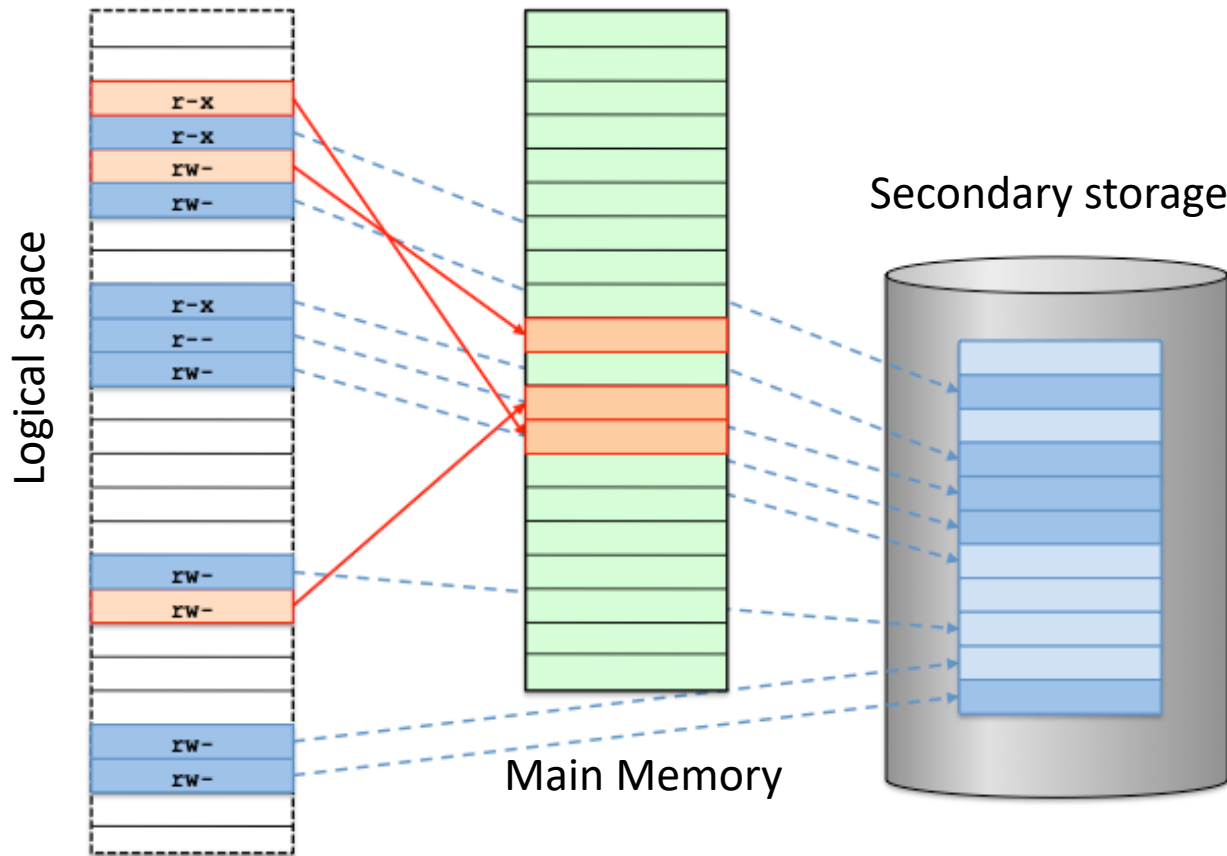
- Memory requirements change due to dynamic regions
 - The stack grows due to function calls
 - Dynamic memory allocation creates the *heap* region that can grow and shrink due to *malloc* and *free* calls
 - Creating new threads produce new thread stack regions



- **Reference locality principle**
 - Independently of the logical size of a process its memory accesses are localted into **hot spots**
 - Given a time interval a small set of instructions is accessed particularly inner loops, something similar happens with data access
 - Along the process lifetime hot spots of code and data move



- **Memory management based on virtual memory**
 - The OS manages memory allocation to processes in such a way that **only their hot spots are allocated on physical memory**
 - The **remaining logical space** content is allocated on **secondary storage** (swapping area)



- **Virtual memory base technologies**

- It combines physical memory (RAM) with secondary storage (hard disk or SSD)
- Main memory is made of words (i.e. 32-bit, 64-bit) addressable by the CPU, access time is a couple of nanoseconds
 - Every instruction cycle performs one or more accesses to main memory
- Secondary storage is made of blocks (i.e. 512-byte, 4096-byte), access time is a couple of milliseconds ($1\text{ ms} = 10^6\text{ ns}$)
 - A page transfer is made in one single I/O operation that requires the execution of several instructions by the CPU

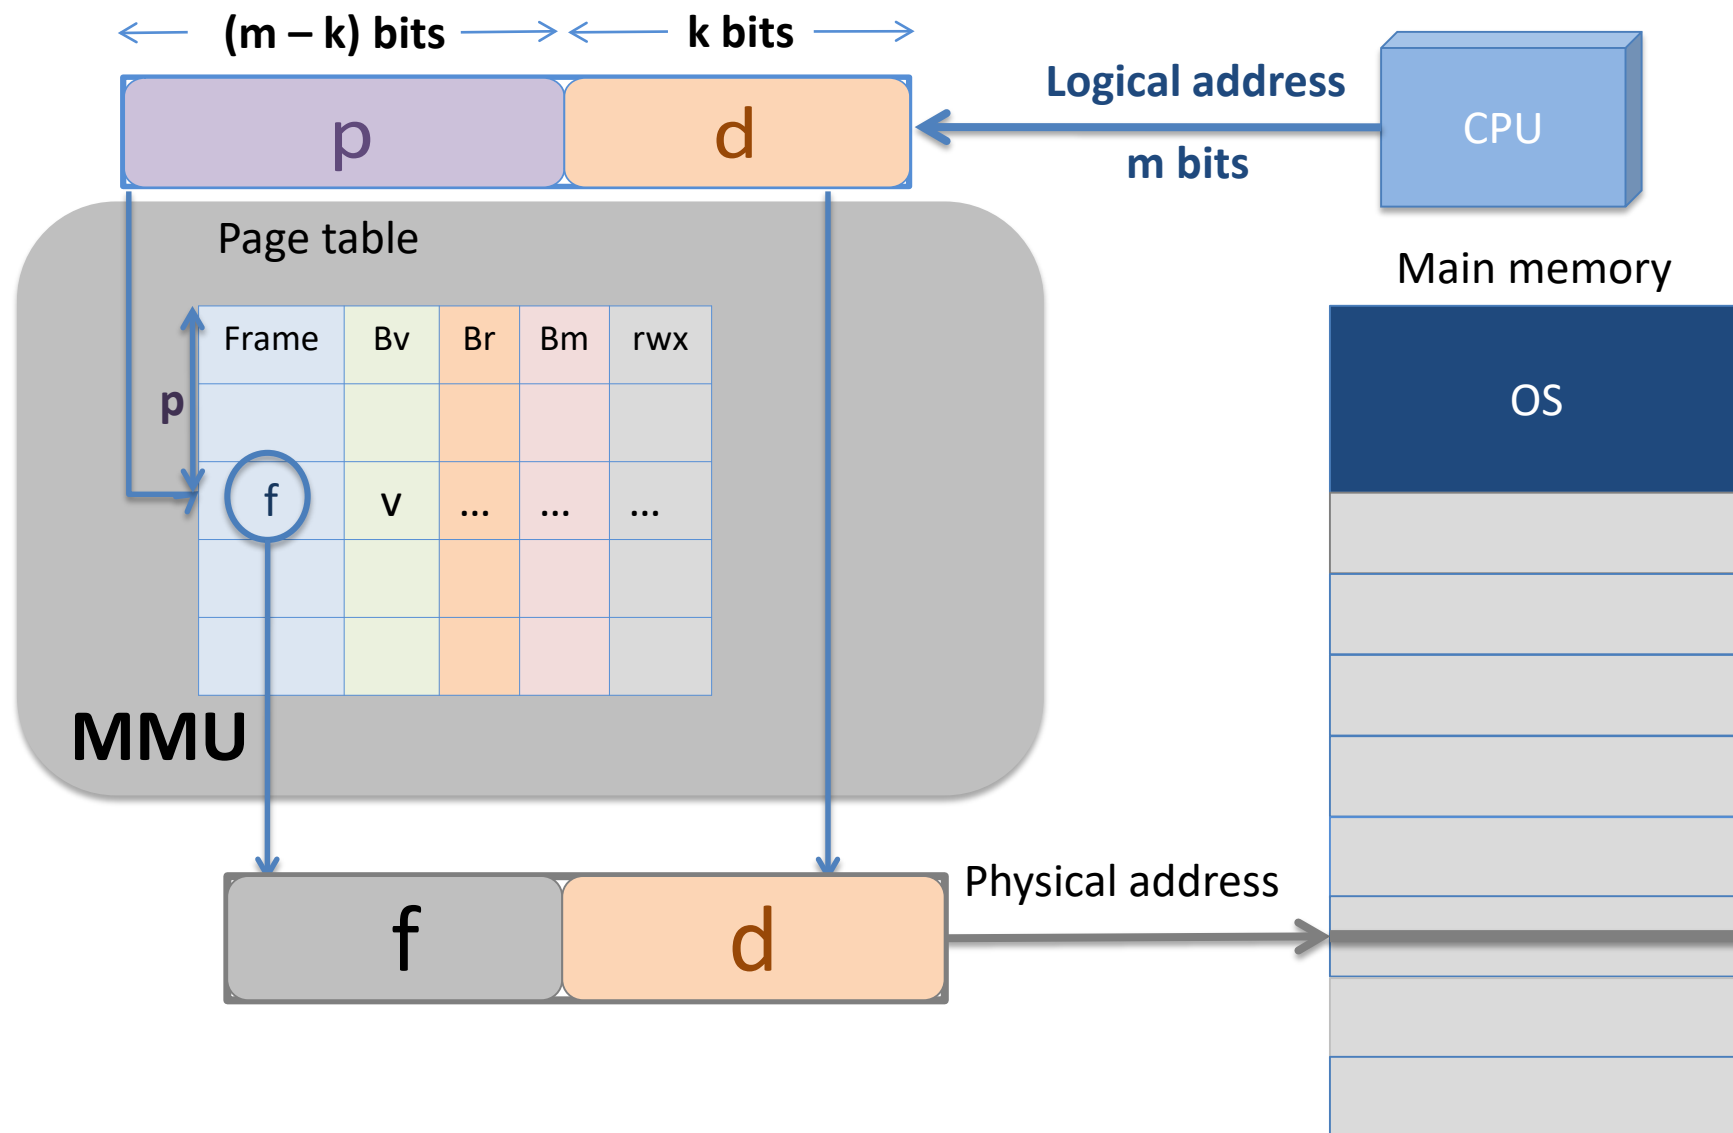
- **Virtual memory scheme**

- The OS manages memory following a sparse allocation approach, typically paging
- A swap area is available on disk, as a partition or as a system file
- For every page in use by a process it can be in two states:
 - **Valid:** page allocated into a main memory frame
 - **Invalid:** page not allocated in main memory but on the swap area
- For every memory access it can happen a:
 - **Hit** (most common): reference to a valid page
 - **Fault:** reference to an invalid page -> the page table has to be updated and one or more pages are transferred between main memory and the swap area

- Benefits:
 - The same as paging: **page sharing and protection**
 - **It saves memory** and increases the multiprogramming level
 - It **allows bigger executable process size**
 - It **easies dynamic memory** management
- Penalties:
 - **Turnaround time** can increase due to page faults
 - **Workload of secondary storage** increases
 - Greater OS design **complexity**


- Virtual memory concept
- **Virtual memory support**
- Demand paging
- Virtual Memory and process management
- FIFO replacement algorithm
- Optimal replacement algorithm
- LRU replacement algorithm

- Address translation as paging (UT08)



- It is every page table entry
- Every page descriptor has the following fields:
 - **Frame number** where the page is allocated in physical memory
 - **Valid bit**: it indicates if a page is mapped in memory, it supports demand paging
 - **Reference bit**: it indicates page access done. It is required for a second opportunity algorithm
 - **Modified bit**: it indicates page write access done. Trouble with shared pages.
 - **Access mode bits**: read only, read-write, execution, ...

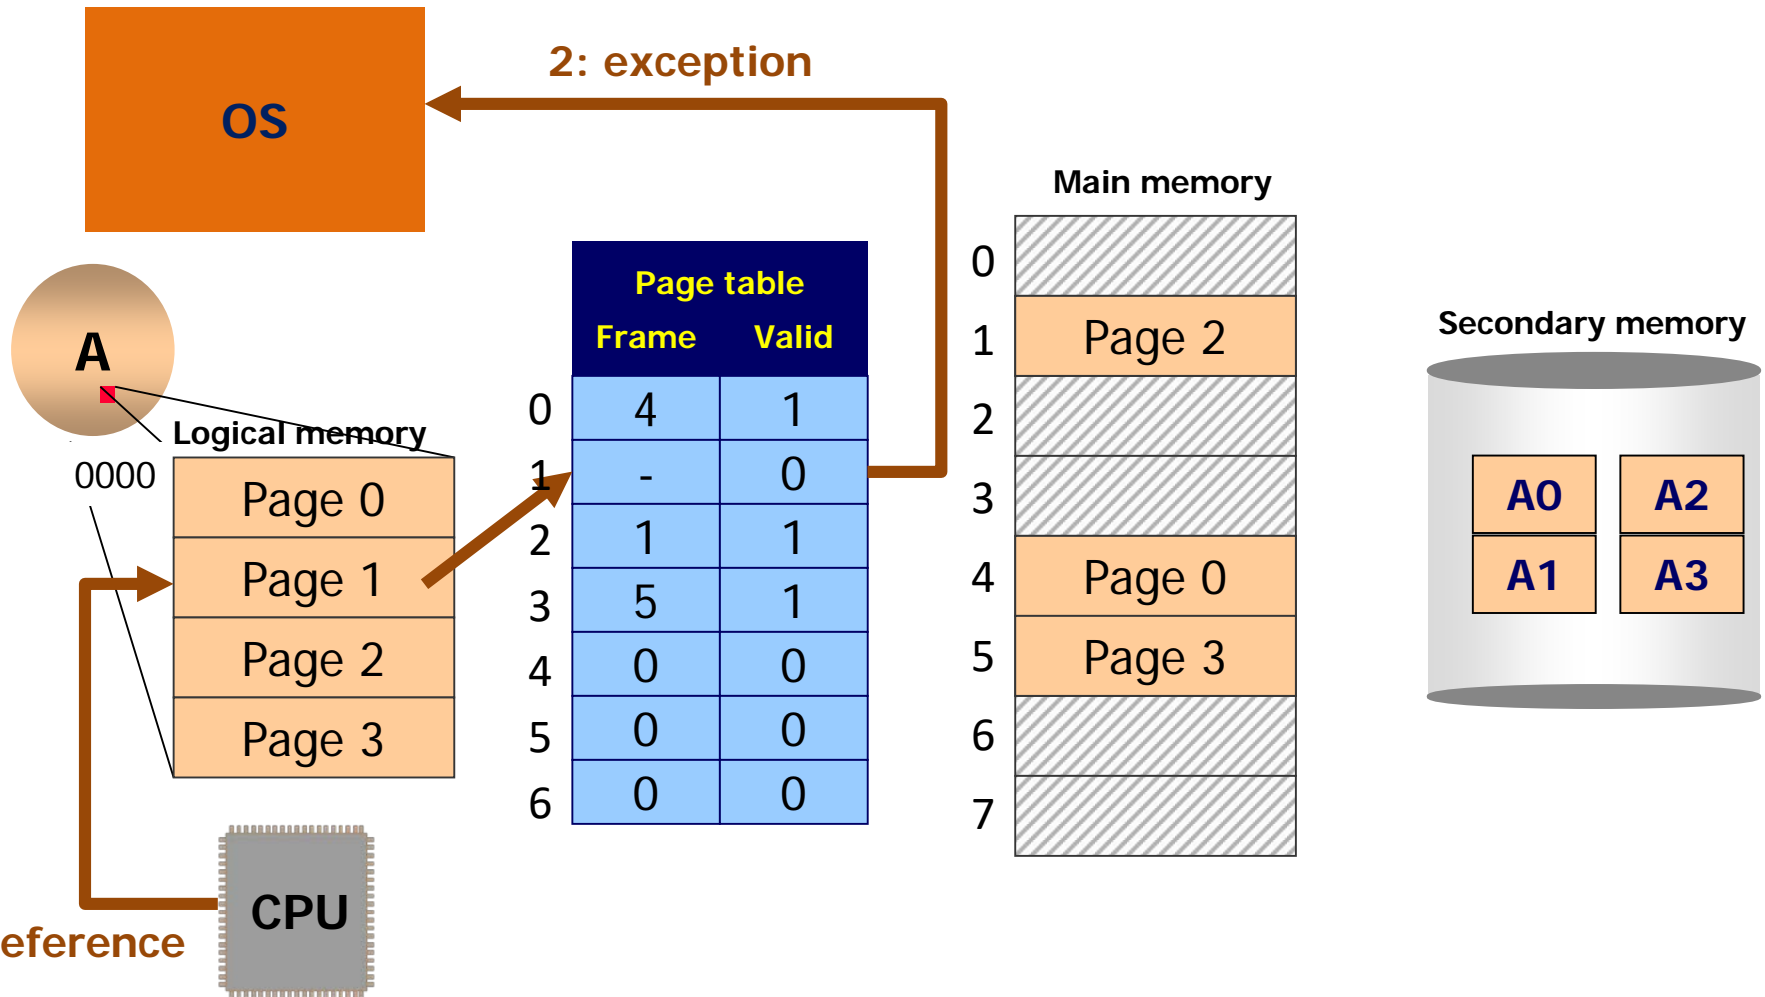
Page descriptor



- **Page fault -> MMU exception**
 - It happens when a **non memory allocated page is referenced**, that is **with valid bit=0**
 - Page fault cases:
 - **Page on disk:** The reference page belongs to the process → it is allocated into a memory frame
 - **Access error:** The reference page doesn't belong to the process and cannot be assigned to it → the process is aborted
 - **Process growing:** The process asks for new pages, if the OS permits it, a new page is assigned to the process with its valid bit set and the new page is allocated into a memory frame

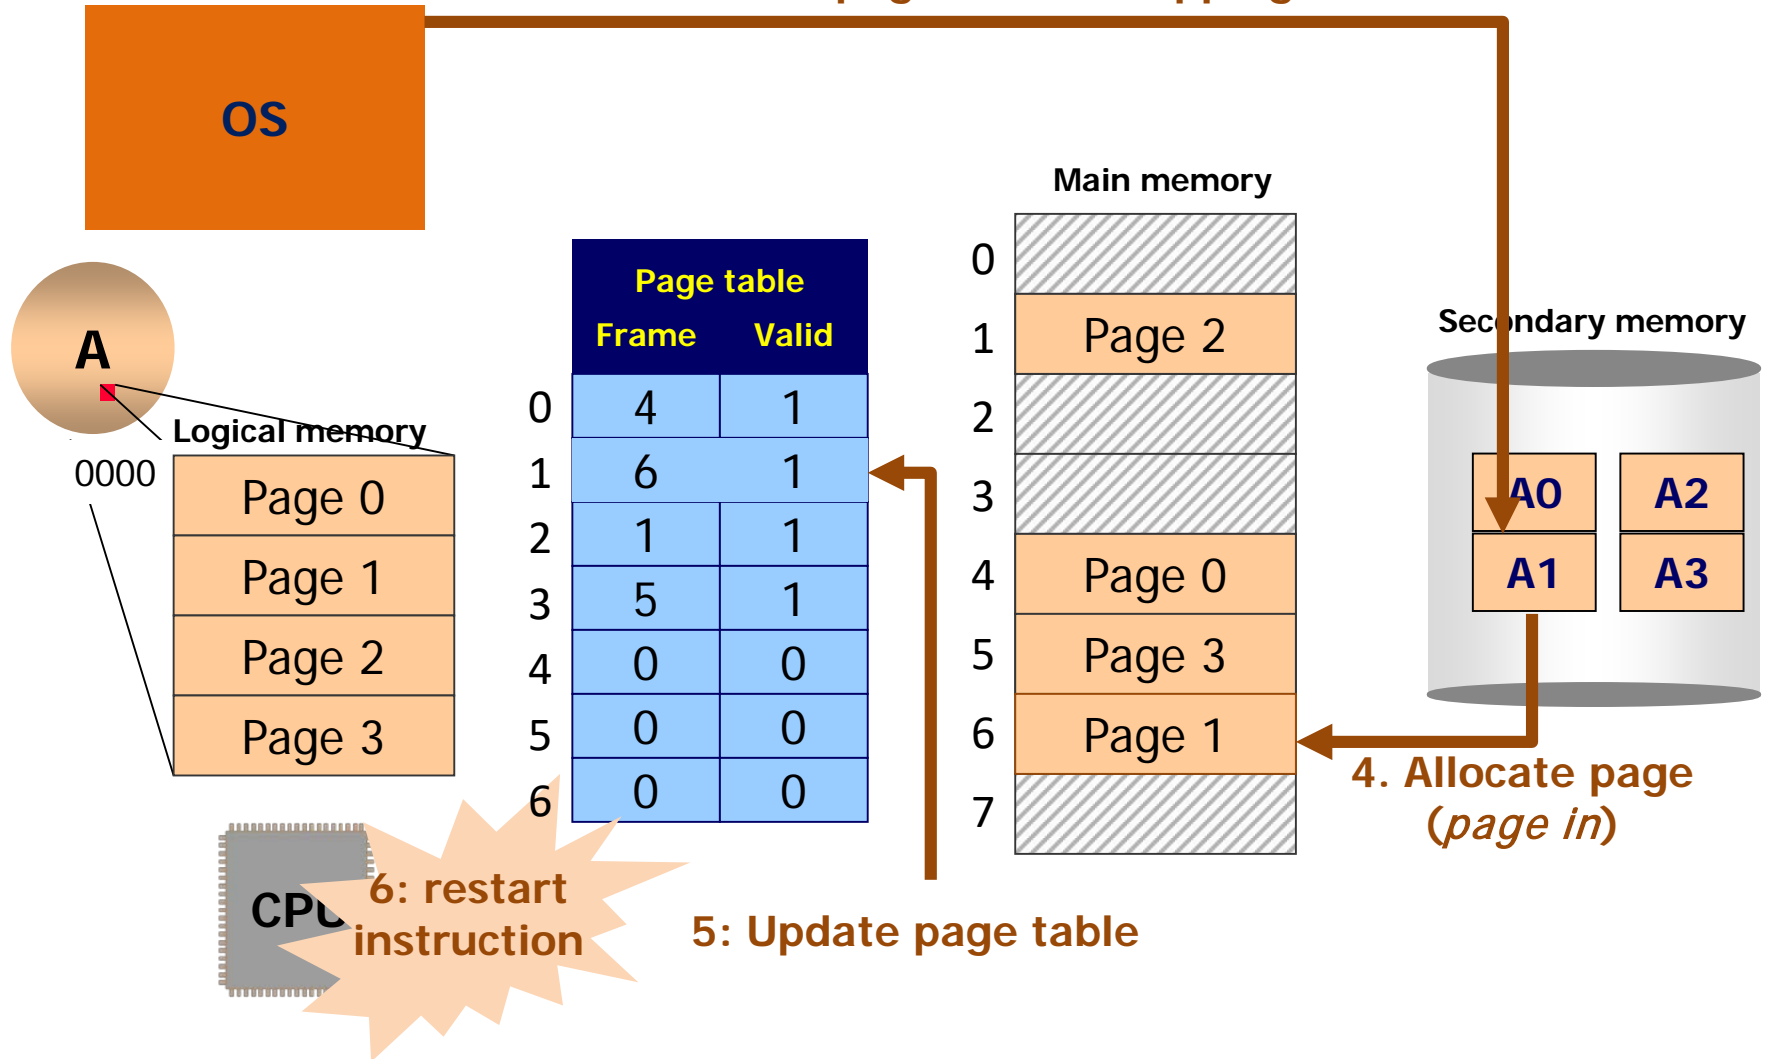
- Virtual memory concept
- Virtual memory support
- **Demand paging**
- Virtual Memory and process management
- FIFO replacement algorithm
- Optimal replacement algorithm
- LRU replacement algorithm

- Page fault: Page on disk case



- Page fault: Page on disk case

3: look for the page in the swapping area



- **Page fault algorithm: page on disk case**
 - Find demanded page on disk
 - Find a free frame:
 - If there is a free frame, then use it
 - If there are no free frames apply **page replacement algorithm** (next slide)
 - **Read demanded page on disk** (page in) and allocate it into the free frame
 - **Update page table** of the process that generates the page fault
 - Update free frame table
 - Transfer control to user process
 - Restart the instruction that has produced the page fault


- **Page replacement:** It is required when main memory is full and there is a page fault
 - **Select an allocated page in main memory, named victim, to leave its frame**
 - There are several victim selection algorithms
 - If the victim page has its modified bit equal to 1 then **save the victim page to disk (page out)**
 - The victim page entry on the page table is updated with valid bit = 0
 - Update the free frame table

- **Reference string:** Sequence of accessed pages along a certain time period
 - From every logical address sent by the processor its page number is obtained
 - Repetitions are removed: several consecutive references to the same page are replaced by a single reference
- **Example:**
 - Page size = 1000 words
 - Referenced logical addresses from process P
2500, 5100, 5234, 1800, 1432, 4388, 2124, 8216, 8498
 - Referenced pages
2, 5, 5, 1, 1, 4, 2, 8, 8
 - The corresponding reference string is:
2, 5, 1, 4, 2, 8

- Virtual memory concept
- Virtual memory support
- Demand paging
- **Virtual Memory and process management**
- FIFO replacement algorithm
- Optimal replacement algorithm
- LRU replacement algorithm

- Virtual memory **favors concurrency**
 - It keeps in memory only processes hot spot areas, so memory capacity is optimized.


Example. Process A has 11 pages and B has 8 pages, but actually they only use 10 frames.



- Virtual memory allows **frame sharing**
 - Executable code pages from common dynamic libraries are shared between processes, as well as shared variables and mapped files

Example. Processes A and B, created from executable files ELF_A and ELF_B, share four allocated pages:

- Three pages from library libc-2.11.so
- Un page without support with shared variables




- **fork** and **exec** calls
 - An *exec* call invalidates all process pages
 - Coming page faults will update page descriptors with the new executable code
 - A *fork* call clones the parent's page descriptor table into the new process
 - Parent and child access read only pages (executable code and constants) on shared frames
 - Pages with writing permissions (i.e. variables, stack and heap) will be allocated on different frames, following *copy-on-write* policy: the new frame allocation is done when the first writing operation happens


- Virtual memory concept
- Virtual memory support
- Demand paging
- Virtual Memory and process management
- **FIFO replacement algorithm**
- Optimal replacement algorithm
- LRU replacement algorithm

FIFO replacement algorithm

- Victim page: the one that is longer loaded in memory
- It suffers from Belady's anomaly



With 3 frames: 9 faults (6 faults with replacement)




With 4 frames: 10 faults (6 faults with replacement)


FIFO replacement algorithm

- Victim page: the one that is longer loaded in memory
- It suffers from Belady's anomaly

Belady's anomaly



With 3 frames: 9 faults (6 faults with replacement)



With 4 frames: 10 faults (6 faults with replacement)

- **Stack algorithms**

- A page set kept into N frames is **always** a subset of the one kept into $N+1$
- Do not suffer from Belady's anomaly

- Virtual memory concept
- Virtual memory support
- Demand paging
- Virtual Memory and process management
- FIFO replacement algorithm
- **Optimal replacement algorithm**
- LRU replacement algorithm

Optimal replacement algorithm

- Victim page: the one which take longer to be referenced
- Minimum number of faults -> impossible to implement (future is unknown)

		1	2	3	4	1	2	5	1	2	3	4	5
Frame													
0		1	1	1	1	1	1	1	1	1	3	3	3
1			2	2	2	2	2	2	2	2	2	4	4
2				3	4	4	4	5	5	5	5	5	5

With 3 frames: 7 faults (4 faults with replacement)

		1	2	3	4	1	2	5	1	2	3	4	5
Frame													
0		1	1	1	1	1	1	1	1	1	1	4	4
1			2	2	2	2	2	2	2	2	2	2	2
2				3	3	3	3	3	3	3	3	3	3
3					4	4	4	5	5	5	5	5	5

With 4 frames: 6 faults (2 faults with replacement)

- Virtual memory concept
- Virtual memory support
- Demand paging
- FIFO replacement algorithm
- Optimal replacement algorithm
- **LRU replacement algorithm**

- Virtual memory concept
- Virtual memory support
- Demand paging
- Virtual Memory and process management
- FIFO replacement algorithm
- Optimal replacement algorithm
- **LRU replacement algorithm**

LRU replacement algorithm

- Victim page: the one that lasted more without being referenced or **least recently used**
- It is a stack algorithm

		1	2	3	4	1	2	5	1	2	3	4	5
Frame													
0		1	1	1	4	4	4	5	5	5	3	3	3
1			2	2	2	1	1	1	1	1	1	4	4
2				3	3	3	2	2	2	2	2	2	5

With 3 frames: 10 faults (7 faults with replacement)

		1	2	3	4	1	2	5	1	2	3	4	5
Frame													
0		1	1	1	1	1	1	1	1	1	1	1	5
1			2	2	2	2	2	2	2	2	2	2	2
2				3	3	3	3	5	5	5	5	4	4
3					4	4	4	4	4	4	3	3	3

With 4 frames: 8 faults (4 faults with replacement)

- LRU implementations
 - Using counters:
 - Every page has an associated counter
 - Using an ordered list of referenced pages
 - When a page is referenced it is put at the end
 - The victim page is the first one
- Performance analysis
 - Advantages
 - Good approach to the optimal algorithm
 - Disadvantages
 - Costly implementation → requires **hardware support**
 - Solution
 - LRU approximation algorithms → based on **reference bit**