#### Fundamentos de los Sistemas Operativos (FSO)

Departamento de Informática de Sistemas y Computadoras (DISCA) *Universitat Politècnica de València* 

Part 3: Memory management

Unit 10
Virtual memory (II)





#### Goals

- To understand 2nd chance page replacement algorithm as an LRU approximation
- To know the thrashing problem and its solutions
- To analyse memory frame management techniques

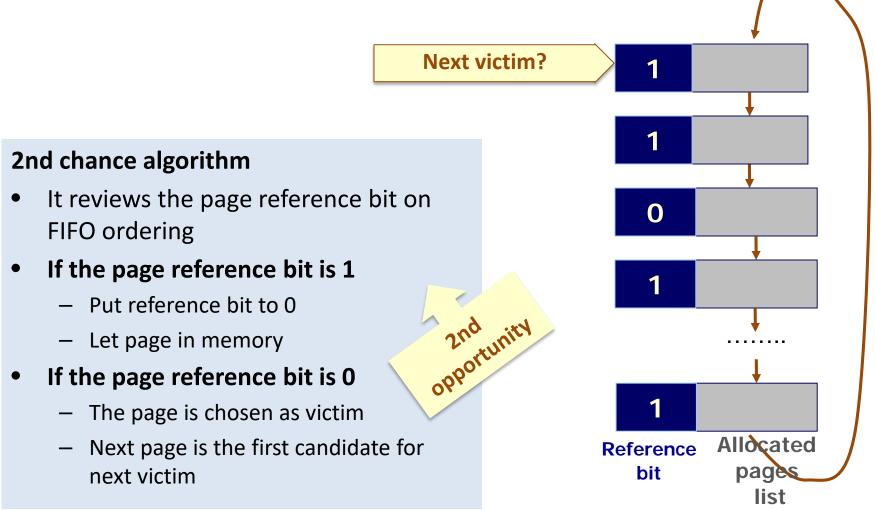
## Bibliography

- Silberschatz, chapter 9

- 2nd chance replacement algorithm
- Frame allocation
- Thrashing
- Frame reservation

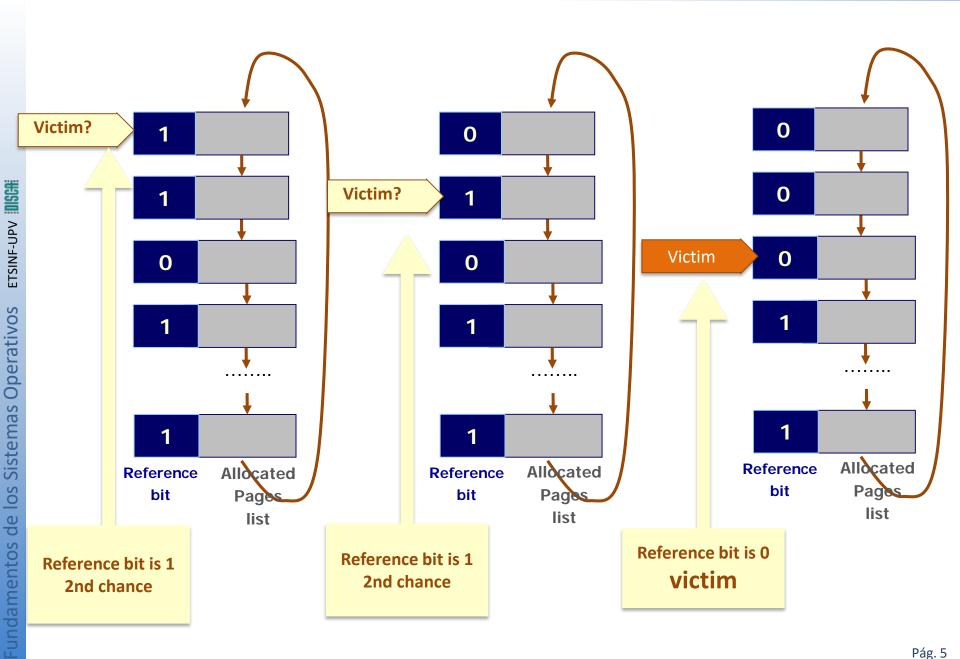
## 2nd chance replacement algorithm

- Supporting LRU as replacement algorithm is too expensive
- Solution: To approximate LRU using the reference bit



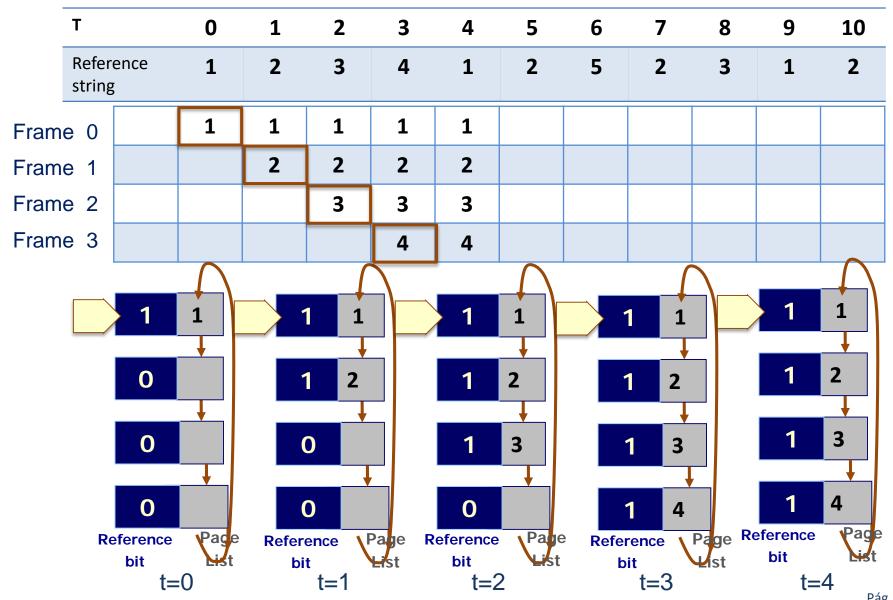
**fSO** 

## 2nd chance replacement algorithm



# 2nd chance replacement algorithm


## Example: Main memory with 4 frames



# ETSINF-UPV IIISCH Fundamentos de los Sistemas Operativos

# 2nd chance replacement algorithm


## Example: Main memory with 4 frames



# ETSINF-UPV IISCH Fundamentos de los Sistemas Operativos


### Example: Main memory with 4 frames


2nd chance replacement algorithm



## Example: Main memory with 4 frames

2nd chance replacement algorithm





Total page faults 6 (with replacement 2)

- 2nd chance replacement algorithm
- Frame allocation
- Thrashing
- Frame reservation

- Frame allocation problem
  - Free frame list:
    - Frame management requires a data structure where free frames are kept
  - Frame to process delivery policy and the OS
    - The OS gets the required number of frames to execute itself
    - Processes receive the minimum initial number of frames and the remaining ones on demand
    - The minimum number of initally assigned frames depends on the indirection level in the CPU instruction set → To execute an instruction all its operands must be allocated in main memory

- Frame delivery policies given m frames and n processes
  - Fair allocation: All processes allocate  $A_i$  frames equally  $A_i = m/n$
  - Proportional allocation: A process P<sub>i</sub> with size S<sub>i</sub> allocate A<sub>i</sub> frames computed as:

$$A_{i} = \frac{S_{i}}{\sum S_{i}} * m$$

Priority allocation: higher priority processes allocate more frames

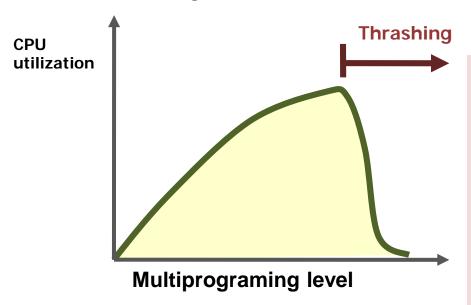
## Replacement policy scope

- Local replacement: only pages allocated to the process that generates the page fault can be replaced
  - It cannot choose a victim from another process
  - The number of process allocated frames does not change
  - A process execution is not affected by the remaining processes
    - Advantage: Sensible response time
    - Disadvantage: Worse global memory management
- Global replacement: the victim is chosen between all allocated pages in main memory
  - The victim can belong to another process
  - The number of process allocated frames can change
    - Disadvantage: Response time sensitive to system load
    - Advantage: Better global memory management

- 2nd chance replacement algorithm
- Frame allocation
- Thrashing
- Frame reservation

#### The thrashing problem:

- Memory becomes scarce and processes generate a lot of page faults
   → I/O time becomes dominant
- The OS allows more processes entering for execution regarding the low level of CPU use
- This makes things worse because the same amount of memory has to be shared with more and more processes → page fault rate keeps increasing



#### How to solve thrashing?

- —To anticipate an to prevent the problem
  - Using a working set model
  - –Controlling page fault rate
- Once thrashing is detected
  - Swap out processes with a medium term scheduler

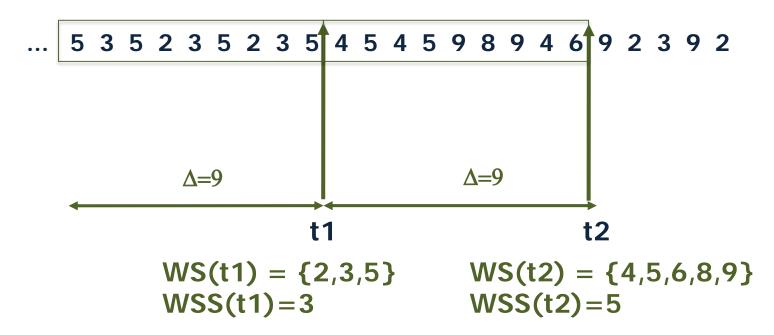
- Reference locality principle
  - Instructions and data processed recently (and the ones close to them) have a high probability of being processed in the near future
  - Locality:
    - Set of pages that a process uses as a whole
    - It is hard to identify
  - Thrashing happens when
    - locality sizes > total main memory size

- Working set model: preventive technique
  - It assumes reference locality principle
  - Obtain the number of pages that a process needs to have simultaneously in memory to avoid thrashing
  - Working set (WS):
    - Set of pages accessed in a process last referenced logical addresses
    - Working set window:
      - Fixed consecutive number of references  $\Delta$  used to compute WS
      - WS is made of the set of pages accessed in the last  $\Delta$  references
    - Working set size (WSS): Number of different pages that belong to the WS
    - In a system with **m frames and n processes**  $P_1...P_n$  there is thrashing when  $\sum WSS_i > m$

#### Example of WS model


- WS window  $\Delta = 9$
- Compute WS and WSS in t1 and t2

#### Sequence of referenced pages



#### Page fault rate control

 Preventive technique that analyses directly the page fault rate to guess if thrashing is near to happen



- 2nd chance replacement algorithm
- Frame allocation
- Thrashing
- Frame reservation

#### Concept

 Modern OS keep a percentage of main memory as a store of free frames (reservation frames)

#### Goals

- To reduce the time taken to serve a page fault
  - Attempt to have free frames available
  - The replacement algorithm is used:
    - Only when the free frame level gets too low
    - To look for victims to amortize its use
  - Page out
    - Several pages are written at once to disk
- To avoid thrashing

### Reservation frame management

- The OS guarantees that there will be always a set of free frames
- Some thresholds are set:
  - Minimum number of free frames (M<sub>MIN</sub>)
  - Recommended number of free frames (M<sub>RFC</sub>)

$$M_{REC} >> M_{MIN}$$

- Very efficient replacement algorithms are NOT required
  - The first VMS systems used FIFO because their MMU did not have reference bit
  - It is common to use a 2nd chance algorithm (Windows, UNIX SVR4, UNIX 4.4BSD, Linux, HP OpenVMS, ...).

## Monitor process

- There is an internal process that periodically accounts for the number of free frames (frame\_free):
  - If frame\_free > Mrec then do nothing
  - If **frame\_free < Mmin** then:
    - Swap out some processes until reaching REC free frames
    - Victims are process that spent more time suspended
  - If Mmin <= frame\_free <= Mrec then:</li>
    - Seek for processes with too many frames (very low frame rate) to "steal" them some frames applying a replacement algorithm
    - Several victims are selected in every process, the actual number differs on every OS

## Frame reservation management in thrashing

- When thrashing happens the number of free frames decreases quickly
- The process monitor detects it when:
  - Between two monitor activation **frame\_free** decreases too much
- Solutions:
  - Swap out whole processes until reaching frame\_free = Mrec
    - OpenVMS and Windows NT
  - Free a constant number of frames in every monitor activation if frame\_free < Mrec</li>
    - The monitor activation frequency increases
    - UNIX SVR4

#### Reservation frames content

- The frame of a victim selected by the OS is not allocated immediately to another page, instead the frame goes to the reservation stock
  - If victim pages are referenced again soon by the process:
    - There is a high probability that they be in the reservation stock and then they can be relocated without having to read them on disk
    - The OS remembers which is the content of every frame in the reserve stock
  - If frames included in the reservation stock correspond to modified pages:
    - They are not considered to solve page faults immediately because its content has to be written on disk (into a file or paging area)
      - » When a threshold is reached all these pages are written at once
      - » Page out overhead is amortized -> the number of global page writing is minimized
    - After writing on disk, frames become free and can serve page faults