

Fundamentos de los Sistemas Operativos (FSO)

Departamento de Informática de Sistemas y Computadoras (DISCA)
Universitat Politècnica de València

Part 4: File systems and I/O

Unit 11

Implementing file systems

fSO

DISCA



UNIVERSITAT
POLITÈCNICA
DE VALÈNCIA

- **Goals**

- To know file system functionality
- To know the abstraction levels of a file system
- To understand a file as a secondary storage abstraction and the access ways to its content
- To analyse the file block allocation techniques

- **Bibliography**

- Silberschatz, chapters 10 and 11

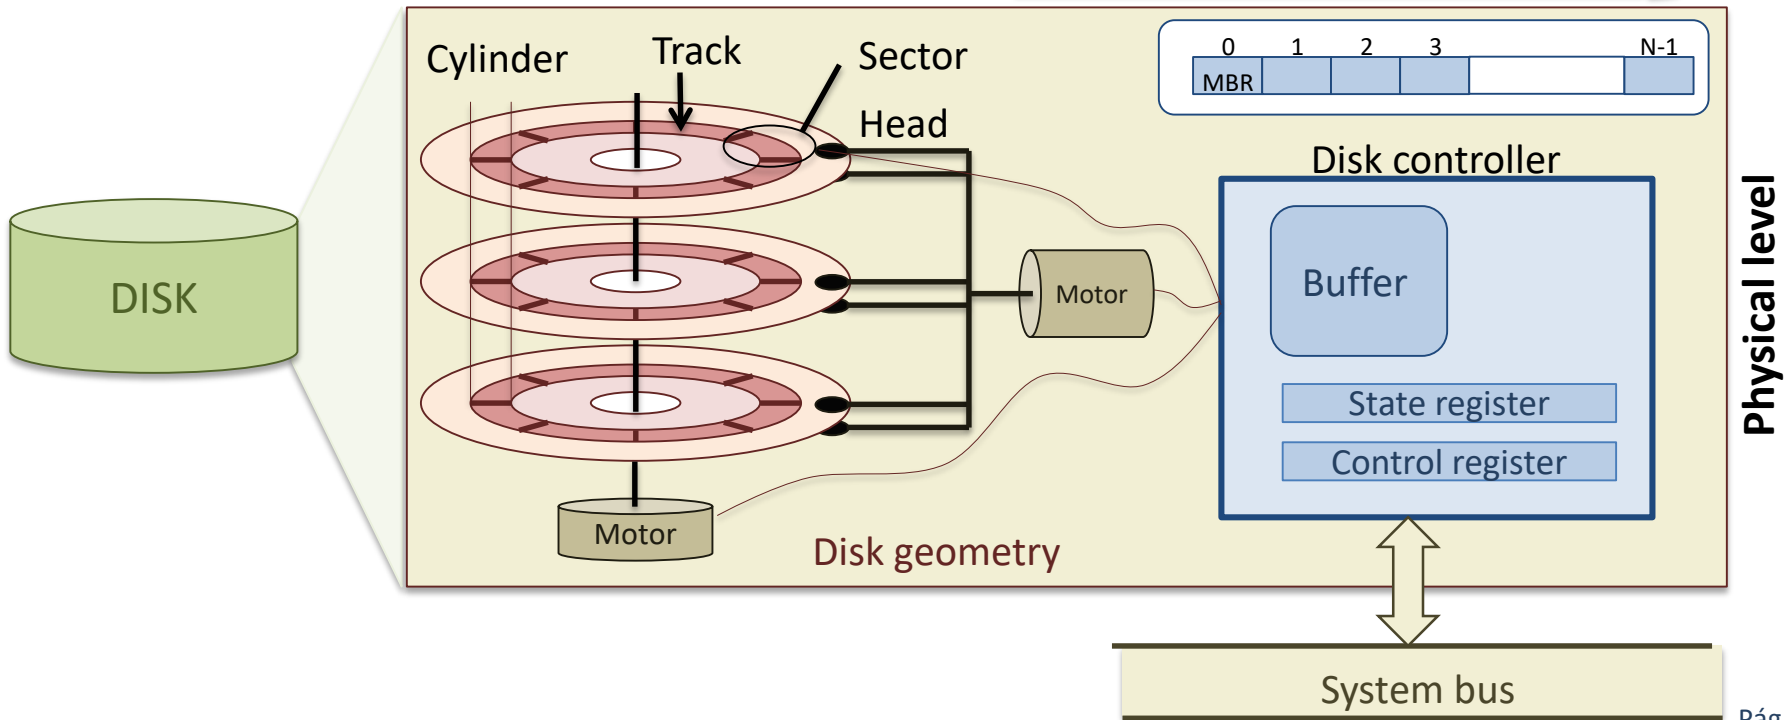
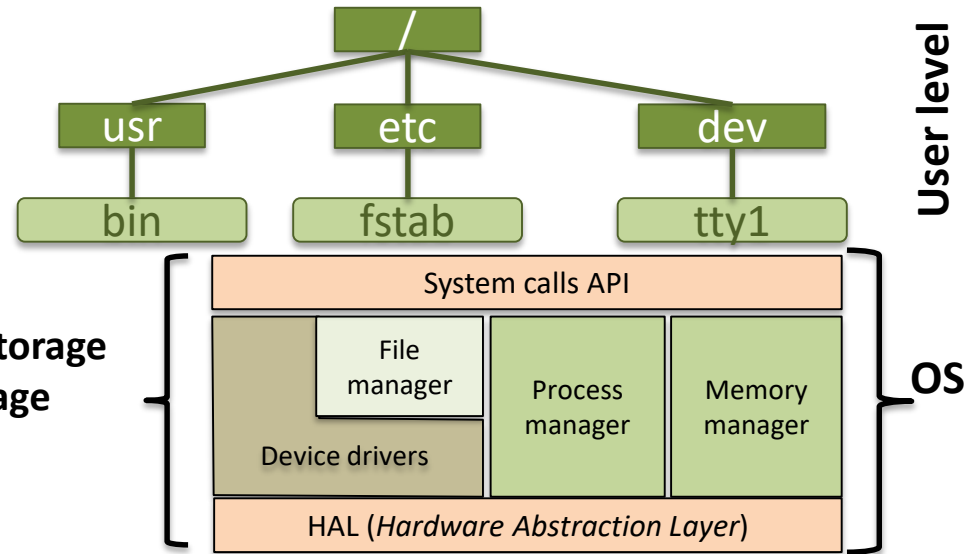
- **File system architecture**
- File concept
- File block allocation

- File system architecture has three aspects:
 1. **How files are organized** from the user point of view
 2. **How files are stored** in secondary storage (commonly disks)
 3. **How file operations** like reading, writing, positioning, etc, are done.

File

OS provided abstraction

The OS hides the **physical properties of the storage system** and provides a **uniform vision of storage devices**



- The **file system** provides mechanisms to:
 - **Persistent storage** of information
 - Information stays in the computer when it is switched off, the most common device nowadays is the hard disk
 - **Access to information**
 - A **user interface** made of:
 - **Files**: logical unit for persistent storage of data
 - **Directories**: mechanism (container) to organize files
- **Importance**
 - It keeps system critical data
 - It constrains global system performance
 - It is the most visible and used aspect of an OS

System calls (user library)

It does file and directory management from the programmer sight

File manager

- It uses a device driver to do information transfer between disk and memory
- It implements mechanisms to provide **coherency**, **security** and **protection**
- It **optimizes** performance
- It creates basic user interface elements: **files** and **directories**

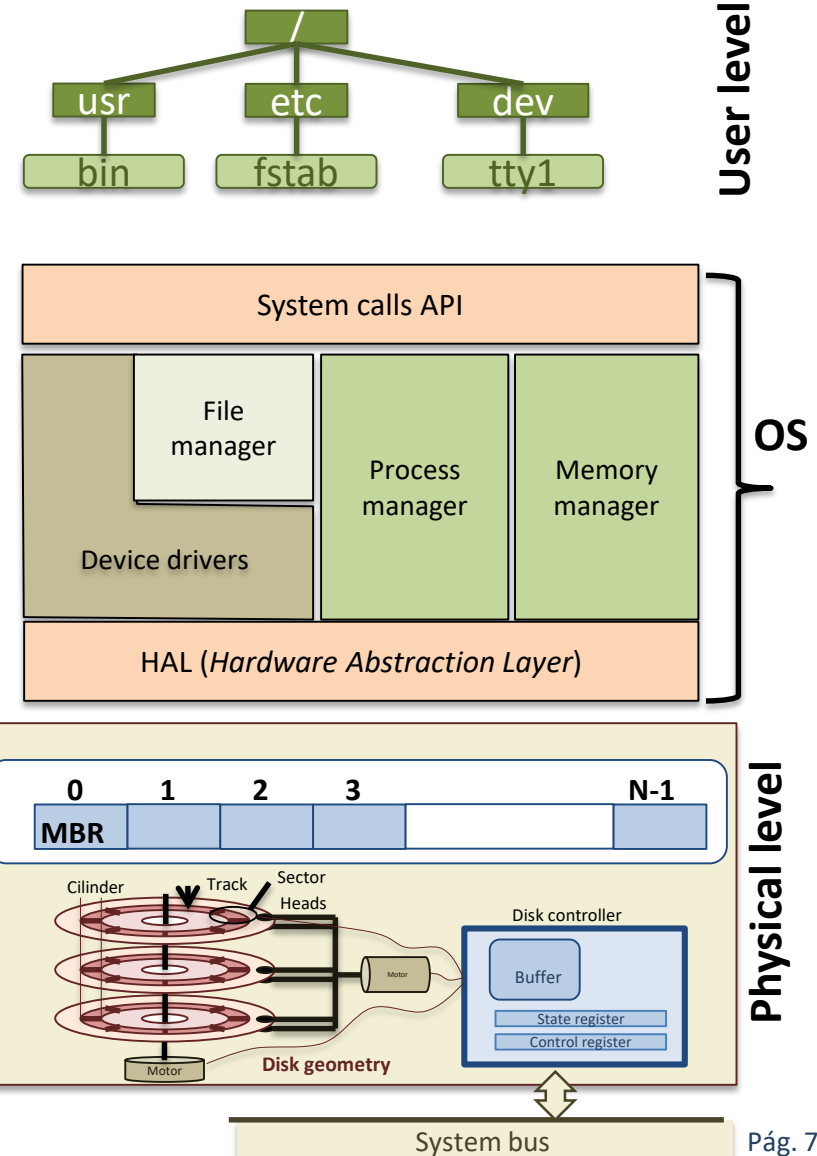
Device driver

- It dialogs with the device controller
- It starts physical operations and processes the end of I/O

Physical level

Device + controller

- Block device
 - Disk geometry
- Disk = Block vector**



- **User sight**

User libraries (to operate with files)

API with system calls related to files and directories

File operations:

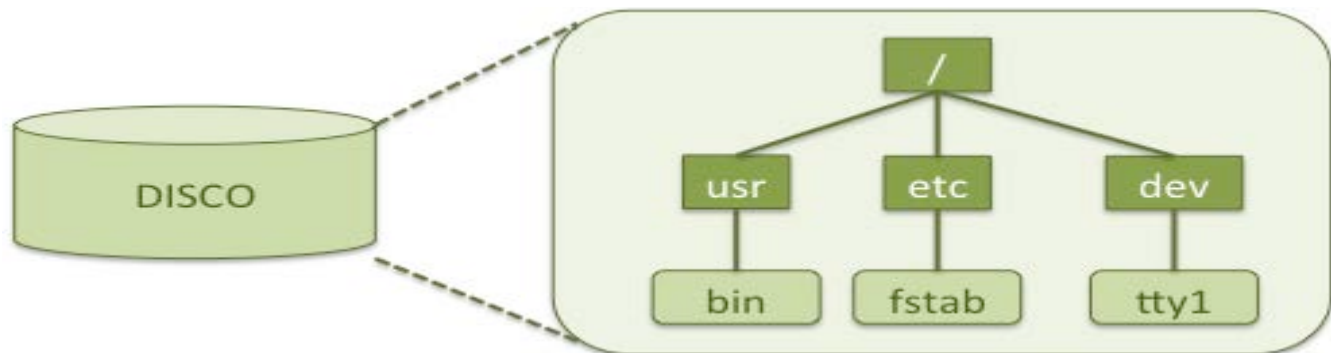
- Open/close
- Read/write
- Seek within a file

Directory operations:

- Create/remove
- Rename
- Searching
- Navigating through the file system

Hierarchical view

Files and directories have a tree organization



User level

File and directory abstractions

- **Open file call**
 - System call that allows accessing file content
 - Processes must open a file before reading or writing on it
 - The calling process receives a file handler that will reference the file in future reading or writing operations
 - It establishes the file access modes (read and/or write) and the initial position on the file of the access pointer
 - For instance, when accessing a file in writing mode
 - Locating the access pointer at the beginning (0) to overwrite the whole file content
 - Locating the access pointer at the end (filesize – 1) to add new content to the file
 - Upcoming accesses will start at the access pointer location left in the last access and will update it
 - File access permissions are checked
 - Open call will fail if the mode specified is inconsistent with the file permissions

- **File operations**

- The OS provides a set of system calls to work with files
 - **Create**
 - It requires free space on disk and to create a new directory entry
 - **Open**
 - Required operation before reading or writing
 - **Read**
 - It requires a file identifier and a reading location pointer
 - **Write**
 - It requires a file identifier, data to write and a writing location pointer
 - **Seek**
 - It sets file reading and writing pointers
 - **Close**
 - It frees OS internal structures that support file access
 - **Remove**
 - It frees disk space allocated to a file and removes the corresponding directory entry

- File system architecture
- **File concept**
- File block allocation

- **A file is:**

- An abstract data type
- An interrelated information collection established by its author
- The required element to write information in secondary storage

```
#include <stdio.h>

main() {
    int x; /*variable entera*/
    int y; /*variable entera*/
    int *px; /* puntero a entero*/
    x=5;
    px=&x; /*px=direccion de x*/
    y=*px; /* y contiene el dato apuntado por px*/

    printf("x=%d\n",x);
    printf("y=%d\n",y);
    printf("*px=%d\n",*px);
    printf("px = %p\n", px);
}
```

Content of a C course code file

- **File = Attributes + Data**

METADATA Attributes

required to
manage the file
system

DATA

File content,
like for instance
text, binary
code, etc

- **File attributes**

- The change from system to system
 - Type
 - Size
 - Protection info
 - Owner
 - Creation date and time

- **File data**

- The OS sees the file content as a byte vector, it is up to the application to give meaning to them
- A file can store different information types: program source, text data, binary code, graphics, sound, etc
- The file data can have a certain structure set by the file type (i.e. wav files, jpeg files, etc)
- An executable file is a file made up by a sequence of code sections that the OS is able to load and execute

File list showing file
attributes

```
gandreu@shell-labs:~/sisop/FS0/ejemplosC$ ls -lhl
total 138K
11928522 -rw-r--r-- 1 gandreu disca-upvnet 470 sep 20 2013 aritmetica_punteros.c
11930469 -rw-r--r-- 1 gandreu disca-upvnet 453 sep 26 2011 aritmetica_punteros.c-
11930470 -rwxr-xr-x 1 gandreu disca-upvnet 0,0K sep 26 2011 cir
11928236 -rw-r--r-- 1 gandreu disca-upvnet 193 sep 22 2011 circulo.c
11930472 -rwxr-xr-x 1 gandreu disca-upvnet 0,9K sep 26 2011 cua
11928246 -rwxr-xr-x 1 gandreu disca-upvnet 8,3K sep 16 16:41 cuad
11928433 -rwxr-xr-x 1 gandreu disca-upvnet 8,9K sep 20 2013 cuadrado
11928435 -rw-r--r-- 1 gandreu disca-upvnet 214 sep 22 2011 cuadrado.c
11928418 -rw-r--r-- 1 gandreu disca-upvnet 193 sep 22 2011 cuadrado.c-
11928437 -rwxr-xr-x 1 gandreu disca-upvnet 0,9K sep 22 2011 cuadro
11933192 -rw-r--r-- 1 gandreu disca-upvnet 80 sep 18 2013 error
11930471 -rw-r--r-- 1 gandreu disca-upvnet 579 sep 26 2011 hipotenusa.c
11930468 -rw-r--r-- 1 gandreu disca-upvnet 453 sep 26 2011 hipotenusa.c-
11927803 -rwxr-xr-x 1 gandreu disca-upvnet 424 jun 27 2014 hola.c
11928409 -rwxr-xr-x 1 gandreu disca-upvnet 241 jun 20 2014 hola.c-
11930473 -rwxr-xr-x 1 gandreu disca-upvnet 8,8K sep 26 2011 punt
11928436 -rwxr-xr-x 1 gandreu disca-upvnet 0,0K sep 22 2011 puntero
11928438 -rw-r--r-- 1 gandreu disca-upvnet 315 sep 22 2011 punteros.c
11928434 -rw-r--r-- 1 gandreu disca-upvnet 214 sep 22 2011 punteros.c-
11928243 -rw-r--r-- 1 gandreu disca-upvnet 525 sep 22 2011 variables.c
```

```
#include <stdio.h>

main() {
    int x; /*variable entera*/
    int y; /*variable entera*/
    int *px; /* puntero a entero*/
    x=5;
    px=&x; /*px=direccion de x*/
    y=*px; /* y contiene el dato apuntado por px*/


    printf("x=%d\n",x);
    printf("y=%d\n",y);
    printf("*px=%d\n",*px);
    printf("px = %p\n", px);
}
```

File content

- **Access methods to file data**
 - There are three access modes to file information:
 - **Sequential**
 - Information is accessed (reading or writing) in order
 - In every read/write operation the location pointer is implicitly updated
 - **Direct**
 - The file is made up of logical registers
 - In every operation an argument indicates the working register
 - **Memory mapping**
 - The file is allocated in a logical memory range of one or several processes
 - In this way file read/write ops are transformed into main memory read/write ops
 - The OS is in charge of updating information into disk

- File system architecture
- File concept
- **File block allocation**

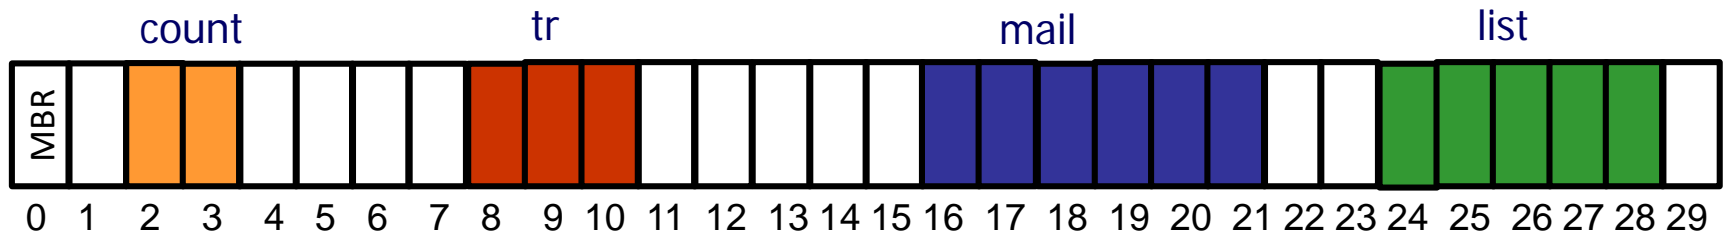
- ¿How to allocate disk space to files?
 - Modern OSs view hard disks as a numbered set of fixed size byte blocks
 - Common sizes are between 512 Bytes y 4 KB (i.e. 1KB)
 - It requires:
 - Efficient use of disk space
 - Fast access to files
 - Allocation schemes:
 - Contiguous
 - Linked
 - Indexed



- **Contiguous allocation**

- A file is allocated as a set of consecutive disk blocks
- It is defined for every file as the first allocated block address and the file length in blocks

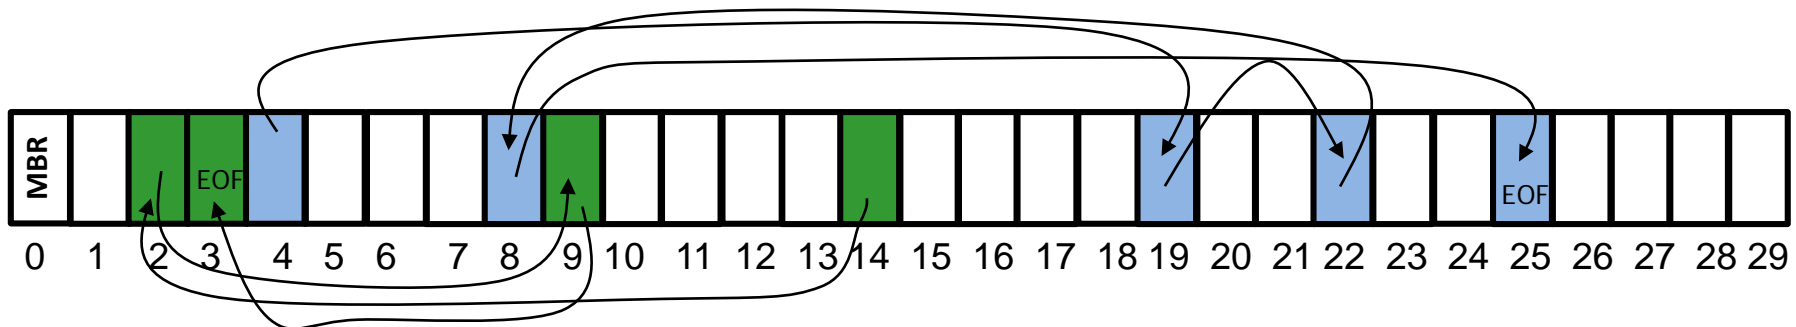
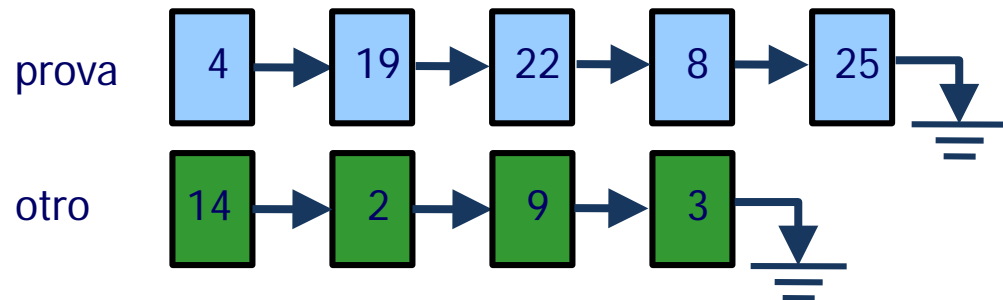
Directory		
File	Start	Length
count	2	2
tr	8	3
mail	16	6
list	24	5



- **Linked allocation**

- File allocated blocks do not need to be contiguous, then every block is linked to the next by means of a pointer

Directory	
File	Start
prova	4
otro	14

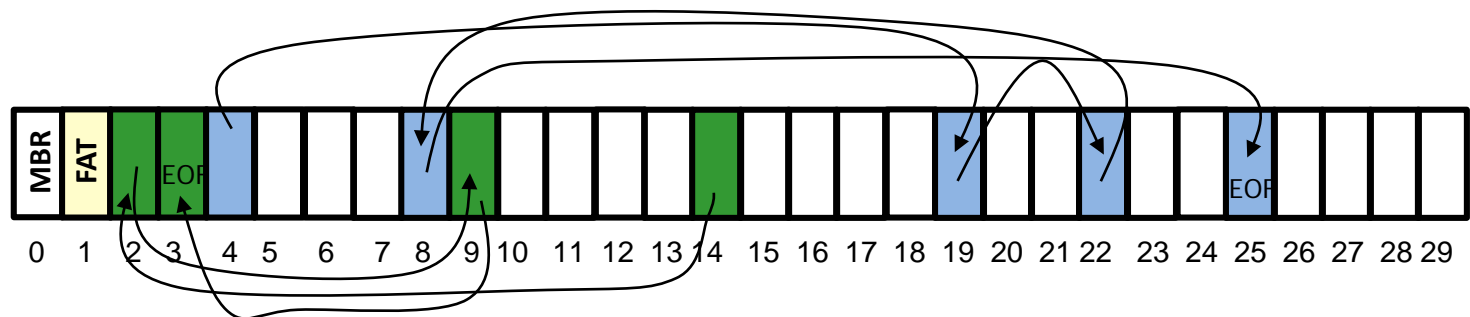
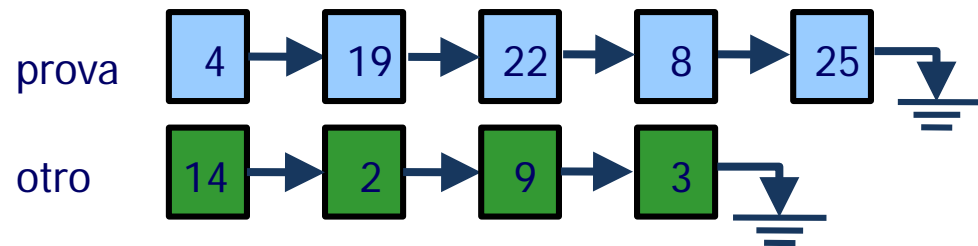


• FAT – variation of linked allocation

- Pointers are not inside the disk blocks but in a disk dedicated area (File Allocation Table)
- EOF marks the list end

0	reservado
1	reservado
2	9
3	Eof
4	19
5	-
6	-
7	4
8	25
9	3
10	-
11	-
12	-
13	-
14	2
15	-
16	-
17	-
18	-
19	22
20	-
21	-
22	8
23	-
24	-
25	Eof
26	-
27	-
28	-
29	-



Directory	
File	Start
prova	4
otro	14



- **Indexed allocation**

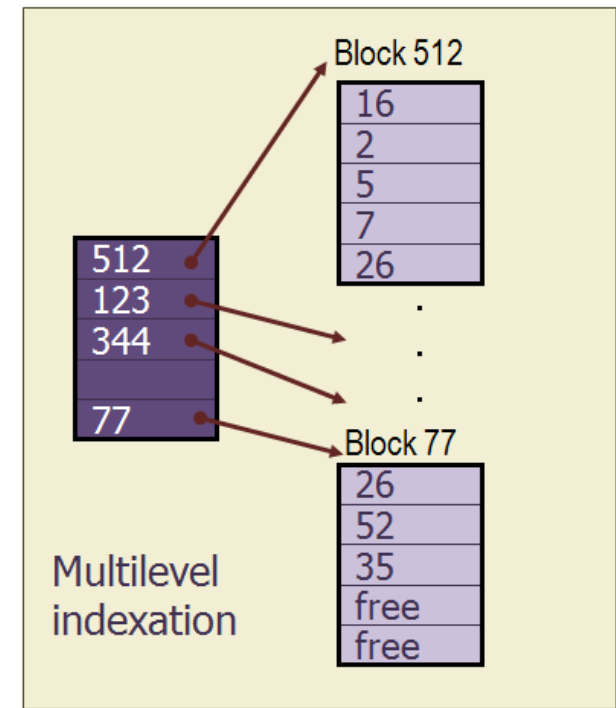
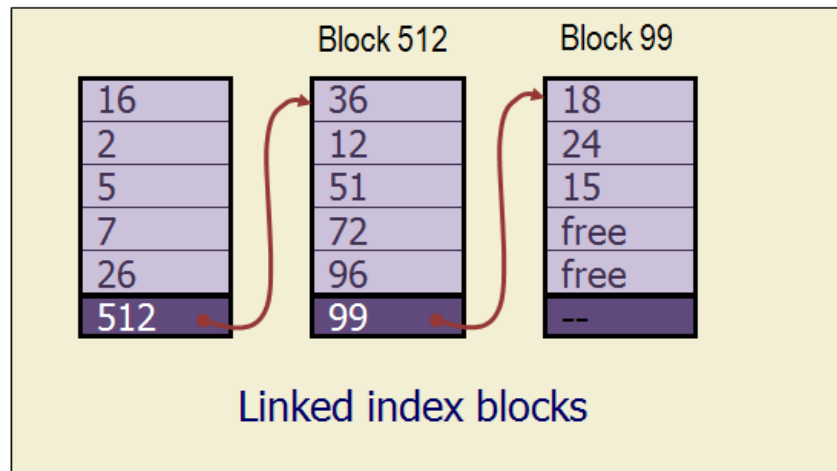
- A block could be index block or data block, a index block contains pointers to data blocks

Directory	
File	Start
prova	5
otro	10



- **Multilevel indexed allocation**

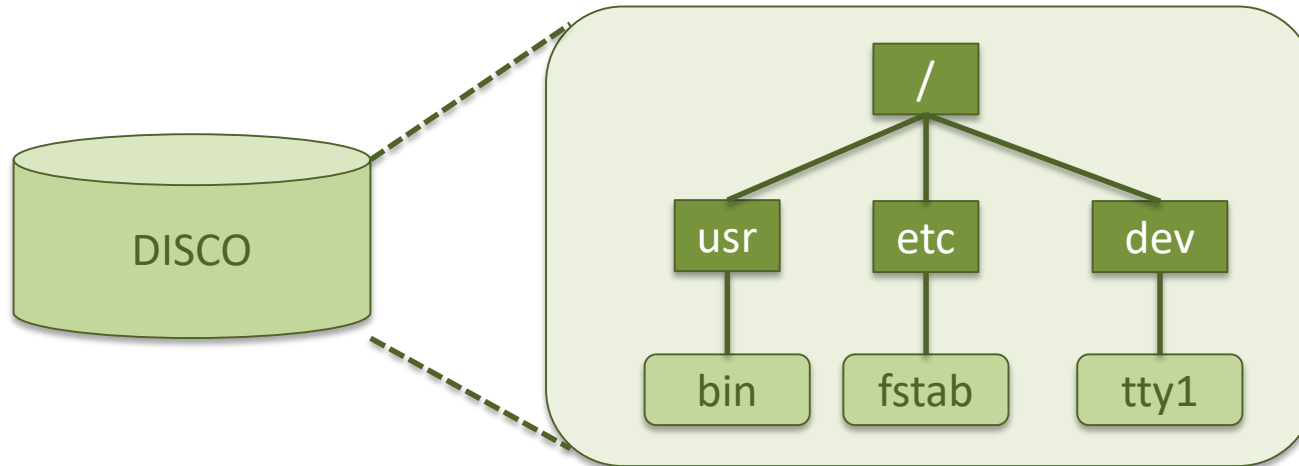
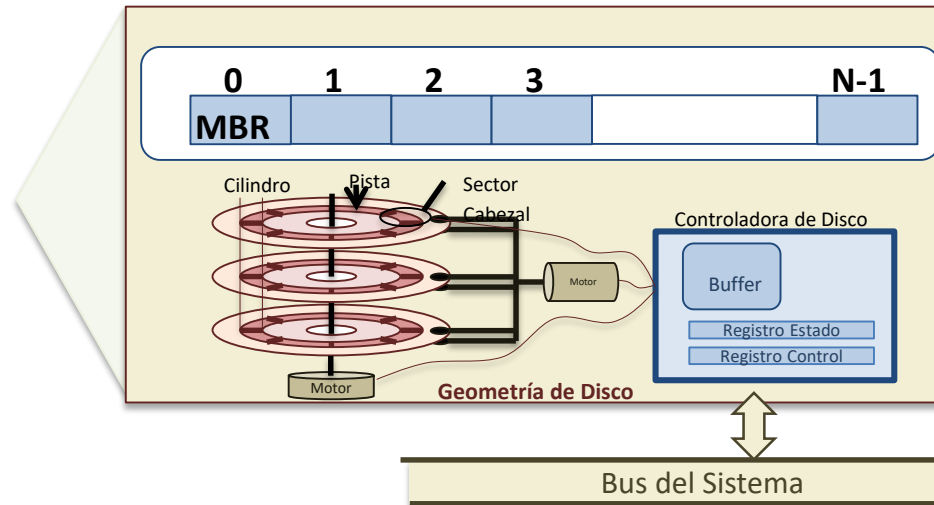
- It is a variation of indexed allocation
- Motivation
 - Supporting big files requires several index blocks
- Solution
 - A pointer can point to a data block or to another index block



- Allocation types analysis

	Advantages	Disadvantages
Contiguous	<p>It is the more efficient</p> <p>It supports sequential and direct access</p> <p>Stable access speed</p> <p>Perfect for read only devices (CD, DVD, etc)</p>	<p>Complex space management (i.e. finding the best gap, relocation due to file growth, etc)</p> <p>It suffers from external fragmentation (computation required from time to time)</p>
Linked	<p>It doesn't constrain file growing</p>	<p>It doesn't support direct access</p> <p>It is little robust against failures</p>
FAT	<p>If FAT is copied in memory then it supports direct access</p> <p>It makes easy free space management</p>	<p>It FAT doesn't fit in main memory then it lacks from any advantage -> only useful for low capacity devices</p> <p>It is little robust against failures</p>
Indexed	<p>It supports sequential and direct access</p>	<p>It constrains file growing (index block size)</p>
Multilevel Indexed	<p>It doesn't constrain file growing</p>	<p>To locate a block several disk accesses may be required</p>

NOTE. In every case there is **internal fragmentation**, half of last block is wasted in average



- Disco

