

Unit 3.- Synchronization Primitives

Concurrency and Distributed Systems



Teaching Unit Objectives

- ▶ Explain the concurrent programming model provided by monitors.
- ▶ Solve the problem of conditional synchronization by means of using monitors
- ▶ Build monitors in concurrent programming languages, avoiding their potential problems.
- ▶ Evaluate existing monitor variants.



Content

- ▶ Concurrent Programming Languages
- ▶ Monitor: Concept
- ▶ Monitor Variants
- ▶ Nested calls



Concurrent Programming Languages

- ▶ Concurrent Programming must solve the needs of communication and synchronization between threads
- ▶ We can design solution strategies at different levels:

Without support of the Operating System

- Busy-waiting
- Disabling interruptions (in supervisor mode)
- TestAndSet, Swap primitives

Using Concurrent Programming Languages

- Monitors

With support of the Operating System

- Semaphores
- Other tools



- Pros: Efficient and flexible
- Cons: Wrong usage cannot be detected when programming

Example.- Ants


- ▶ The ants share a territory
 - ▶ Matrix of cells, each with values free/occupied
 - ▶ Maximum one ant per cell
 - ▶ There is a *lock* that protects this territory
- ▶ Each ant is modeled by a thread
- ▶ When an ant moves from (x,y) to (x',y') executes

close lock;

// if occupied[x',y'] the ant must wait until the cell gets free

occupied[x',y']=true; occupied[x,y]=false //updates matrix

open lock;



How can you
implement this
waiting?



Monitor.- Motivation

- ▶ Primitives ***open lock*** and ***close lock*** ensure secure access to shared variables
- ▶ You also need **other primitives** to safely wait **until certain logic condition is met** (synchronization)

- ▶ A wait loop (i.e. an empty loop that repeatedly checks the condition) does not work. Why not?

close lock

```
while (occupied[x',y']) {}
```

//This DOES NOT work

```
occupied[x',y']=true; occupied[x,y]=false
```

open lock



Monitor.- Motivation

- ▶ Most modern Programming Languages:
 - ▶ They are Object-Oriented Programming (OOP) Languages
 - ▶ Programmers can define types of data (classes)
 - ▶ A class allows defining variables (objects)
 - ▶ Distinction between interface/implementation
 - **Interface** (visible part) corresponds to its behavior (set of methods)
 - **Implementation** (hidden part).- set of attributes, methods code
 - ▶ They require concurrency



Monitor.- Motivation

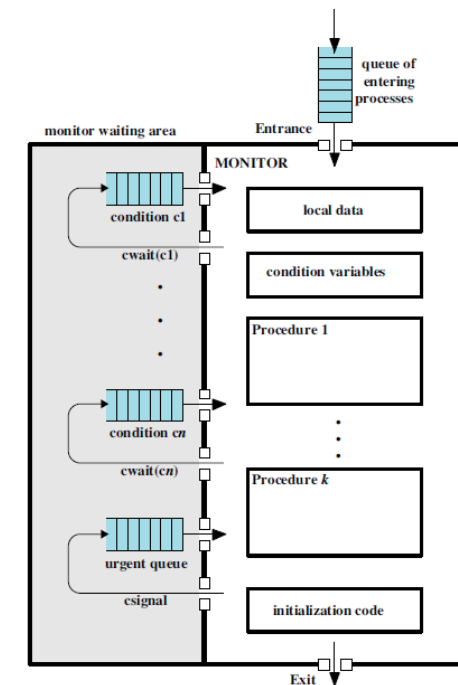
- ▶ Idea.- Mixing OOP and Concurrent Programming
 - ▶ Threads are coordinated by means of **shared objects**.
 - ▶ Details of mutual exclusion and synchronization are hidden inside the classes that represent the shared objects.

- ▶ Advantages:
 - ▶ Simplifies development, maintenance and understanding of code
 - ▶ Facilitates debugging (you can test each piece separately)
 - ▶ Facilitates reusing code
 - ▶ Improves documentation and readability



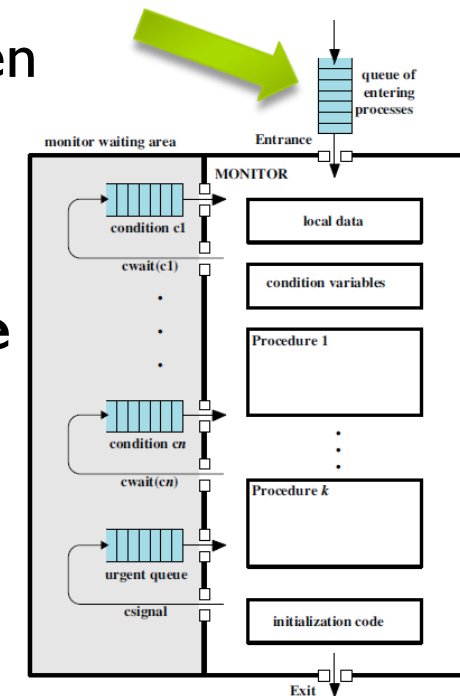
Monitor.- Concept

- ▶ **Monitor** = class to define objects that can be safely shared between different threads
- ▶ Its methods are executed in Mutual Exclusion
- ▶ Solves synchronization



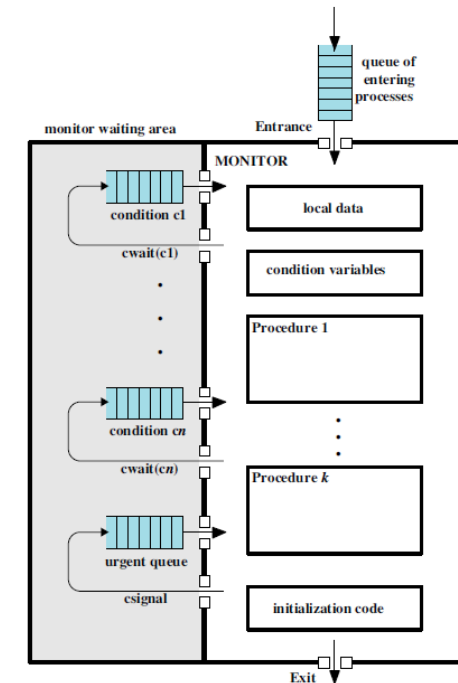
Monitor.- Concept

- ▶ **Monitor** = class to define objects that can be safely shared between different threads
- ▶ Its methods are executed in **Mutual Exclusion**
 - ▶ It has an **Entry queue** for waiting there those threads that want to use the monitor when another thread is using it
- ▶ **There are no race conditions inside the monitor**

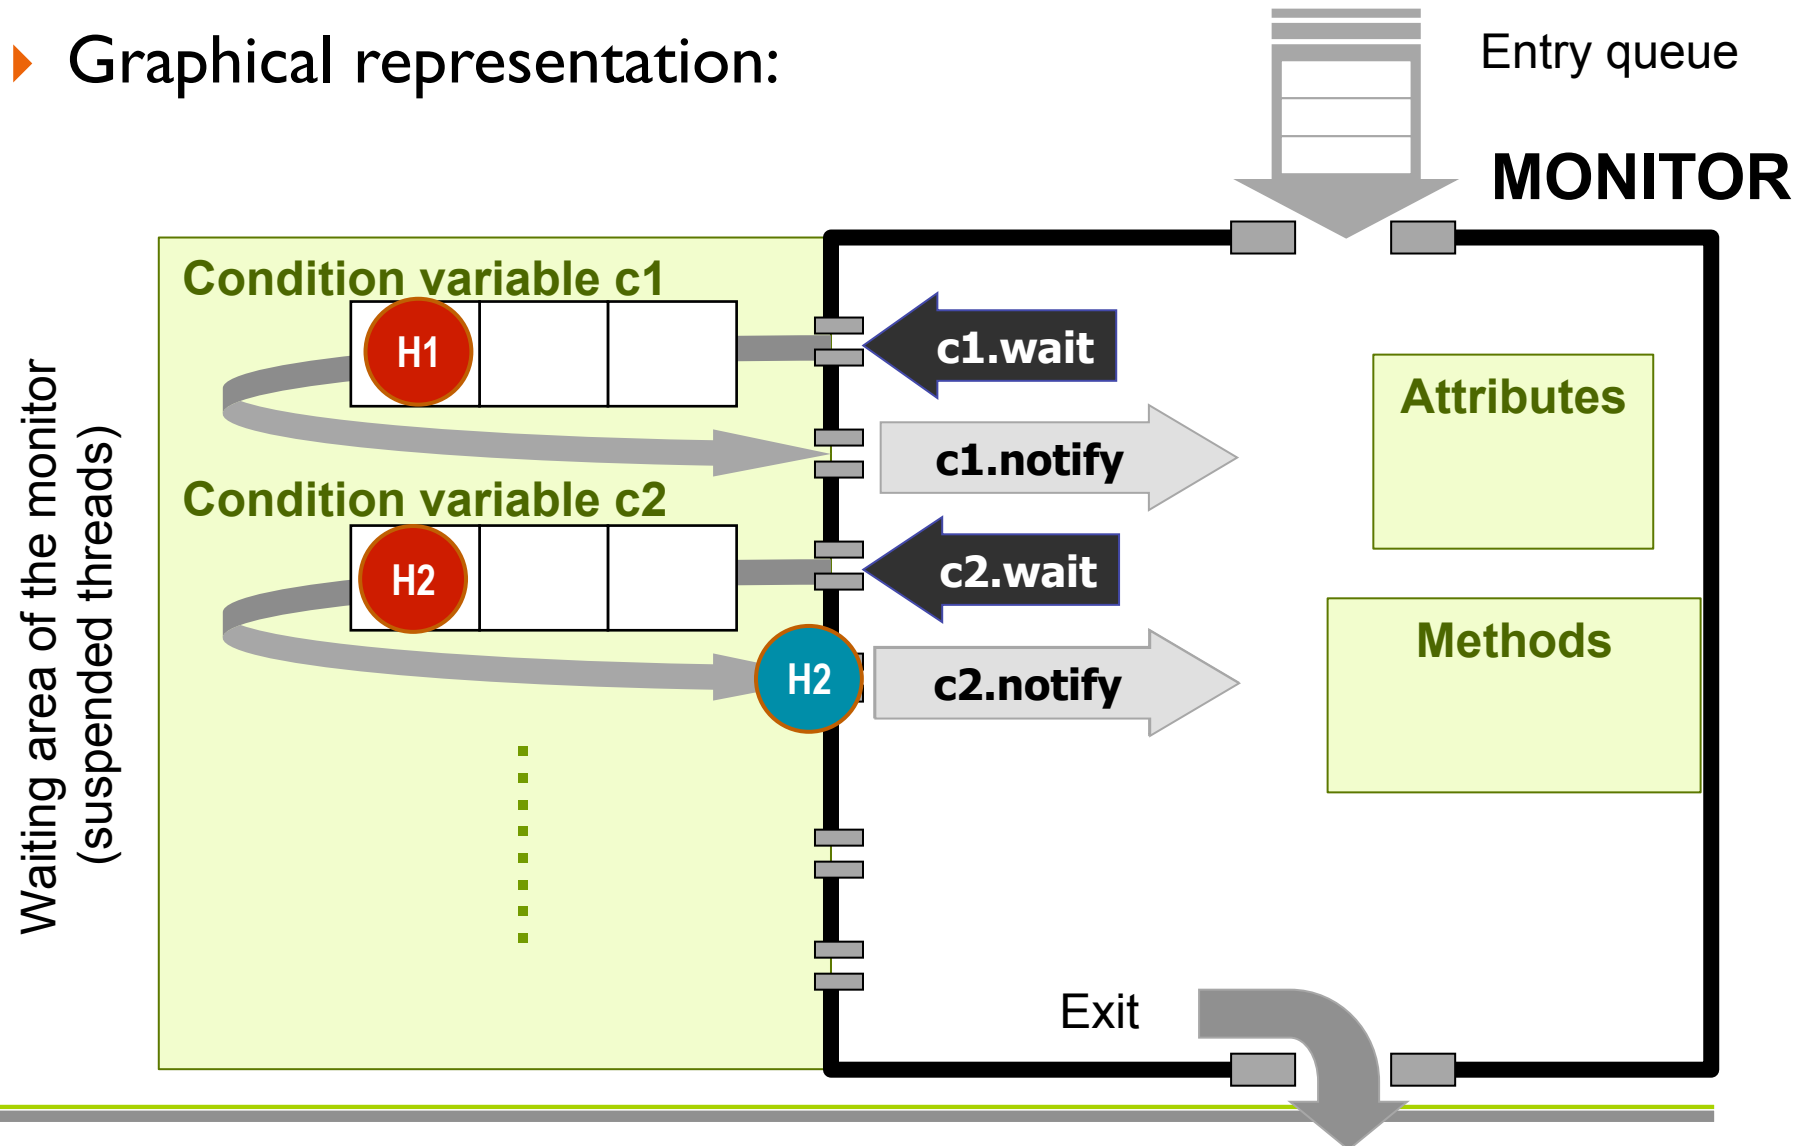


Monitor.- Concept

- ▶ **Monitor** = class to define objects that can be safely shared between different threads
- ▶ **Solves synchronization**
 - ▶ You can define waiting queues (named '**condition variables**') inside the monitor
 - ▶ If a thread executes code inside the monitor it has to wait until a particular logic condition meets
 - Executes **c.wait()** → frees the monitor and waits at the **condition c** queue
 - ▶ When another thread **modifies the state** of the monitor
 - Executes **c.notify()** → reactivates a thread that is waiting in the **condition c** queue



► Graphical representation:





Monitor Example.- Producer/Consumer

- ▶ The shared object is the **buffer**
 - ▶ You design interface and implementation (attributes and method codes) like any other class

- ▶ You need to fill up a table:

Method	Waits when	Notifies to
int get()	Buffer empty	Who waits when buffer full
void put(int e)	Buffer full	Who waits when buffer empty
int numItems()	--	--

- ▶ You define a waiting queue (**condition variable**) for each case of waiting detailed in the table
 - ▶ **condition notFull, notEmpty;**



Monitor Example.- Producer/Consumer

//IMPORTANT.- THIS IS NOT JAVA, but pseudo-code

Monitor Buffer {

.... //attributes for implementing the monitor

condition notFull, notEmpty; //waiting queues

int elems=0;

public Buffer() {..} //initializing the attributes

entry void put(**int** x) { // entry: public method with access in Mutual Exclusion

if (elems==N) {notFull.**wait**();} // waits in the notFull queue

 ... //code for inserting elements in the buffer

 elems++;

 notEmpty.**notify**(); // reactivates someone from notEmpty queue

}

entry int get() {

if (elems==0) {notEmpty.**wait**();} // waits in notEmpty queue

 ... // code to pick up elements from the buffer

 elems--;

 notFull.**notify**(); // reactivates someone from notFull queue

}

entry int numItems() { return elems;}

}

Buffer b; //and from any thread we can invoke b.numItems(), b.get() or b.put(x)



Monitor.- Summary

- ▶ **Monitor** = Class + Mutual Exclusion + Synchronization
- ▶ Hides details of Mutual Exclusion and Synchronization
 - ▶ Performing methods on the same monitor does not overlap (**mutual exclusion**)
 - ▶ For coordination, waiting queues are employed (named **condition variables**)
 - ▶ Primitives: **wait()** and **notify()**
 - When notifying, if there is nobody waiting, then this notification has no effect
 - ▶ These primitives can only be employed inside the monitor
 - ▶ The programmer is responsible of implementing the *wait/notify* calls in the right places.
- ▶ It provides **abstraction**
 - ▶ The programmer that invokes methods on the monitor might ignore how they have been implemented
 - ▶ The programmer that implements the monitor might ignore how its methods will be later employed



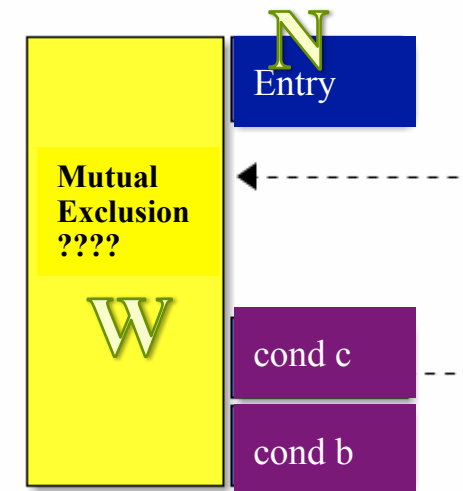
Content

- ▶ Concurrent Programming Languages
- ▶ Monitor: Concept
- ▶ Monitor Variants
- ▶ Nested calls



Monitor.- Variants


- ▶ Monitor ensures **Mutual Exclusion**
 - ▶ Only one thread executes code of the monitor at a given time
 - ▶ If another thread tries to execute code and the monitor is occupied, then this thread waits on the **Entry**.
 - ▶ When the method finishes, then the monitor gets free.
 - ▶ When the active thread (W) inside the monitor executes *c.wait()*, it goes to wait on condition c
 - ▶ **The monitor gets free** (waits outside the CS)
 - ▶ Another thread (N) waiting on the Entry becomes active inside the monitor
- ▶ Problem: if thread N executes *c.notify()*
 - ▶ Reactivates W
 - ▶ But only one thread (W or N) can be active inside the monitor. Which one?





Monitor.- Variants

- ▶ Monitor ensures **Mutual Exclusion**
 - ▶ Only one thread executes code of the monitor at a given time
 - ▶ When the **active thread (W)** inside the monitor executes **c.wait()**, it goes to wait on condition c
 - ▶ **The monitor gets free** (waits outside the CS)
 - ▶ **Another thread (N)** waiting on the Entry becomes active inside the monitor
- ▶ **Problem:** if thread N executes **c.notify()**
 - ▶ Reactivates W
 - ▶ But only one thread (W or N) can be active inside the monitor. Which one?





Monitor.- Variants


- ▶ Let us assume that:
 - ▶ W waits in cond c
 - ▶ N executes c.notify() and reactivates W

- ▶ Solution alternatives (monitor variants)
 - ▶ Thread N leaves the monitor (**Brinch Hansen model**)
 - ▶ Thread N waits in a special queue (**Hoare model**)
 - ▶ Thread W waits in the entry (**Lampson-Redell model**)



Monitor **Brinch Hansen** model


- ▶ Let us assume that:
 - ▶ *W* waits in cond *c*
 - ▶ *N* executes ***c.notify()*** and reactivates *W*
- ▶ The sentence ***notify*** must be the last sentence of the method (this is compulsory in this model)
 - ▶ *N* leaves the monitor and awakes thread *W*
- ▶ Ensures mutual exclusion
 - ▶ *N* leaves monitor
 - ▶ *W* stays active in the monitor
- ▶ It cannot be always applied
 - ▶ Some complex problems might require doing other actions after ***c.notify()***
 - ▶ “*Waterfall awake*” cannot be applied here





Monitor Hoare model


- ▶ Let us assume that:
 - ▶ W waits in cond c
 - ▶ N executes **c.notify()** and reactivates W
- ▶ Apart from entry, there is a **special queue**
 - ▶ To wait until monitor gets free
 - ▶ Priority over entry queue
- ▶ When N executes **c.notify**
 - ▶ N goes to the special queue
 - ▶ W remains active in the monitor
- ▶ Ensures mutual exclusion
 - ▶ N waits outside the monitor
 - ▶ W stays active inside the monitor





Monitor **Lampson-Redell** model

- ▶ Let us assume that:
 - ▶ *W* waits in cond *c*
 - ▶ *N* executes *c.notify()* and reactivates *W*
- ▶ When *N* executes *c.notify*
 - ▶ *W* goes to the entry queue
 - ▶ *N* remains active inside the monitor
- ▶ Ensures mutual exclusion
 - ▶ *W* waits outside the monitor (in the **Entry**)
 - ▶ When it enters again, the state might have changed: the thread must recheck the condition
 - ▶ *N* stays active inside the monitor






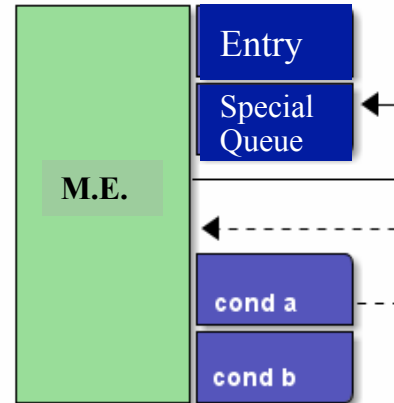
Monitor Variants: Summary

Keeps W (finds condition OK)

Brinch Hansen




Hoare



Keeps N (W will recheck condition when enter again)

Lampson-Redell





Monitor.- Ant example

- ▶ The territory is modeled by a monitor
 - ▶ **Attributes:**
 - ▶ A matrix of logical values, which indicate whether a cell is free/occupied
 - `boolean[N][N] occupied; // This is NOT Java`
 - ▶ **Methods:**
 - ▶ *entry void moves(x,y,x',y')* The ant moves from cell (x,y) to cell (x',y')
- ▶ Two solution alternatives:
 1. A waiting queue for each cell to wait till this cell gets free
 - ▶ `condition[N][N] free; // This is NOT Java`
 2. A single waiting queue named *free*
 - ▶ `condition free; // This is NOT Java`



Monitor.- Ant example

Alternative 1

```
Monitor Territory {  
    boolean[N][N] occupied;  
    condition[N][N] free;  
  
    entry void moves(int x,y,x',y') {  
        if (occupied[x',y'])  
            free[x',y'].wait();  
        //updates matrix  
        occupied[x',y']=true;  
        occupied[x,y]=false;  
        //notifies who wants to go to x,y  
        free[x,y].notify();  
    }  
}
```

Alternative 2

```
Monitor Territory {  
    boolean[N][N] occupied;  
    condition free;  
  
    entry void moves(int x,y,x',y') {  
        while (occupied[x',y']) {  
            free.wait();  
            free.notify();  
        }  
        //updates matrix  
        occupied[x',y']=true;  
        occupied[x,y]=false;  
        //notifies who wants to go to x,y  
        free.notify();  
    }  
}
```



Monitor.- Ant example

- ▶ Alternative 1:
 - ▶ It is much more efficient
 - ▶ Only reactivates an ant if the cell it was waiting for has become free.
 - ▶ *if* clause only possible in some monitor variants (to be discussed later). The other variants require a *while* clause.
- ▶ Alternative 2:
 - ▶ It is less efficient
 - ▶ Employs “**waterfall awake**”. At worst case, all suspended ants can be reactivated when a cell gets free.
 - ▶ It cannot be applied in the Brinch Hansen monitor variant.
 - ▶ It is the only option if you cannot define multiple waiting queues (i.e. multiple condition variables).
- ▶ Whenever possible, you must choose alternative 1.



Java.- Monitor

- ▶ Two possible levels:
 - ▶ Basic support in language
 - ▶ Extended support (by means of *java.util.concurrent* library)

- ▶ In this unit we focus on the basic support
 - ▶ *Java.util.concurrent* is detailed in Unit 5



Java supports the monitor concept

- ▶ Every object has in an **implicit way** (you do not need to declare):
 - ▶ A lock
 - ▶ If we label a method with the word **synchronized**, this ensures execution with mutual exclusion
 - This is equivalent to call *close lock()* before its first instruction and call *open lock()* after its last instruction
 - ▶ A waiting queue with primitives:
 - ▶ **wait()** wait on the *waiting queue*
 - ▶ **notify()** reactivates one of the threads that waits on the *waiting queue*
 - ▶ **notifyAll()** notifies all threads that are waiting
- ▶ But here you **cannot** declare other locks nor other waiting queues.



Java.- How to define a Monitor

- ▶ A class that defines objects to be shared among threads should:
 - ▶ Define all its attributes as **private**
 - ▶ Synchronize all its non-private methods (using *synchronized*)
 - ▶ In the implementation of each method, access only to class attributes and to local variables (defined in this method)
 - ▶ Use *wait()*, *notify()*, *notifyAll()* inside synchronized methods
- ▶ **IMPORTANT:** The compiler does not check anything (it is all the responsibility of the programmer)
 - ▶ There is no warning or error if there are non-private attributes, or public methods without the synchronized label, or we employ *wait()*, *notify()*, *notifyAll()* inside a non-synchronized method.



Java.- How to define a Monitor (cont.)

- ▶ In an ideal monitor, threads that wait for **different** logical conditions have to wait in **different queues**
 - ▶ E.g.- in Producer/Consumer you can have *notEmpty*, *notFull* queues.
 - ▶ Producers that find the buffer full have to wait in the *notFull* queue
 - ▶ Consumers that wait because the buffer is empty stay in the *notEmpty* queue
- ▶ But Java only employs **one condition variable** per monitor
 - ▶ Threads that wait for different logical conditions **have to wait in a unique queue** (thus, in the same queue) and not in different queues
 - ▶ So when reactivating a thread you do not know if this thread was waiting for one condition or another
 - ▶ Except in very simple cases, it is recommended to wake up all threads and each of them has to recheck its condition
 - ▶ *java.util.concurrent* library (see Unit 5) solves this limitation



Java.- How to define a Monitor (cont.)

- ▶ The typical scheme of a method in a shared object in Java is:

```
while (logicalCondition){ // has to wait while the condition is true
    wait();           // try { wait } catch (InterruptedException e) {...}
};

... // normal code

notifyAll(); // if we have modified the state of the object, we have to
notify all threads that are waiting
```



Java.- Ant example

```
public class Territory{
    private boolean[][] occupied;

    public Territory(int N) {
        occupied=new boolean[N][N];
        for (int i=0; i<N; i++)
            for (int j=0; j<N; j++)
                occupied [i][j]=false; //free
    }

    public synchronized void moves(int x0, int y0, int x, int y) {
        while (occupied[x][y])
            try { wait(); } catch (InterruptedException e) {};
        occupied [x0][y0]=false; occupied[x][y]=true;
        notifyAll();
    }
}
```



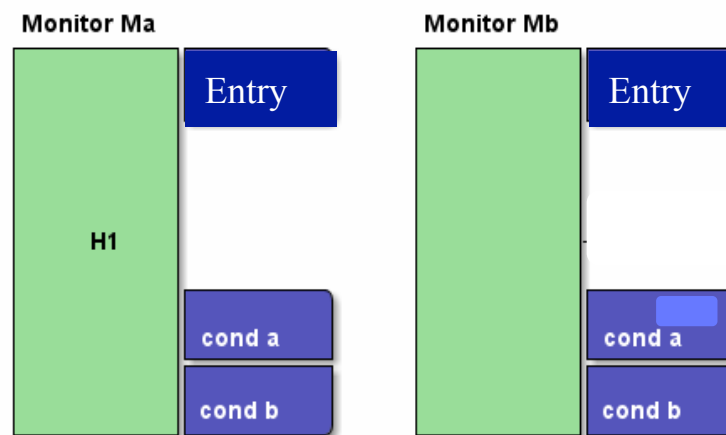

Content

- ▶ Concurrent Programming Languages
- ▶ Monitor: Concept
- ▶ Monitor Variants
- ▶ Nested calls



Monitor.- nested calls


- ▶ Calling from one monitor to a method of another monitor can:
 - ▶ reduce concurrency
 - ▶ and even cause deadlocks.
- ▶ Example: Let us suppose...
 - ▶ 2 monitors Ma and Mb:
 - ▶ from one method of Ma we invoke a method of Mb and vice versa
 - ▶ 2 threads H1 and H2





Monitor.- nested calls

- ▶ H1 active in Ma, invokes a method of Mb, in which it executes **a.wait()**
 - ▶ Goes to the waiting queue “a” of monitor Mb
 - ▶ Releases monitor Mb, but not Ma
 - ▶ Nobody else can use Ma.- we reduce concurrency
- ▶ If H2 enters in Mb (which was free) and invokes a method of monitor Ma
 - ▶ Waits in the entry queue of Ma (since Ma is occupied)
 - ▶ It does not release monitor Mb
 - ▶ We have reached a **deadlock**






Nested calls.- 1st Example of deadlock

```
public class Problem {  
    public synchronized void hello() {...}  
    public synchronized void test (Problem x) { x.hello(); }  
}
```

- ▶ Two threads H1, H2
- ▶ Two variables Problem p1, p2

```
H1 { p1.test(p2) }  
H2 { p2.test(p1) }
```



What happens here?



Nested calls.- 2nd Example of deadlock

- ▶ We define two monitors (p, q) of type Bcell
- ▶ We assume 2 concurrent threads H1 and H2
 - ▶ H1 invokes **p.swap(q)**, obtains access to monitor p and starts execution of *p.swap*
 - ▶ H2 invokes **q.swap(p)**, obtains access to monitor q and starts execution of *q.swap*

What happens here?

```
class BCell {  
    int value;  
    public synchronized void getValue() {  
        return value;  
    }  
    public synchronized void setValue(int i)  
    {  
        value=i;  
    }  
    public synchronized void swap(BCell x)  
    {  
        int temp= getValue();  
        setValue(x.getValue());  
        x.setValue(temp);  
    }  
}
```



Nested calls.- 2nd Example of deadlock

- ▶ There is a deadlock:
 - ▶ Inside *p.swap*, H1 invokes *q.getValue()*, but it must wait because monitor q is not free
 - ▶ Inside *q.swap*, H2 invokes *p.getValue()*, but it must wait because monitor p is not free
 - ▶ Both are waiting to each other, and the situation cannot evolve
→ **DEADLOCK**

```
class BCell {  
    int value;  
    public synchronized void getValue() {  
        return value;  
    }  
    public synchronized void setValue(int i)  
    {  
        value=i;  
    }  
    public synchronized void swap(BCell x)  
    {  
        int temp= getValue();  
        setValue(x.getValue());  
        x.setValue(temp);  
    }  
}
```



Learning results of this Teaching Unit

- ▶ At the end of this unit, the student should be able to:
 - ▶ Program efficient solutions to the conditional synchronization problem, using monitors.
 - ▶ Design a new monitor suitably, attending to the conditions to be managed.
 - ▶ Compare the existing monitor variants.
 - ▶ Classify concurrent programming languages depending on the monitor variant that they support.