



UNIT 5: MAIN MEMORY

Estructura de Computadores (Computer Organization)

Course 2018 / 2019

ETS Ingeniería Informática

Universitat Politècnica de València

Unit goals

- To provide a global view of the computer's memory system and the most relevant features of memories
- To measure memory performance
- To justify the need for a hierarchical memory organization
- To describe the most relevant structural, functional and timing properties of synchronous dynamic RAM (SDRAM)
- To understand how memory modules are built
- To define the concept of memory map and to introduce the design of their decoding circuits
- To understand the role played by dynamic memory controllers

Unit contents

- 1 The memory subsystem
- 2 Basic characteristics of memory
- 3 Main memory
- 4 Dynamic RAM
- 5 Memory modules
- 6 Memory map
- 7 The memory controller

Bibliography and links

- D. Patterson, J. Hennessy. *Computer organization and design. The hardware/software interface*. 5th edition. 2014. Elsevier
- C. Hamacher, Z. Vranesic, S. Zaky. *Computer Organization*. 5th edition. 2001. McGraw-Hill
- Web links
 - Wikipedia: “computer memory”
 - www.micron.com (memory chips)
 - www.kingston.com (memory modules)
 - www.tomshardware.com

1. THE MEMORY SUBSYSTEM

What is memory and where is it located in the computer

Which technologies are more commonly used in computers

The concept of memory hierarchy

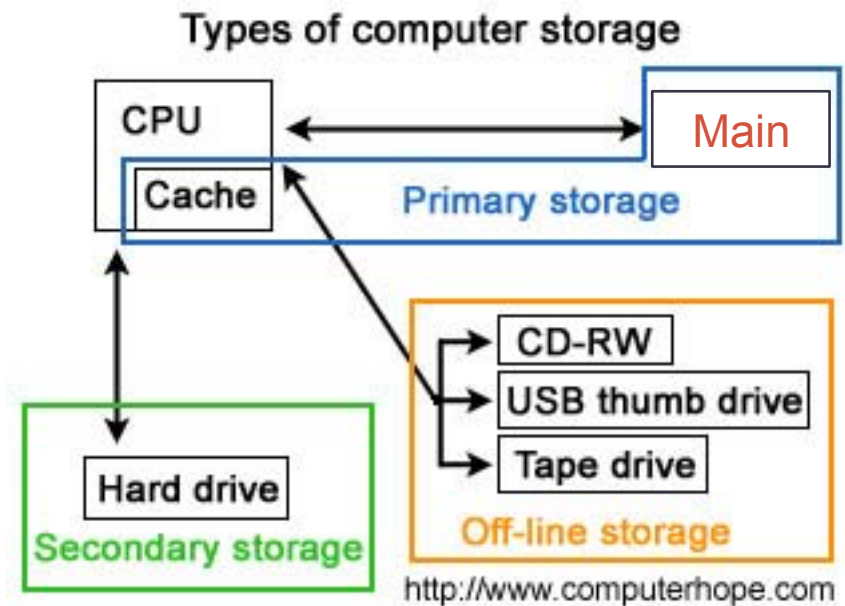
What is memory in a computer

Memory:

Physical support capable of retaining information, namely programs and data, either permanently or temporarily

Where is memory located

- Memory pervades the computer:
 - Primary storage
 - Cache memory
 - Main memory
 - Secondary storage
 - Provides extended capacity and enables multiprogramming through Virtual Memory techniques
 - Off-line storage
 - Backup, portable storage
- Main memory is AKA *RAM*
 - Although this is not precise



Memory technologies

- Four primary technologies are used in today's computers:
 - Main memory uses **synchronous dynamic RAM** (SDRAM)
 - Cache memory is based on **static RAM** (SRAM)
 - Secondary storage is supported by **magnetic disks**
 - **Flash** memory has different uses:
 - As a main memory in portable/embedded devices (replacing SDRAM)
 - As high-performance secondary storage (replacing magnetic disks)
- Off-line storage is supported by a variety of devices using different technologies
 - Optical (CD-ROM, DVD...)
 - Flash (pendrives)
 - Magnetic (external disks, tape)
- Each technology provides different features and performance

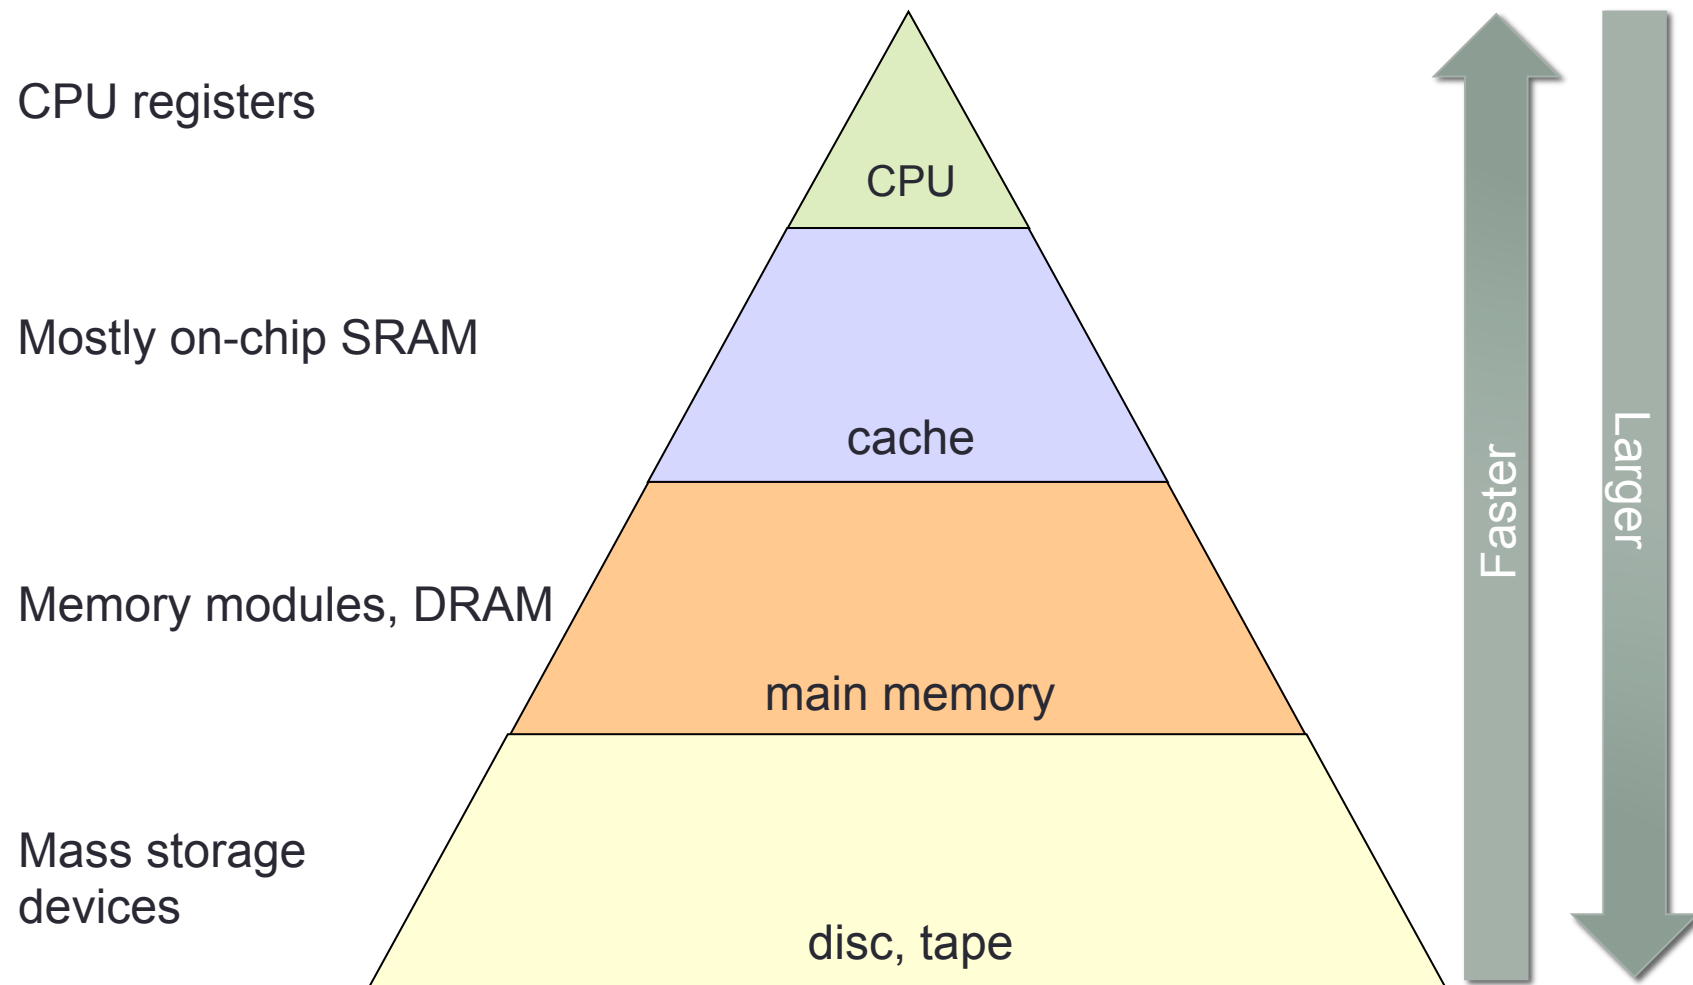
Memory hierarchy

- Ideally, computer memory should be both **large** and **fast**
 - However, the larger the slower; the faster the smaller
- Other features need also be considered
 - Cost per bit, power consumption, reliability...
- No single device is ideal under all requirements

Technology	Access time (typ.)	Cost/GB (yr. 2012)
SRAM	0.5 .. 2.5 ns	\$500 .. \$1000
DRAM	50 .. 70 ns	\$10 .. \$20
Flash	5,000 .. 50,000 ns	\$0.75 .. \$1
Magnetic disk	5,000,000 .. 20,000,000 ns	\$0.05 .. \$0.10

- A **hierarchical memory** organisation provides the illusion of a large, fast, and affordable memory
 - Each hierarchy level uses the most convenient technology

Memory hierarchy



2. BASIC CHARACTERISTICS OF MEMORY

Access modes

Capacity

Access and cycle times

Bandwidth

Access modes

- **Random access:** access time is independent from data location
 - Typical of semiconductor memories. E.g., cache and main memory, ROM memory
- **Sequential access:** access time is proportional to the location of data in the media
 - E.g., magnetic tape
- **Direct access:** access comprises two positioning movements of the reading device, one direct and one sequential
 - Typical of magnetic discs
- **Associative access:** you ask whether some data is in memory
 - AKA contents-addressable memory
 - Access time is independent from data location – but larger than typical *random access memories*, due to the need for the hardware to find the searched data
 - Examples: cache and virtual memory directories (see next unit)

Accessing memory contents

- **Word access**
 - Word size is defined by the width of the memory bus (typical sizes range between 1 and 16 bytes, with 4 bytes in MIPS32)
 - Word (or sub-word) access requires a single bus operation
 - Examples: traffic between cache and CPU, or between main memory and CPU (in systems with no cache)
- **Block access**
 - A block is formed by a set of words
 - Block access may require a number of bus operations
 - Examples:
 - hard disks, optical disks (blocks of 256 B .. 1024 B)
 - transfers between main memory and cache (blocks of 16, 32, 64, 128 B...)
 - Busses for the first cache levels are often block-wide, requiring only one transfer

Memory capacity

- Information (data or code) is typically stored in bytes (8×bit)
- B = Byte; b = bit
 - Each byte is located at a **unique address**
- Most memory chips enable **word** access
 - ...but each byte preserves its address: a word takes as many addresses as bytes it contains
- The **capacity** of a memory device denotes the total number of bytes (sometimes *bits*) it can store

Prefixes	Name	value (2^n)	value (10^n)
	Kilo (K)	2^{10}	10^3
	Mega (M)	2^{20}	10^6
	Giga (G)	2^{30}	10^9
	Tera (T)	2^{40}	10^{12}
	Peta (P)	2^{50}	10^{15}

Memory capacity: examples

- Total capacity, expressed in bytes
 - $1024 \text{ Bytes} = 2^{10} \text{ Bytes} = 1 \text{ KB}$
- Number and size of words
 - $128 \text{ K words of 2 Bytes (16 bits)} = 128 \text{ K} \times 2 \text{ B} = 256 \text{ KB}$
- Total capacity and word size
 - $8 \text{ MB in 32-bit words} = 8 \text{ MB in 4-Byte words} = 2 \text{ M} \times 4 \text{ B}$
 $= 2 \text{ M} \times 32 \text{ b}$
- Other examples
 - $64 \text{ Kb} = 2^{16} \text{ b} = 2^{13} \times 2^3 \text{ b} = 8 \text{ KB}$
 - $256 \text{ Mb} = 2^8 \times 2^{20} \text{ b} = 2^5 \times 2^{20} \times 2^3 \text{ b} = 32 \text{ MB}$
- Example: calculate the number of 16 b words in a memory of 1 Gb


Timing of memory accesses

- **Access time** of a memory: is the **maximum** time elapsed between the start of the (read or write) operation and the arrival or storage of:
 - the data word (in word accesses), or
 - the first word of the block (in block accesses)
- Measured in seconds
 - Prefixes:

Name	Value
milli (m)	10^{-3}
micro (μ)	10^{-6}
nano (n)	10^{-9}
pico (p)	10^{-12}

Timing of word access

- **Cycle time**: **minimum** time required by the memory device between two consecutive accesses
- In general, $t_{\text{cycle}} \geq t_{\text{access}}$
- Example (based on a semiconductor memory)



Transfer speed: Bandwidth

- **Bandwidth**: amount of data transferred per time unit

$$B = \frac{D}{t}$$

- D : amount of data transferred, in bytes
- t : time in seconds it takes the memory to transfer those D bytes
- B : bandwidth in B/s
- B is either measured in B/s or b/s (KB/s, MB/s,... Kb/s, Mb/s...)
 - Caution: when time is involved, prefixes denote powers of 10
- Example: find the bandwidth (in MB/s) of a memory chip with:
 - Word size, $W = 32$ b
 - Cycle time, $t_c = 20$ ns
 - Access time, $t_a = 15$ ns

Transfer speed: Bandwidth

- Another way to calculate bandwidth:


$$B = \text{Bytes_transferred_per_cycle} \times \text{Transfer_frequency}$$

- *Bytes_transferred_per_cycle* depends on data bus width
- *Transfer_frequency* is the inverse of cycle time
- In our previous example:
 - Bytes_transferred_per_cycle = 4 (word size was 32 b = 4 B)
 - Transfer_frequency = 1/20ns = 50 MHz (remember 1 Hz = 1 s⁻¹)
 - Hence, 4 B × 50 MHz = 200 MB/s

Timing of block accesses

- Typical block accesses involve two steps:
 - Access the first word in the block (imposes an initial latency t_L)
 - Access subsequent words in block at regular intervals (t_W)
- Bandwidth of a block transfer is measured starting at the arrival of the first word in the block (B is **asymptotic**)

$$B = \frac{\text{WordSize}}{t_W}$$



3. MAIN MEMORY

External interface

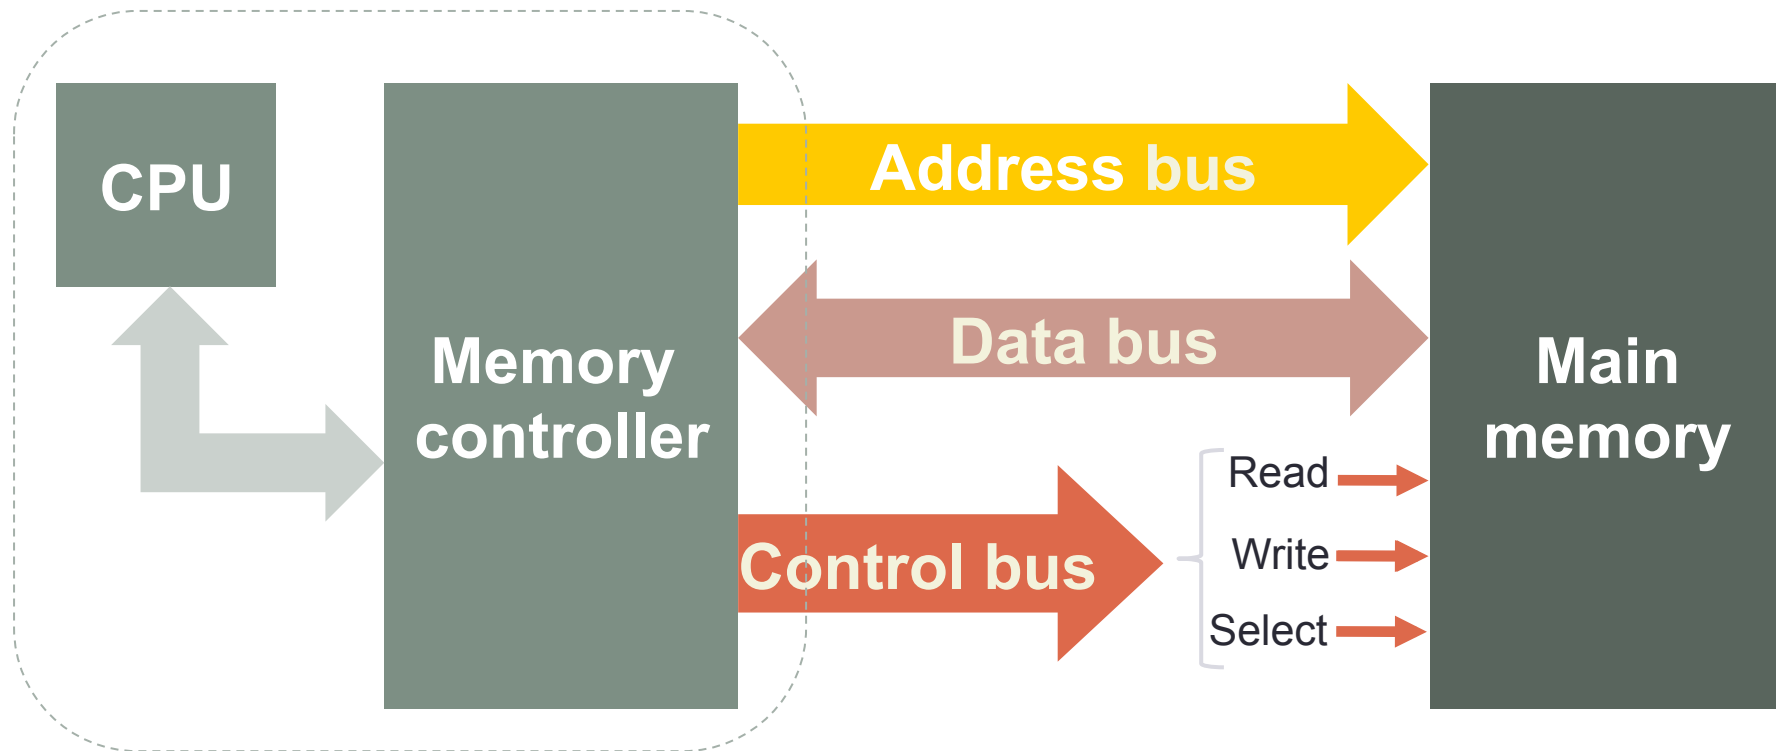
Addressing main memory

Storage arrangement: endianness

Data arrangement in memory (MIPS R2000)

Accessing data

External interface



- Main memory is supported in today's computers by SDRAM chips
- The **memory controller** adapts CPU busses to main memory busses
 - Both in terms of signals and timing

Addressing main memory


- Addresses: each **byte** in memory has a unique address
 - A byte (not a bit!) is the minimum addressable amount of data
- Parameters of Mem-CPU connection
 - **Word width**, in bits or bytes
 - **Addressing capacity**, in number of address bits (n) or in total number of addressable bytes (2^n)
 - Word width \approx data bus size; addressing capacity \approx address bus size
 - **Endianness**: how bytes are arranged in memory
 - Little endian or big endian
- Example: MIPS R2000 addressing space
 - Word width of 32 b = 4 B
 - 32 address bits for a total addressing capacity of 2^{32} B = 4 GB

Storage arrangement: endianness

- With multi-byte words, two arrangements are possible
 - Low-order bytes of the word use lower byte addresses: **little endian**
 - High-order bytes of the word use lower byte addresses: **big endian**
- Example: in the MIPS R2000 ($W = 32$ b)
 - Words are aligned to addresses multiple of 4
 - Half words are aligned to even addresses (multiple of 2)
 - Bytes may be located at any address, with no restriction
 - So when you address a word, where are the most and least significant bytes located?


Data arrangement: endianness

- Little endian example



Data arrangement: endianness

- Big endian example



Data arrangement in assembly

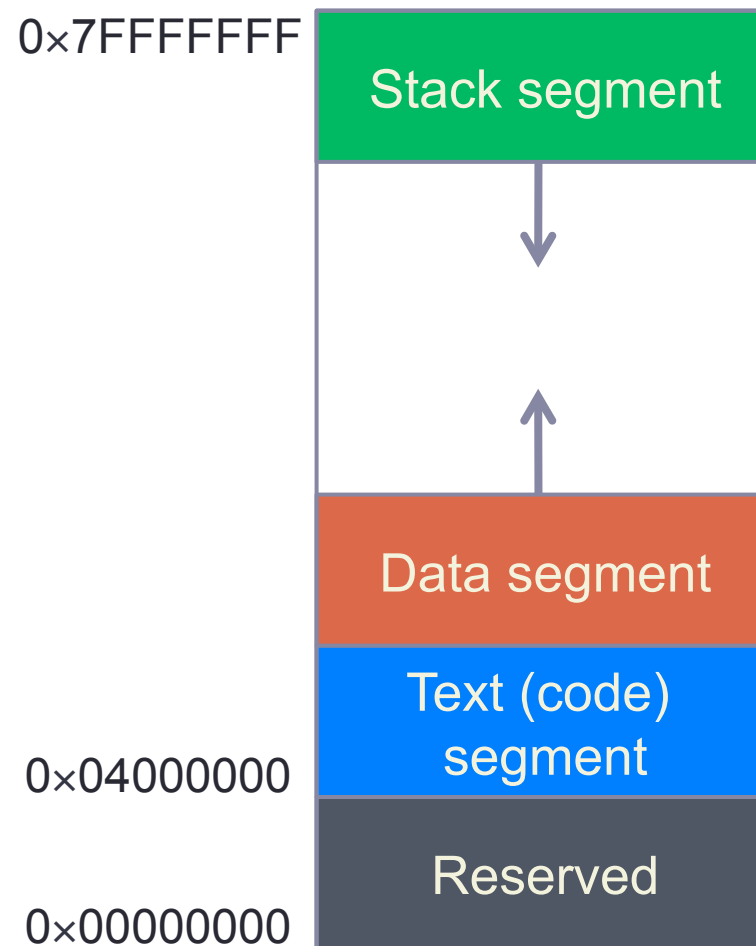
- Assembly directives for data declaration
 - Bytes: **.byte**
 - Half words (2 B): **.half**
 - Words (4 B): **.word**
 - Char strings: **.ascii** and **.asciiz**
- Example

```

.data 0x1000A000
dni:      .word 0x98534169
strz:     .asciiz "K"
codbar:   .half 0x8733
date:     .byte 0x54
name:     .ascii "Pere"
          .word 0x12345678
          .space 2
  
```

0x98	0x53	0x41	0x69	0x1000A000
0x87	0x33	0x00	0x4B	0x1000A004
0x72	0x65	0x50	0x54	0x1000A008
0x00	0x00	0x00	0x65	0x1000A00C
0x12	0x34	0x56	0x78	0x1000A010
		0x00	0x00	0x1000A014

Logical memory arrangement (R2000)




Note that this figure is not drawn to scale

Accessing data in MIPS R2000

- Possible accesses are read/write and byte/half/word
- Read accesses
 - Byte read: `lb $t0, 0($a1)`
 - Half-word read: `lh $t3, 10($a0)`
 - Word read: `lw $t1, 0($a2)`
- Write accesses
 - Byte write: `sb $t2, 27($a3)`
 - Half-word write: `sh $t0, Var($a2)`
 - Word write: `sw $t5, 0x40($a1)`
- The type of access is encoded in the instruction
- Accessing wrong addresses is an error!
 - Cases: half-word at odd address; word at address $\neq 4$
 - Both cases result in an *exception* being raised

Accessing memory: hw details

- The programmer uses **logical addresses** (32 b in R2000)
 - calculated as the sum of the address register and displacement (AKA offset), eg.
 - `lw $t0, 28($zero)` → The address is 28, or 0x1C
- Accessing a word implies accessing its four bytes
 - The word's address is the address of its LSB (in little endian)
- Example ($W = 32$ b): Logical address



Accessing memory: hw details

Addressing in 8-bit CPUs

A₃ A₂ A₁ A₀ 8-bit data

0	0	0	0		0
0	0	0	1		1
0	0	1	0		2
0	0	1	1		3
0	1	0	0		4
0	1	0	1		5
0	1	1	0		6
0	1	1	1		7
1	0	0	0		8
⋮					

← `lb $t0, 2($0)`

← `sb $t0, 5($0)`

Accessing memory: hw details

Addressing in 16-bit CPUs

$A_3 A_2 A_1$			16-bit data	
			$A_0=1$	$A_0=0$
0	0	0	1	0
0	0	1	3	2
0	1	0	5	4
0	1	1	7	6
1	0	0	9	8
1	0	1	11	10
1	1	0	13	12
1	1	1	15	14

`lb $t0, 2($0)`

`sb $t0, 5($0)`

`sh $t0, 10($0)`

Can't make
 $A_0 = 0$ and $A_0 = 1$
simultaneously

Accessing memory: hw details

Addressing in 16-bit CPUs: Use of **Byte Enable** lines (BE^*_i)

$A_3 A_2 A_1$	16-bit data			BE^*_1	BE^*_0
	$BE^*_1=0$	$BE^*_0=0$			
0 0 0	1	0			
0 0 1	3	2	← lb \$t0,2(\$0)	1	0
0 1 0	5	4	← sb \$t0,5(\$0)	0	1
0 1 1	7	6			
1 0 0	9	8			
1 0 1	11	10	← sh \$t0,10(\$0)	0	0
1 1 0	13	12			
1 1 1	15	14			

BE^*_i depend on the address (odd/even) and the type of access (byte/half)
 A_0 is not needed as an address bus line anymore

Accessing memory: hw details

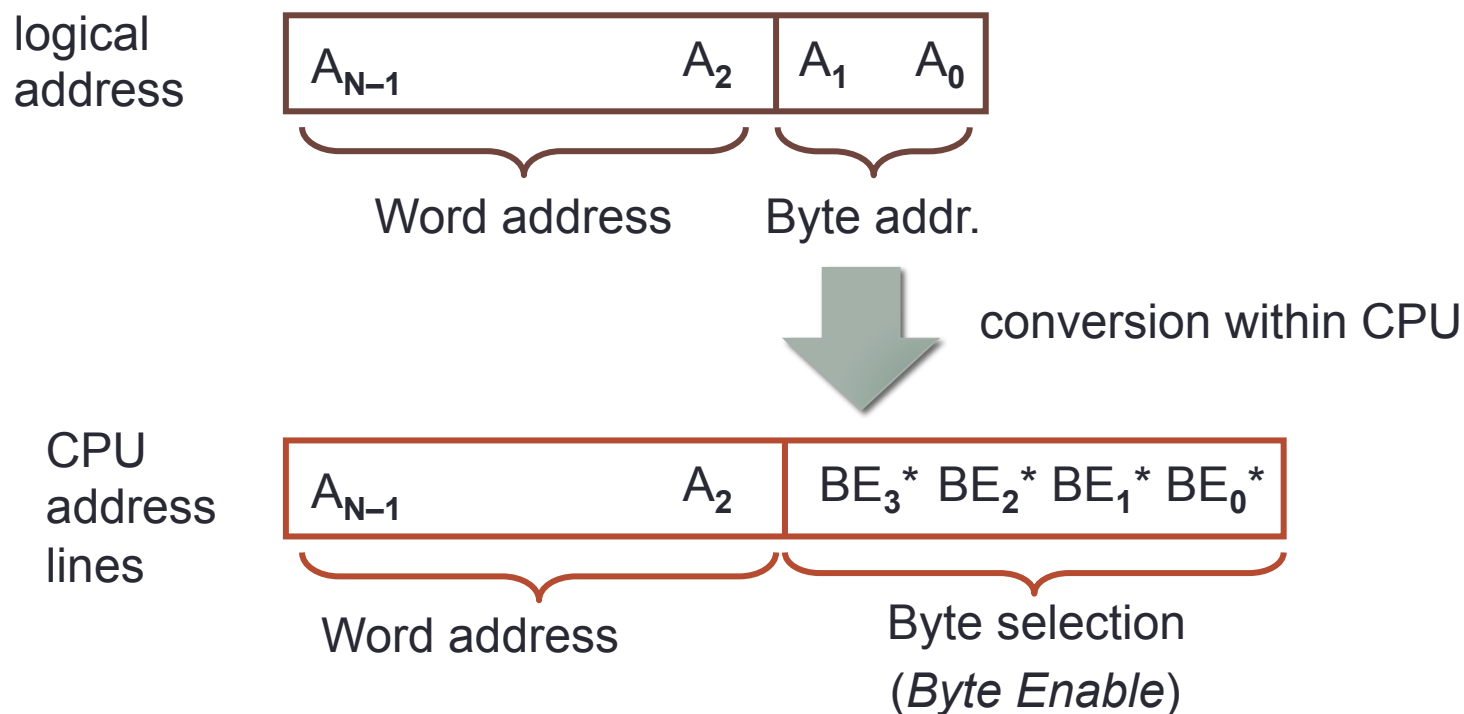
Addressing in 32-bit CPUs: Use of **Byte Enable** lines (BE^*_i)

$A_3 A_2$		32-bit data					$BE^*_3 BE^*_2 BE^*_1 BE^*_0$			
		$BE^*_3=0$	$BE^*_2=0$	$BE^*_1=0$	$BE^*_0=0$					
0	0	3	2	1	0	← lb $\$t0, 2(\$0)$	1	0	1	1
0	1	7	6	5	4	← sh $\$t0, 4(\$0)$	1	1	0	0
1	0	11	10	9	8					
1	1	15	14	13	12	← lw $\$t0, 12(\$0)$	0	0	0	0

BE^*_i depend on the address (2 LSbs) and the type of access (byte/half/word)
 A_0 and A_1 are not needed as address bus lines anymore

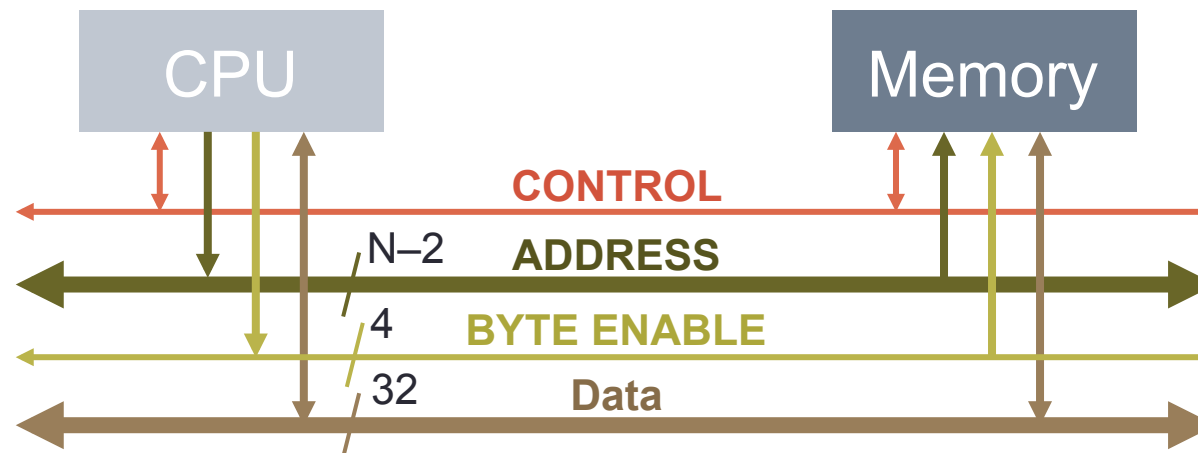
Accessing memory: hw details

- Address lines decode the byte address bits to allow proper access to sub-word items
 - This conversion considers the type of access (byte/half/word)
 - The least significant bits of the logical address are decoded into *byte selection* (or *byte mask*) lines BE_i^*



Accessing memory: hw details

- Memory bus structure, assuming
 - logical addresses of N bits (2^N addressing capacity)
 - $W = 32$ b
 - → there will be
 - $N - 2$ physical address lines, $A_{N-1} \dots A_2$
 - 4 Byte Enable lines, named BE_0^* , BE_1^* , BE_2^* and BE_3^*
 - 32 data lines, $D_{31} \dots D_0$
 - The needed control lines (read, write,...)



Accessing memory: hw details


- Generating BE_i^* signals: example $W = 32$ b
 - The control unit within the CPU is in charge

Inputs			Outputs			
Access type	A_1	A_0	BE_3^*	BE_2^*	BE_1^*	BE_0^*
Byte	0	0	1	1	1	0
	0	1	1	1	0	1
	1	0	1	0	1	1
	1	1	0	1	1	1
Half word	0	0	1	1	0	0
	1	0	0	0	1	1
Word	0	0	0	0	0	0

Accessing memory: hw details

- Examples of BE_i^* activation


```
la $t0, 0x10000000  
lw $a0, 0($t0)  
lb $a1, 5($t0)  
sh $a2, 6($t0)  
sb $a3, 8($t0)
```



	A_1	A_0	BE_3^*	BE_2^*	BE_1^*	BE_0^*
lw	0	0	0	0	0	0
lb	0	1	1	1	0	1
sh	1	0	0	0	1	1
sb	0	0	1	1	1	0

Accessing memory: hw details

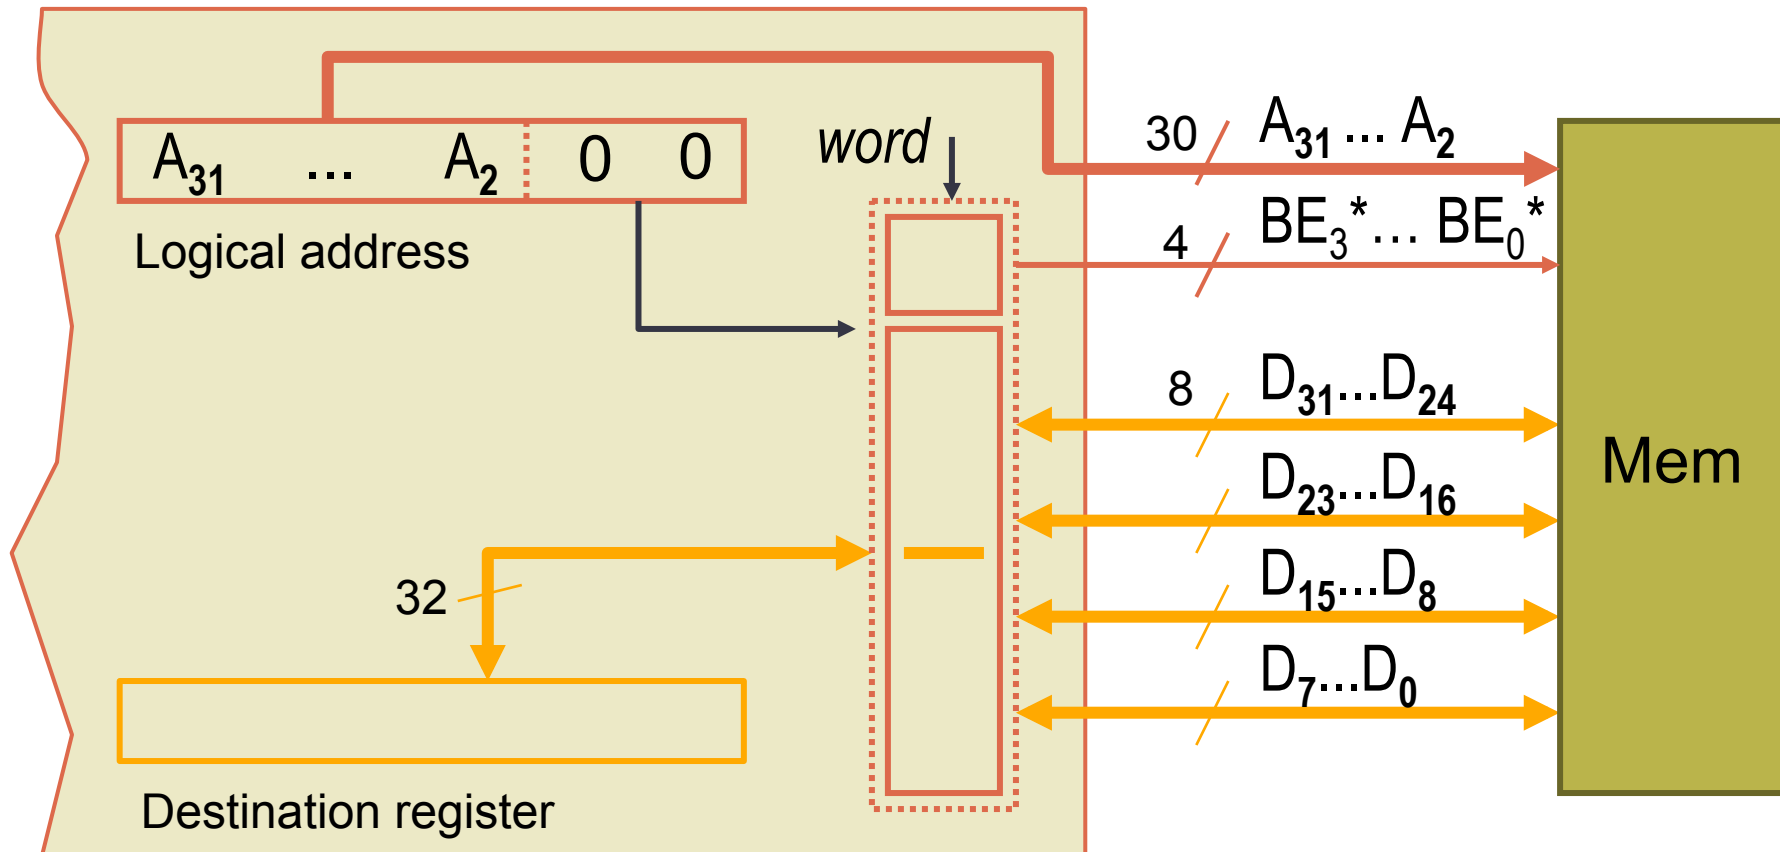
- Endianness dictates the association between BE_i^* lines and data bus bytes



Little endian example

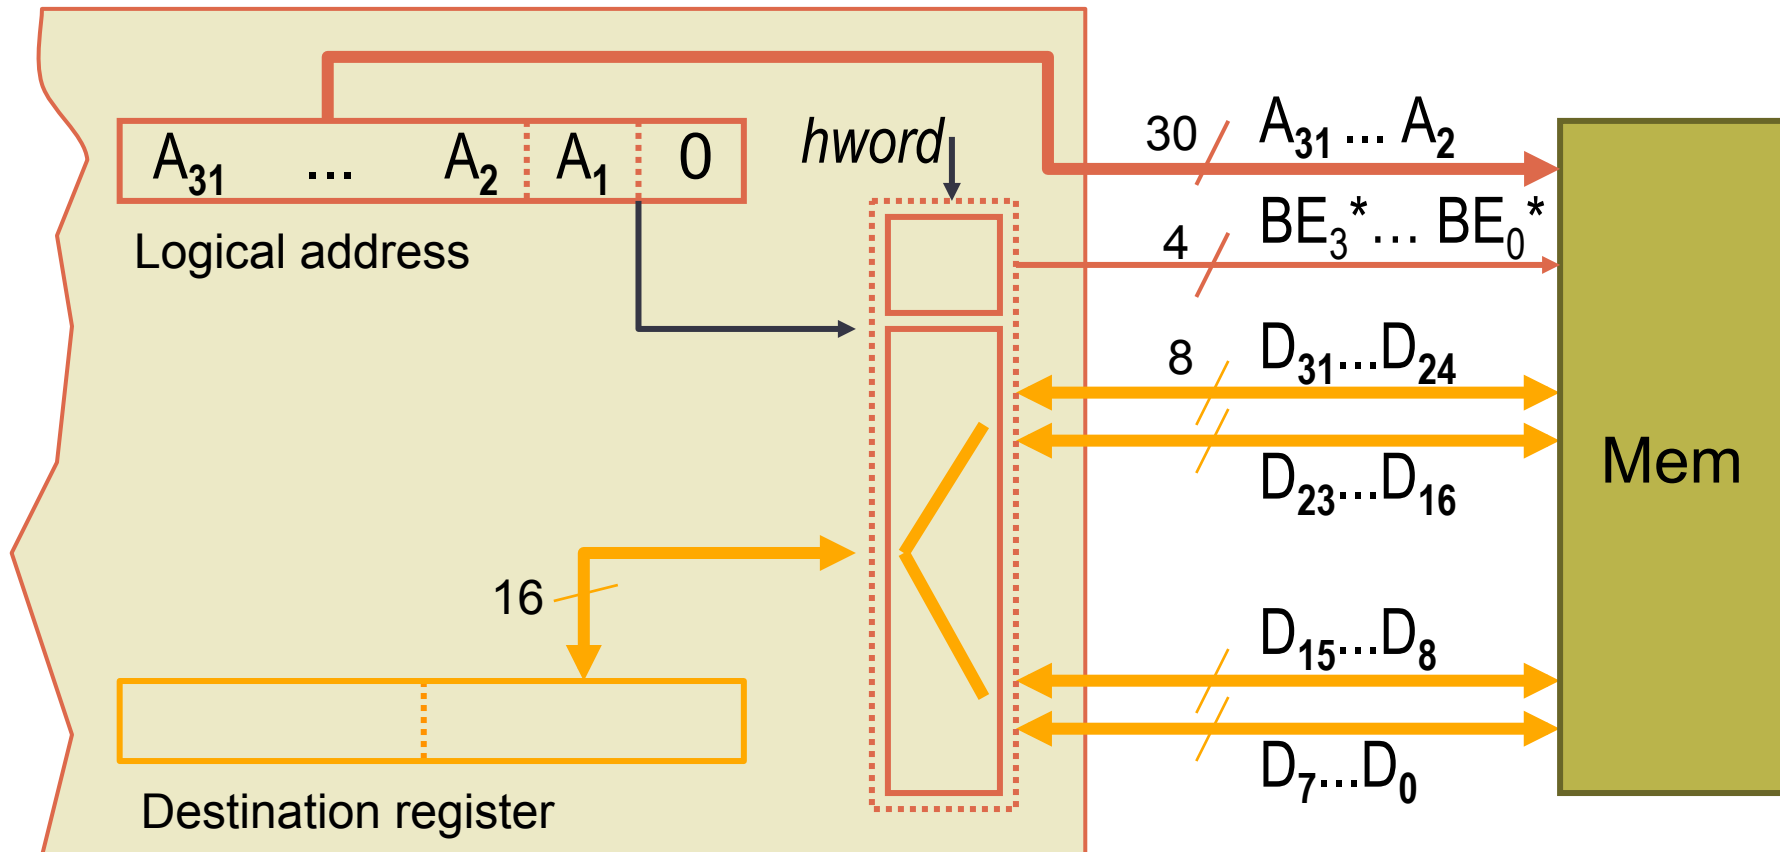
Accessing memory: hw details

- Accessing a word (32 b)
 - Logical address must be multiple of 4 ($A_0=A_1=0$)



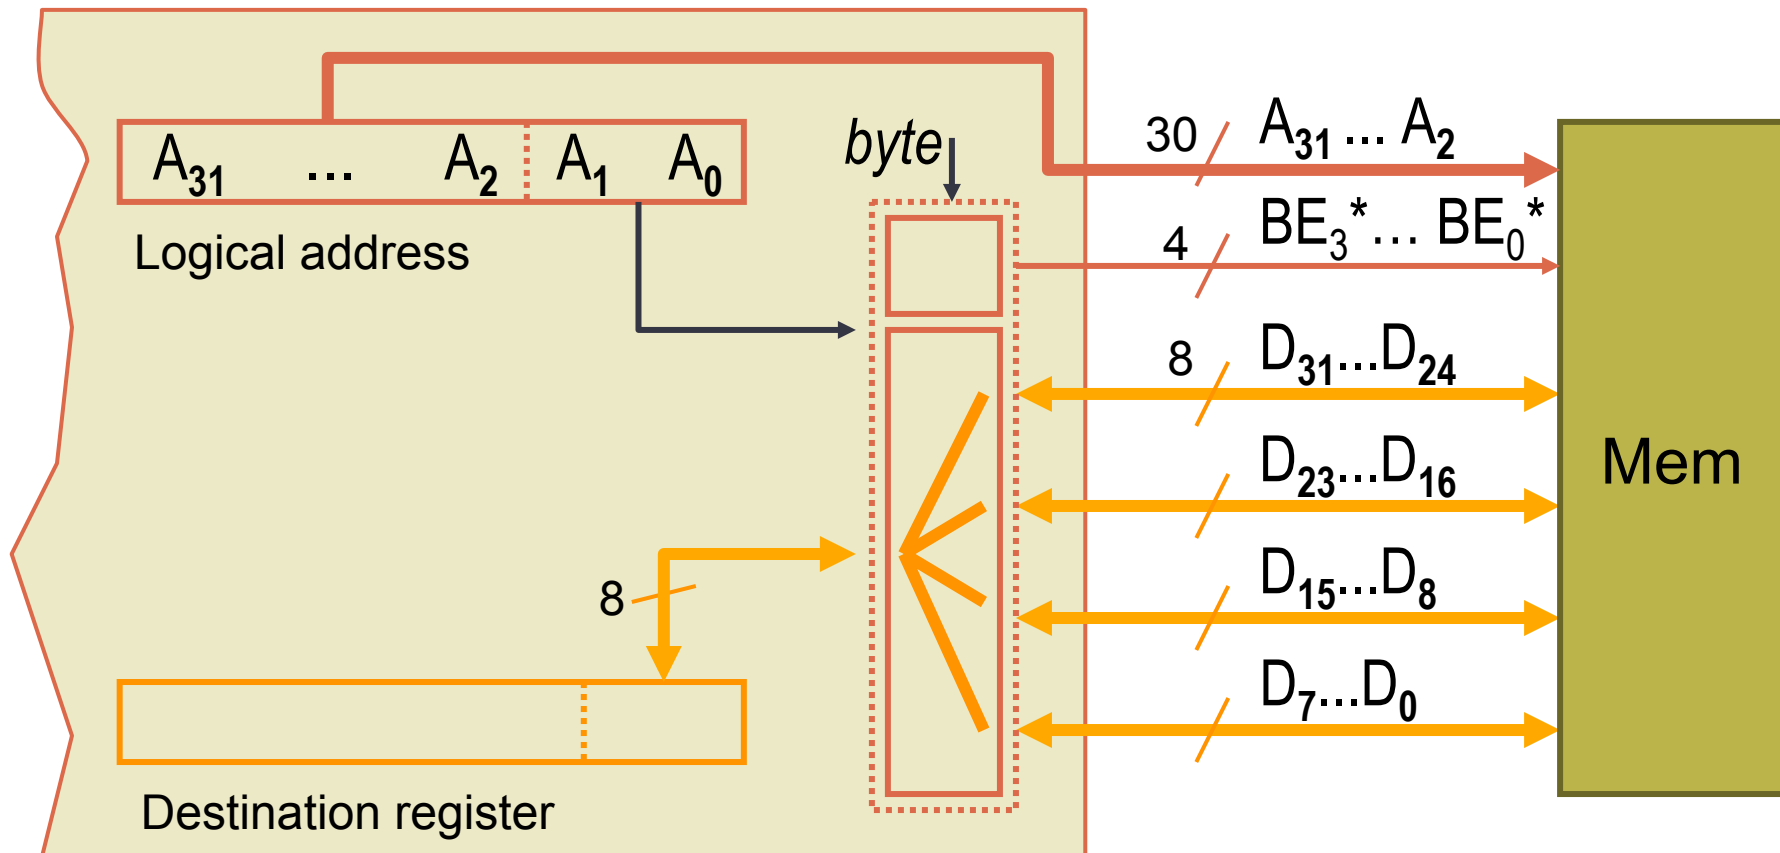
Accessing memory: hw details

- Accessing a half word (16 b)
 - Logical address must be multiple of 2 ($A_0=0$)



Accessing memory: hw details

- Accessing a byte (8 b)
 - Logical address may have an arbitrary value






4. DYNAMIC RAM MEMORY

Description of DRAM / SDRAM memory
Synchronous DRAM (SDRAM) operation
SDRAM commercial example
Commands and timing
DRAM technology: status and trends


DRAM memory basics

- DRAM is a R/W semiconductor memory
- Cells are formed by a *pass transistor* and a capacitor – bits stored as capacitor charge
- Row activation enables
 - Cell write: C is charged through the bit line
 - Cell read: C's charge flows through the bit line
- Charge is very small
 - Bit line is *precharged* to an intermediate voltage and upon row activation, the sense amplifier detects and amplifies slight voltage variations around it
 - Bit line has an associated capacitance
- Charge leak imposes the need for *refresh*
 - Refresh is accomplished by reading the cell and overwriting the same value




DRAM memory basics

- DRAM is preferred to Static RAM for main memory because:
 - DRAM requires less area per bit
 - Static RAM cells use flip-flops (~6 transistors), SDRAM uses only one transistor + capacitor
 - Less area implies higher density (more capacity in the same area)
 - DRAM has much lower cost per bit
 - Although DRAM is slower for individual word access, this is mitigated by *block access*
 - Especially convenient for systems with cache



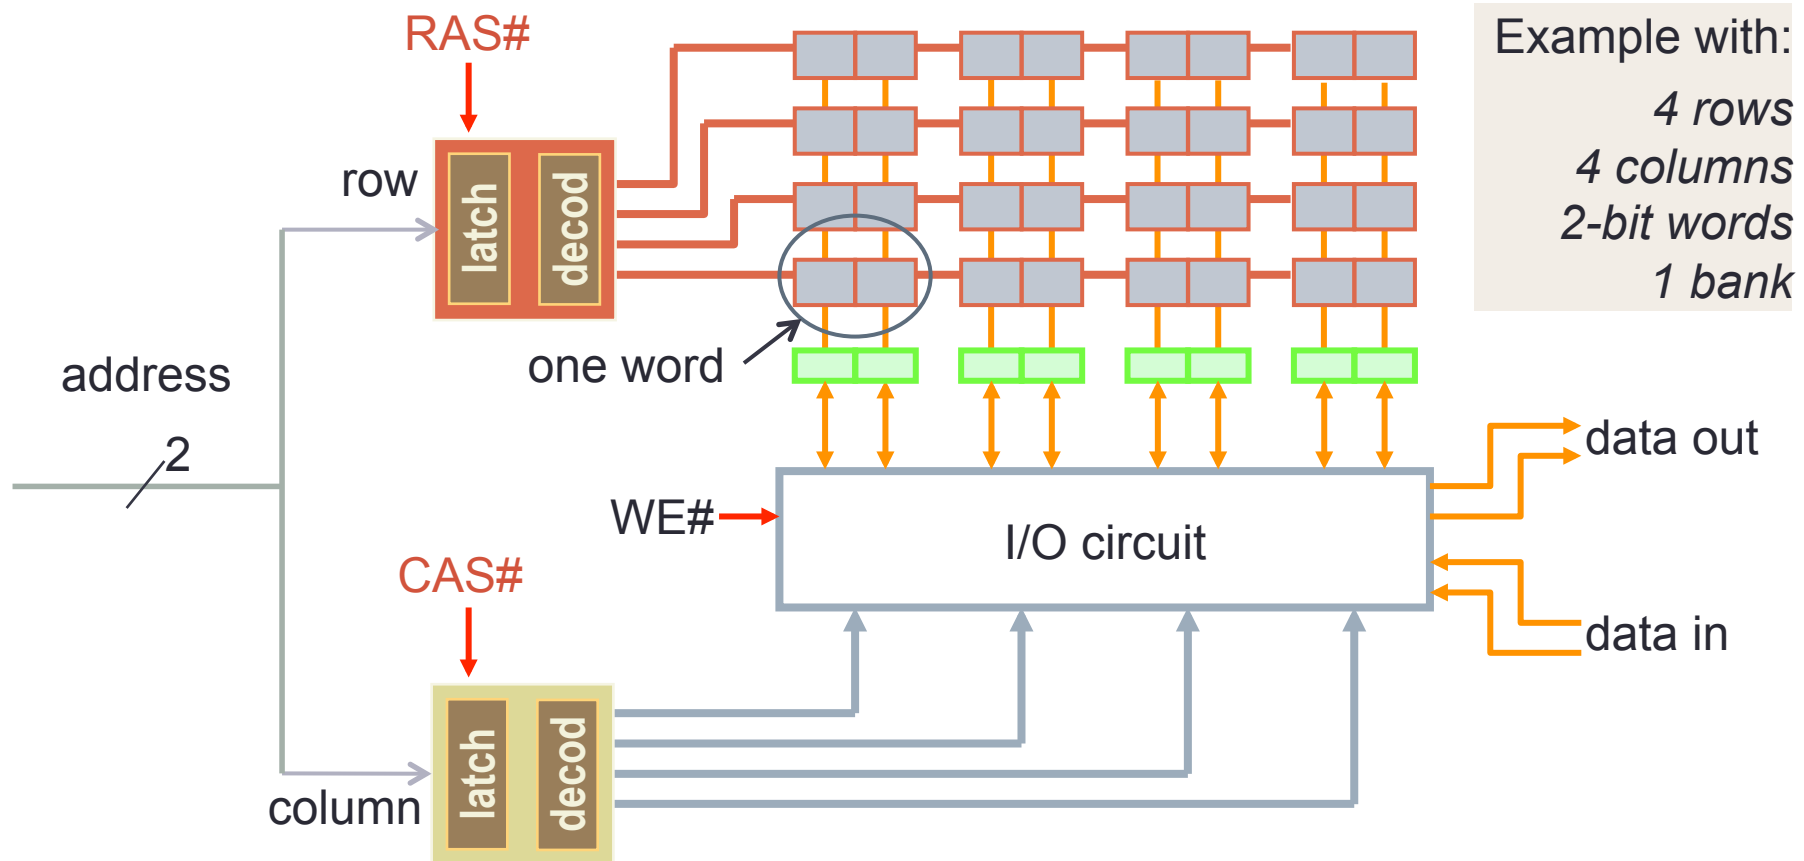
DRAM memory basics

- Cells are internally arranged as a matrix
 - Accessing a cell is a two-step process: row, then column
- All bits in an *open* row are read and latched
- Subsequent bits from the same row can be read from the latches, without a row opening delay
- Additionally, a single refresh cycle is effective for all cells in a row



DRAM memory: internal organization

- Each word is stored in the set of cells in the intersection of a row and a column



DRAM operation: addressing

- Addresses issued to DRAM are structured
 - In a matrix of 2^{r+c} words (2^r rows \times 2^c columns), the $n = r + c$ address bits are organized as:



- Each *linear* address has two components
 - $row = address \text{ div } 2^c$ $column = address \text{ mod } 2^c$
 - Conversely, $address = (row \times 2^c) + column$

row,column form

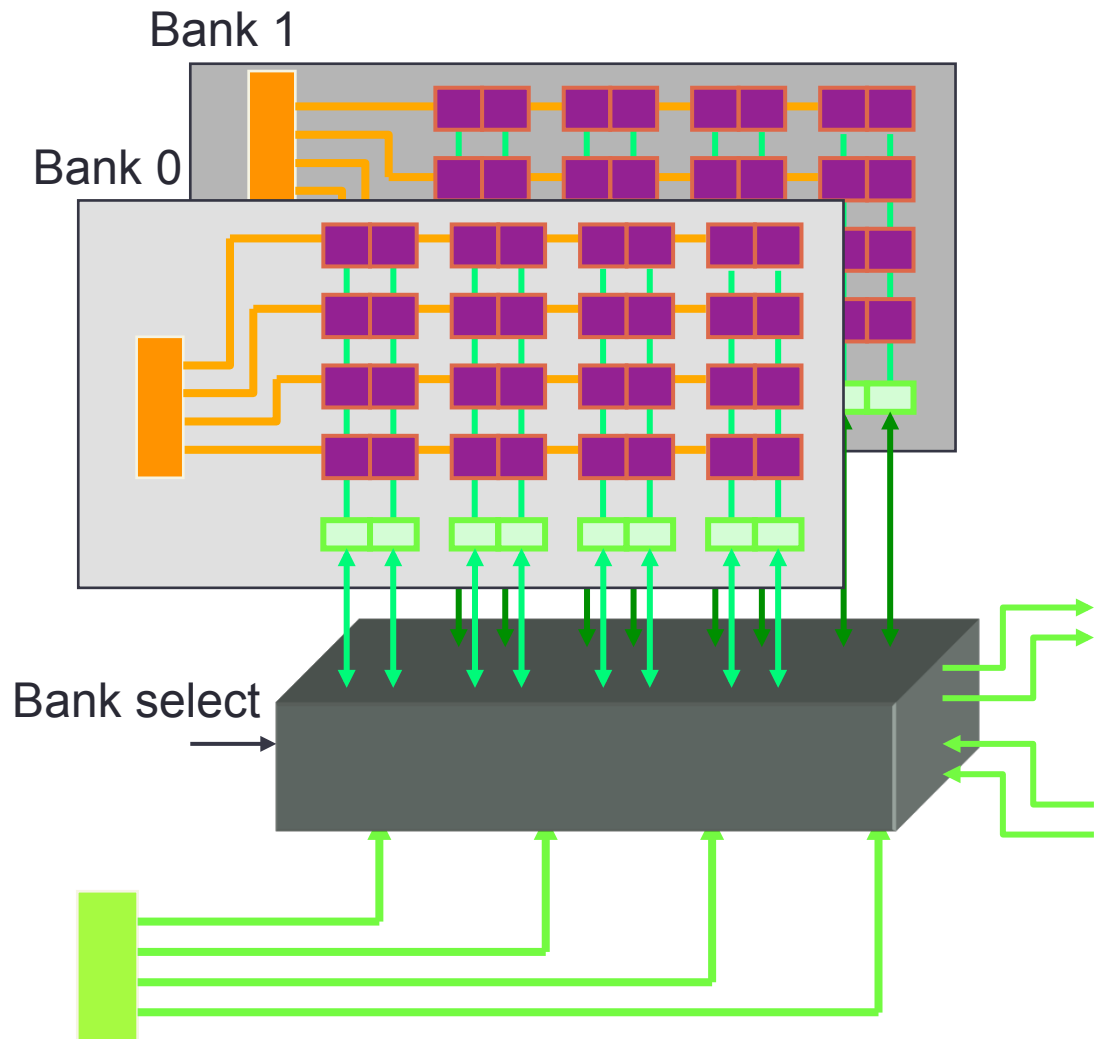
0,0	0,1	0,2	0,2 ^c -1
1,0	1,1	1,2	1,2 ^c -1
2,0	2,1	2,2	2,2 ^c -1
⋮	⋮	⋮		⋮
2 ^r -1,0	2 ^r -1,1	2 ^r -1,2	2 ^r -1,2 ^c -1

linear address form

0	1	2	2 ^c -1
2 ^c	2 ^c +1	2 ^c +2	2 ^{c+1} -1
2 ^{c+1}	2 ^{c+1} +1	2 ^{c+1} +2	2 ^{c+2} -1
⋮	⋮	⋮		⋮
2 ^{n-2^c}	2 ^{n-2^c+1}	2 ^{n-2^c+2}	2 ⁿ -1


DRAM memory: internal organization

- DRAM chips are organized in **banks**
 - Each bank holds a matrix, including amps and row latch
 - Common (shared) I/O circuit
 - Capacity = $B \times 2^r \times 2^c \times w$ bits
 - r row bits, c col bits, B banks
 - Up to B rows may be open at a time
 - There are b bank address lines such that $2^b = B$



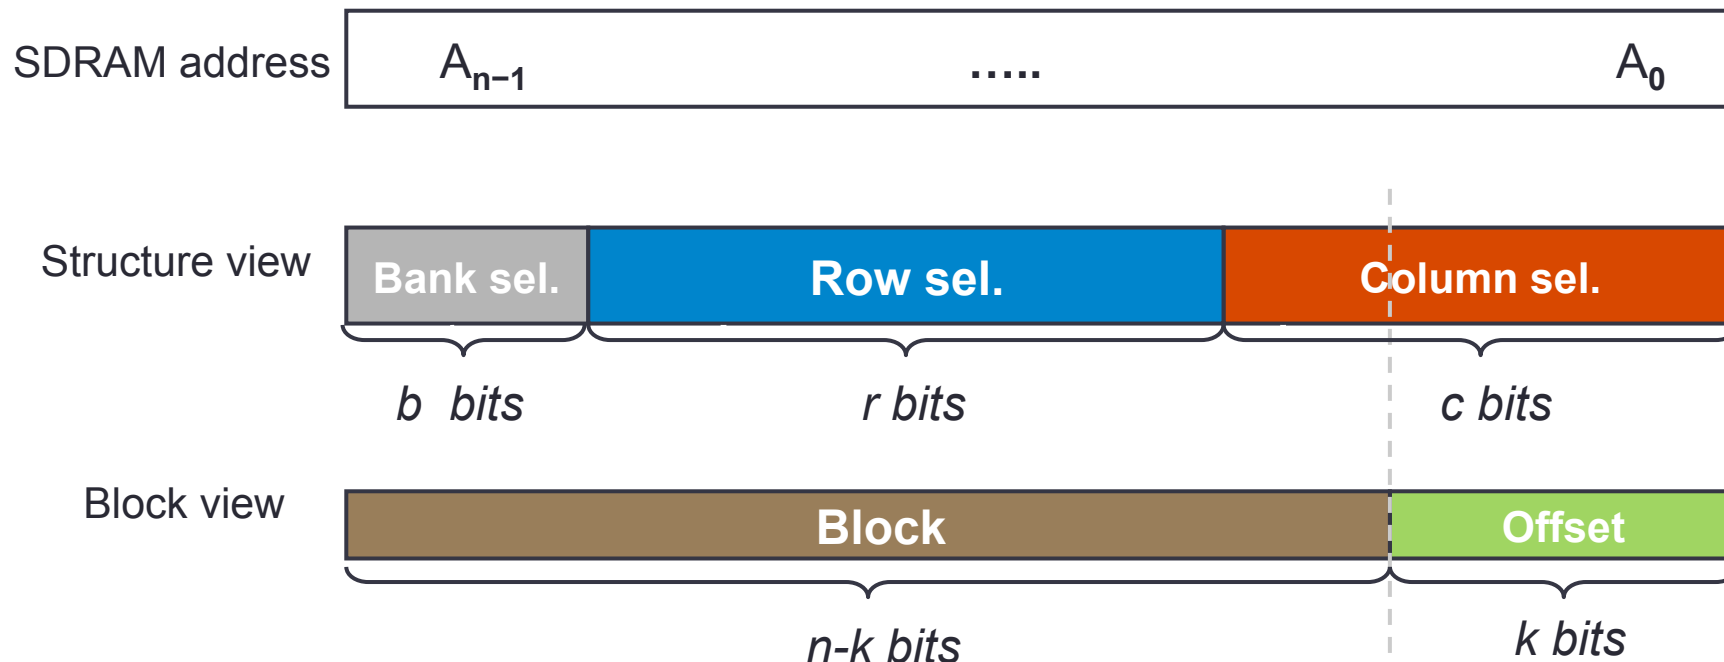
SDRAM memory: external interface

- Global chip select and clock
 - Clock is needed by Synchronous DRAM (**SDRAM**)
- **Bank** and **row/column** form the address
- **Commands** are predefined combinations of three signals
- Byte mask lines allow selective byte enable
- Data I/O: as many lines as the chip's word size



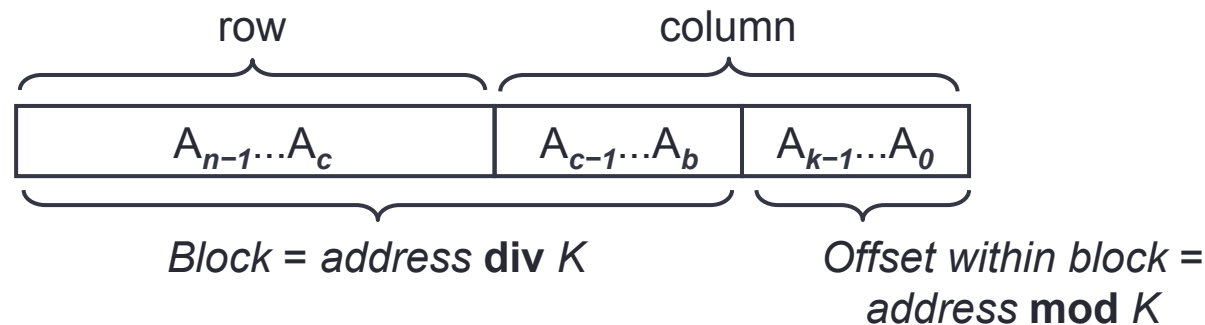
SDRAM addressing

- DRAM size: $2^n \times w$ bits
- Internal structure: 2^b banks \times 2^r rows \times 2^c columns
- Address length: $n = b + r + c$
- Blocks of 2^k words (2^{n-k} blocks in SDRAM)



SDRAM operation: block access

- For efficiency reasons, traffic between cache and main-memory is in the form of blocks
 - Blocks are formed by $K = 2^k$ contiguous words, typically 2, 4 or 8
 - Words in a block use addresses of the form $A_{n-1} \dots A_k x \dots x$
 - Once programmed at setup K remains constant
 - K is typically programmed to fit the *cache line* size (next unit)
 - Blocks always belong in a single row; addresses can thus be further decomposed as:




- Note that a row contains 2^{c-k} blocks

SDRAM operation: block access

• Example

- Capacity 64x8 bits ($W = 8$ b)
- 4 rows of 16 columns (words)
- $K = 4$ (blocks of four 8-bit words)



block 0

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47
48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63

block 9

- Memory components:
 - 16 x 8 sense amplifiers and bit latches
 - 2-bit row latch plus decoder 2:4
 - 4-bit column latch plus decoder 4:16

SDRAM operation: block access

- READ/WRITE commands denote a single word address, but memory operates with the K words of the block containing that word
 - This is also called a **burst access**
- The addressed word goes first (*critical word first*), followed by the rest of the block's words

Example: $K = 4$




Block 0

Address	Sequence
0	0-1-2-3
1	1-2-3-0
2	2-3-0-1
3	3-0-1-2

SDRAM operation: Commands

- Commands are accepted at the rising edges of Clock
- When CS# is high, the command is NOP (No OPeration)



SDRAM operation: Commands

- Commands
 - A combination of CS#, RAS#, CAS#, WE# denote a command
 - The interpretation of address lines depends on the command


Command name	RAS#	CAS#	WE#	Addr
NO OPERATION (NOP)	H	H	H	X
ACTIVATE	L	H	H	Row/Bk.
READ	H	L	H	Column/Bk.
WRITE	H	L	L	Column
PRECHARGE	L	H	L	Code
BURST TERMINATE	H	H	L	X
AUTO REFRESH	L	L	H	X
LOAD MODE REGISTER	L	L	L	Op-Code

SDRAM operation: Commands

- **ACTIVATE**
 - Latch and decode the row address using RAS# and activate the selected row. The row's cells are now connected to the bit lines
 - Sense amps detect and amplify the cells' charge and the whole row contents latched in the row bit latches
 - Access to several columns in the row is possible after a single ACTIVE command
- **READ and WRITE commands**
 - Latch and decode column address using CAS#. WE# then activates the I/O logic for reading or writing the selected column
- **PRECHARGE**
 - Prepares bit lines for next access and takes amps to stand-by. No further access is possible until a new ACTIVATE command


SDRAM operation: block access

- Accessing a memory block requires a sequence of commands
 - ACTIVATE – READ – PRECHARGE
 - ACTIVATE – WRITE – PRECHARGE
 - with additional NOP commands for proper **timing**
- Example: read a 4-word block



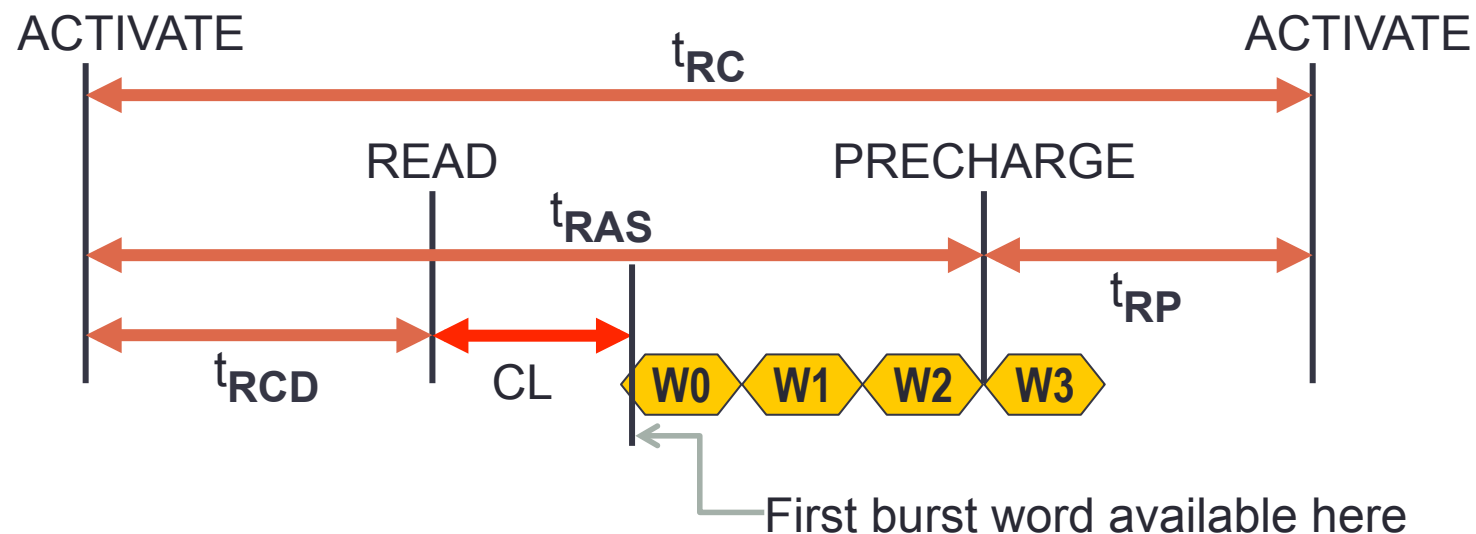
SDRAM Timing

- Main timing parameters
 - t_{RC} : minimum time between two consecutive ACTIVATE commands
 - Equivalent to the memory **cycle time**
 - t_{RCD} : minimum time between ACTIVATE and READ or WRITE
 - t_{RAS} : minimum time between ACTIVATE and PRECHARGE
 - t_{RP} : minimum time between PRECHARGE and next ACTIVATE



SDRAM Timing

- CAS Latency (CL)
 - In a READ operation, CL is the number of cycles by which the first burst item appears in the data outputs
 - CL does NOT apply to WRITE operations
 - Access time = $t_{RCD} + CL$



SDRAM Timing restrictions

- SDRAM chips specify their maximum working frequency (or minimum clock cycle)
 - The chip can be used at lower frequencies
- A proper number of NOP commands must be inserted to comply with timing requirements (t_{RC} , t_{RAS} , t_{RP} , t_{RCD} , CL)
- Example: number of cycles to accommodate timing requirements at different clock speeds (t_{CK} = clock period)

	50 MHz ($t_{CK} = 20$ ns)	100 MHz ($t_{CK} = 10$ ns)	166 MHz ($t_{CK} = 6$ ns)
$t_{RCD} = 18$ ns	1	2	3
$t_{RAS} = 42$ ns	3	5	7
$t_{RC} = 60$ ns	3	6	10
$t_{RP} = 18$ ns	1	2	3


SDRAM Timing

- Clock frequency and CAS latency
 - CL is programmed at memory startup
 - For a given clock frequency the chip has a minimum CL
 - Conversely, there is a maximum frequency for each CL

PARAMETER		SYMBOL	MIN	MAX	UNITS
Clock cycle time	CL = 3	$t_{CK(3)}$	6		ns
	CL = 2	$t_{CK(2)}$	10		ns
	CL = 1	$t_{CK(1)}$	20		ns
ACTIVE to PRECHARGE command		t_{RAS}	42	120k	ns
ACTIVE to ACTIVE command period		t_{RC}	60		ns
AUTO REFRESH period		t_{RFC}	60		ns
ACTIVE to READ or WRITE delay		t_{RCD}	18		ns
Refresh period (4,096 rows)		t_{REF}		64	ms
PRECHARGE command period		t_{RP}	18		ns
ACTIVE bank to ACTIVE bank command		t_{BDP}	12		ns

SDRAM Timing

- RD/WR operations always involve the K words in a burst
 - Once the row is open, column access is relatively fast
- Time to access a block in a burst has two components:
 - Latency to access the first word (independent of K)
 - Time to access subsequent words (depends on K and clock rate)
 - The chip's bandwidth ignores the first-word latency




SDRAM Timing: transfer timing

- First word of a READ appears CL cycles after command
- Write transfers start when WRITE is issued
- The remaining words in the burst are driven by the clock, one word per cycle
- t_{RAS} imposes a minimum delay for issuing PRECHARGE, but it may need to be further delayed when the burst is large
 - A premature PRECHARGE command would abort the burst access
 - PRECHARGE is effective CL cycles after issuing. Optimal time for issuing PRECHARGE is CL – 1 cycles before reading the last word (equivalent to K cycles after READ, for a burst size of K)
- READ and WRITE can be issued in AUTO PRECHARGE mode
 - The chip is then in charge of self-issuing PRECHARGE at the optimum time, based on the programmed burst size, and saves the need for explicit precharge commands


SDRAM Timing: Chronograms

- Example: READ with $f_{\text{CLK}} = 100 \text{ MHz}$; $\text{CL} = 2$



SDRAM Timing: Chronograms


- Chaining two block reads from the same row (with auto-precharge)



Total read latency = $4 \times 10 \text{ ns} = 40 \text{ ns}$
 Bandwidth = $100 \times 10^6 \times w \text{ bits/s}$


SDRAM Timing: Chronograms

- Multiple bank READs
 - Independent banks may receive overlapping commands
 - Data are obtained in request order





SDRAM Timing: Chronograms

- WRITE cycle with $f_{\text{CLK}} = 100 \text{ MHz}$, auto-precharge mode

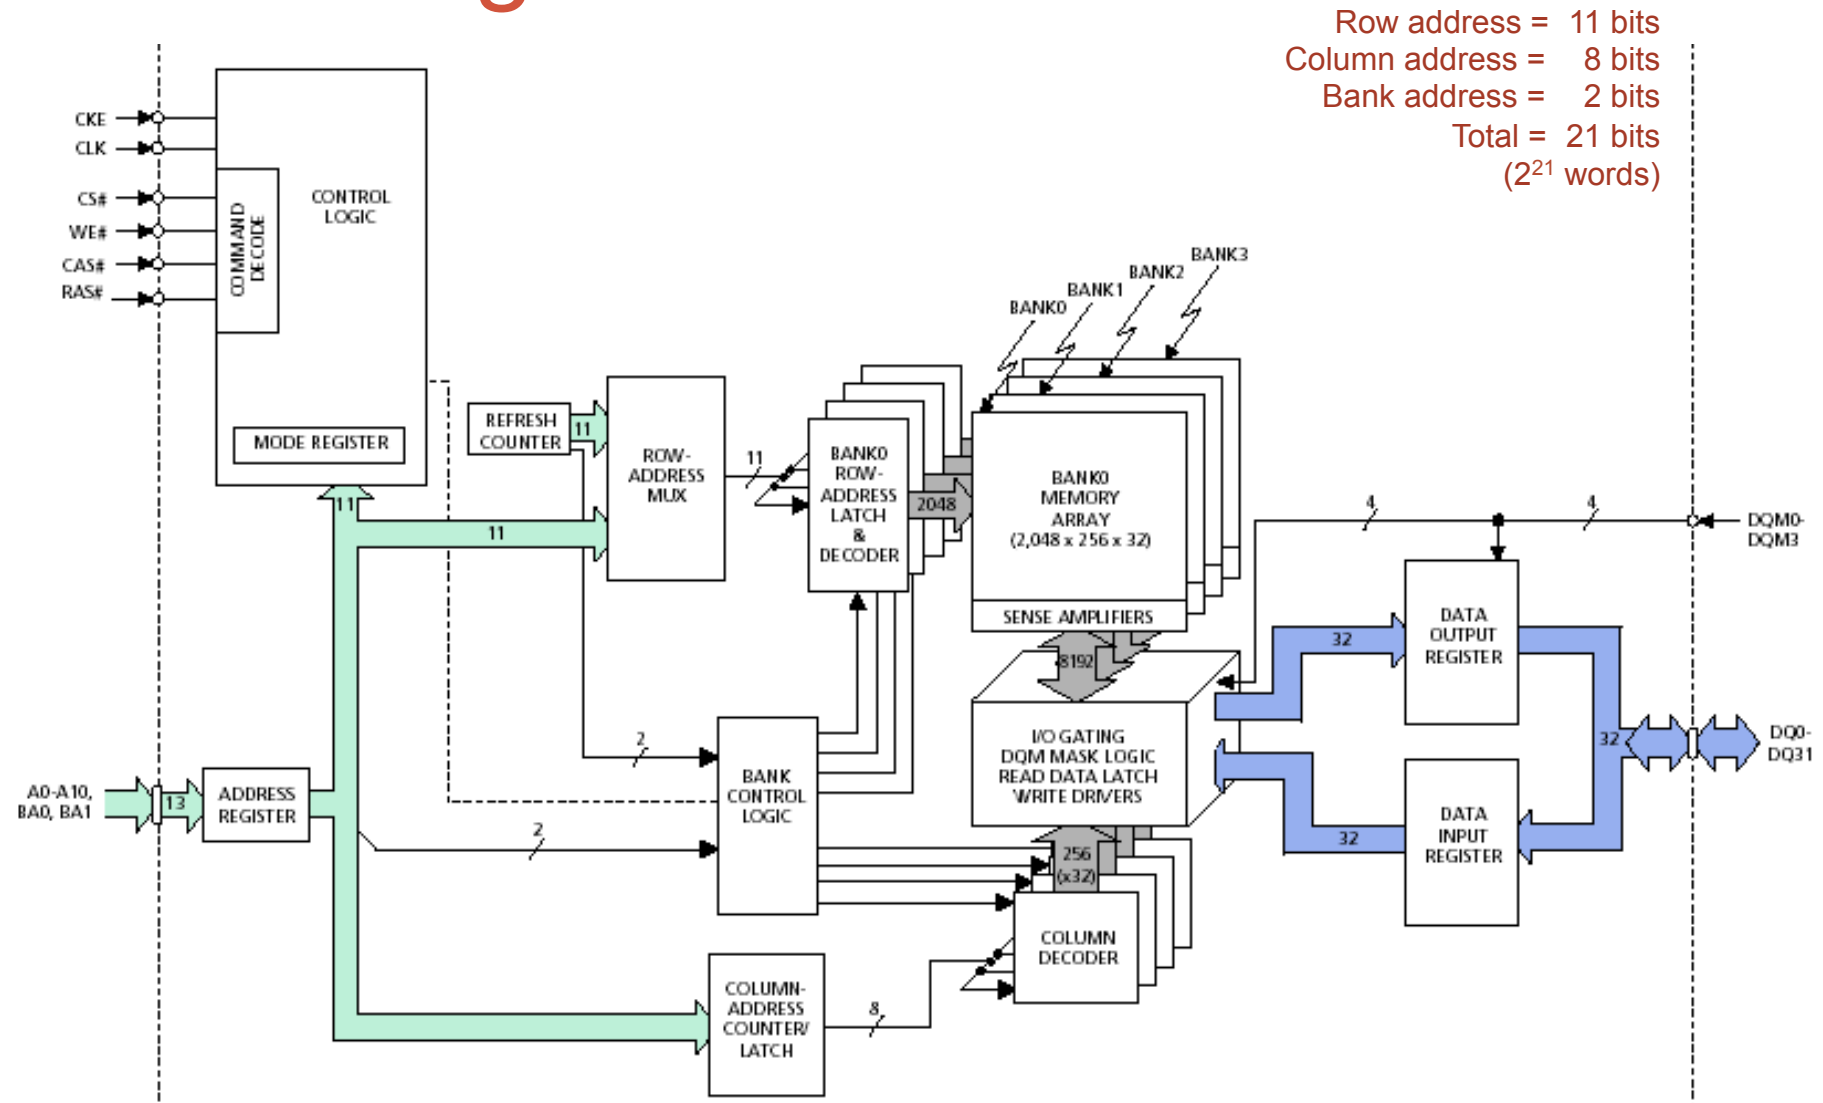


A commercial example: 2M × 32 b

- Model MT49LC2M32B2 (micron)
- 86-pin plastic package
- Clock frequencies and cycle times
 - (143/166/183/200) MHz resp. (7/6/5.5/5) ns
- Programmable CAS latency
 - 143 and 166 MHz → 1, 2 or 3 cycles
 - 183 and 200 MHz → 3 cycles
- Programmable burst length
 - 1, 2, 4, 8 or 256 32-bit words
- Four banks of 2048×256×32 b
 - $2 \text{ M} \times 32 \text{ b} = 2^{26} = 4 \times 2^{24} = 4 \times (2^{11} \times 2^8 \times 32) \text{ b}$
- Refresh period: 64 ms



Block Diagram



Improvements to SDRAM technology

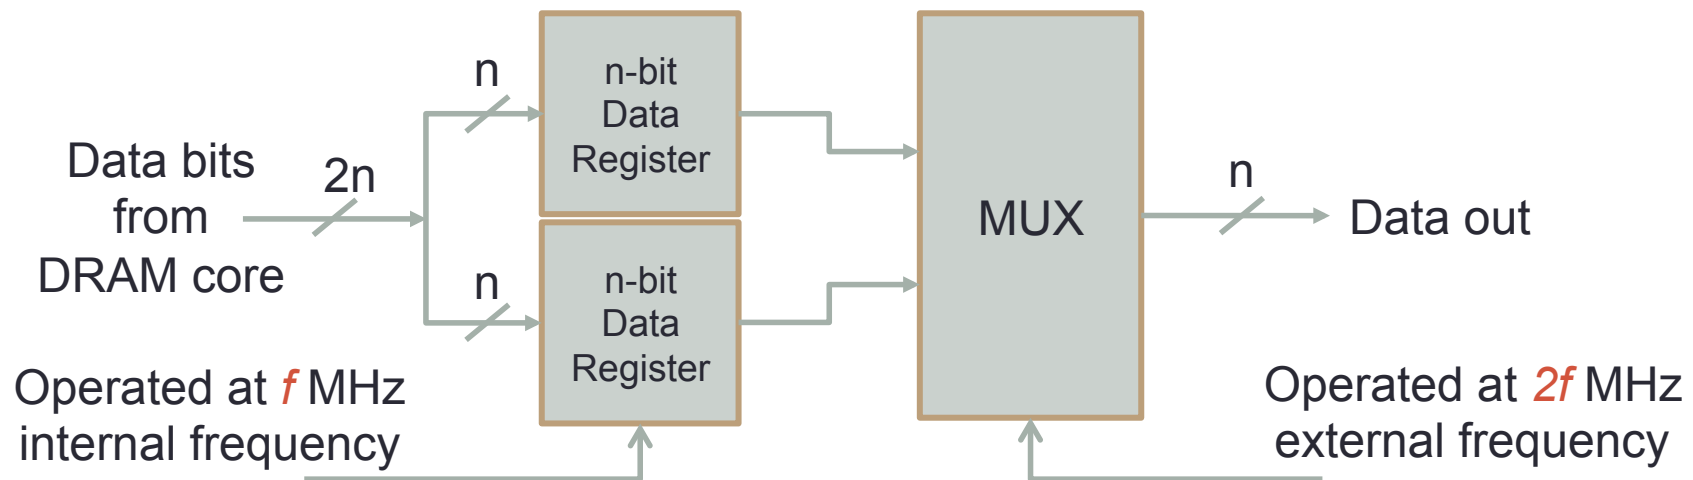
- Memory implementation trends:
 - Higher density
 - Lower power voltage
 - Heat is proportional to V^2 and frequency, and it limits density
 - Higher bandwidth based on structural improvements (DDR1/2/3/4)
 - Latencies to access the internal matrix diminish very slowly with time

Improvements to SDRAM: DDR

- DDR stands for *Dual Data Rate*
 - The JEDEC (Joint Electron Device Engineering Council) maintains standard specification of SDRAM chips
- In DDR, data can be read/written at twice the memory's internal frequency, effectively achieving double bandwidth with respect to SDR SDRAM (Single Data Rate SDRAM)
- DDR is based on the *2n-prefetch* technique
 - A single access to the cell matrix provides $2n$ bits for an n -bit chip
 - The internal data bus is twice the size of the external
 - Data is externally accessed at a rate of 2 words per cycle
- DDR memories operate with two clocks: one internal for accessing the cell matrix and one external for data transferring

Improvements to SDRAM: DDR

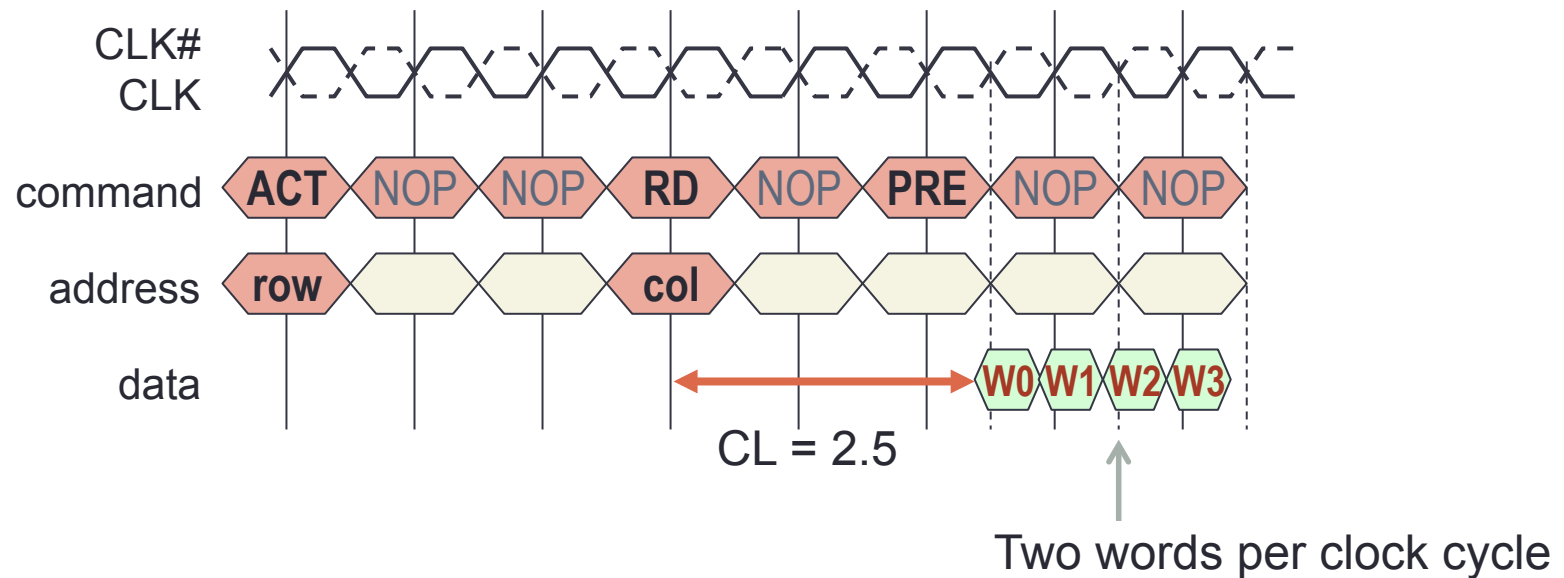
- Simplified block diagram for $2n$ -prefetch READ



- DDR2 and DDR3 use $4n$ - and $8n$ -prefetch for accessing larger bursts at higher frequencies, for larger bandwidth
- DDR4 further reduces power consumption and increases data rate by grouping banks, but still using $8n$ -prefetch

Improvements to SDRAM: DDR

- DDR READ chronogram
 - CAS latency is a multiple of 0.5 clock cycles
 - The external interface includes a CLK# input



SDRAM Standard Naming

- According to the JEDEC standard, SDRAM chips are named after their achievable bandwidth
 - SDR (non-DDR SDRAM): **PC F**
 - F is the chip's clock frequency, hence a direct indication of bandwidth when multiplied by the chip's number of data pins
 - DDR: **DDR_X MTps**
 - X can be "", "2", "3" or "4" for DDR, DDR2, DDR3 and DDR4, resp.
 - MTps (or MT/s) stands for Mega (Millions of) Transactions per Second
 - It is a measure of data pin bandwidth – chip's bandwidth depends on MTps and the number of data pins
 - Examples:
 - Micron's MT46V32M8-5B is a **32M×8b DDR-400**
 - With a 200 MHz clock, it can achieve up to 400 MTps
 - Bandwidth of this 8-bit wide chip is thus: $400 \text{ MT/s} \times 1 \text{ B/T} = 400 \text{ MBps}$
 - Micron's MT41K64M16-125 is a **64M×16b DDR3-1600**
 - 800 MHz clock, 1600 MTps, Bandwidth = $1600 \text{ MT/s} \times 2 \text{ B/T} = 3200 \text{ MBps}$

SDRAM Standard Naming

- Variations of SDRAM chips

Name	Internal clock (MHz)	Prefetch (min burst)	External clock (MHz)	Transfer rate (MT/s)	Voltage (V)
SDR	66-133	1n	66-133	66-133	3.3
DDR	100-200	2n	100-200	200-400	2.5-2.6
DDR2	100-267	4n	200-533	400-1067	1.8
DDR3	100-267	8n	400-1066	800-2133	1.35-1.5
DDR4	133-267	8n	1066-2133	2133-4267	1.05-1.2

(*) Chips exist that deviate from these specifications and limits

5. Memory Modules

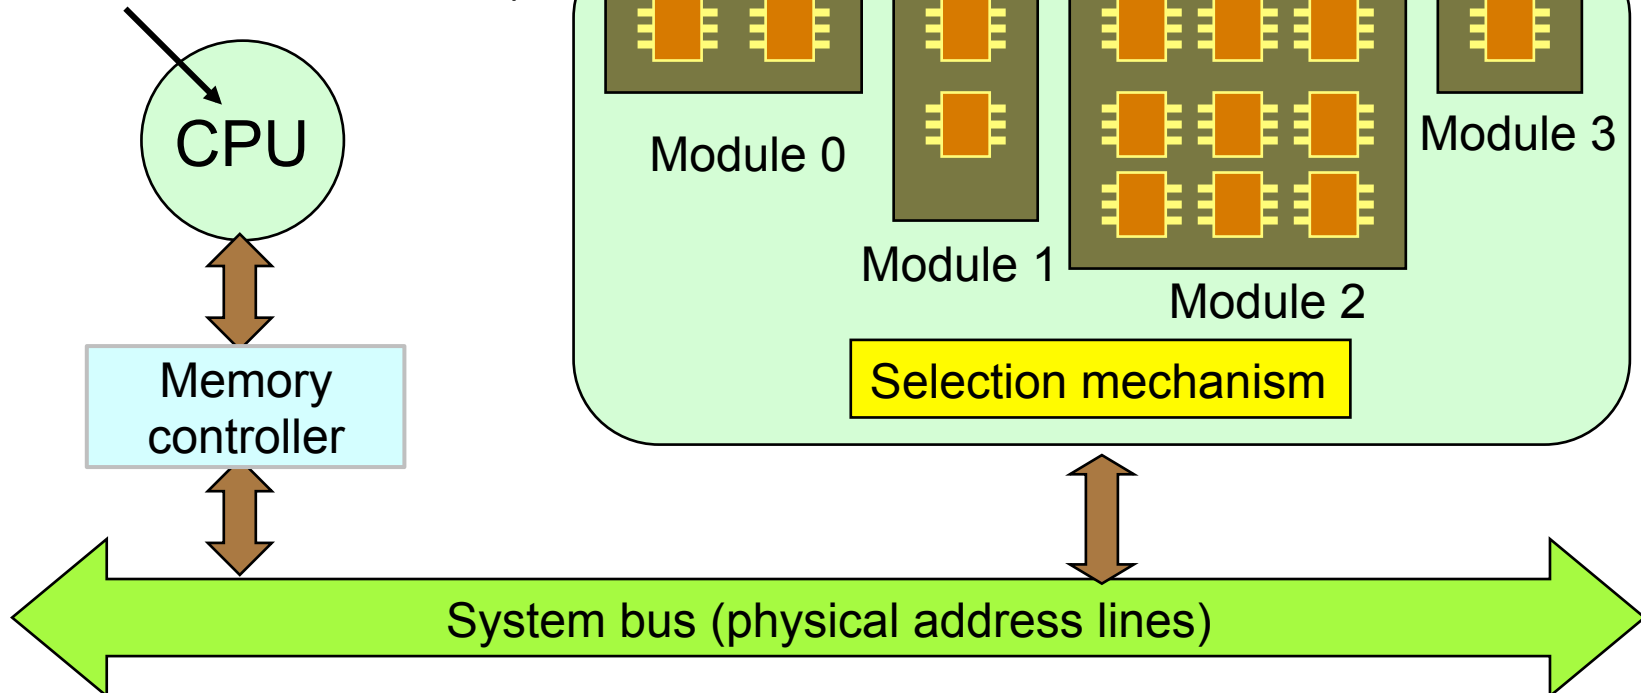
- Basic concepts
- Internal organization
- Commercial DRAM modules

Basic concepts

- The main memory system is typically composed of several memory modules
 - A **memory module** is a circuit composed by memory chips that are grouped to provide an adequate number and size of words
- Each module contains a particular type of memory chips (ROM, SRAM, DRAM, etc.)
- Memory modules tend to comply with a standard specification (see later)
- In principle, only one module may be selected at a time
 - Each module supports a range of addresses
 - A selection mechanism takes care of proper module selection
 - The distribution of addresses per module is called **memory map** (see later)

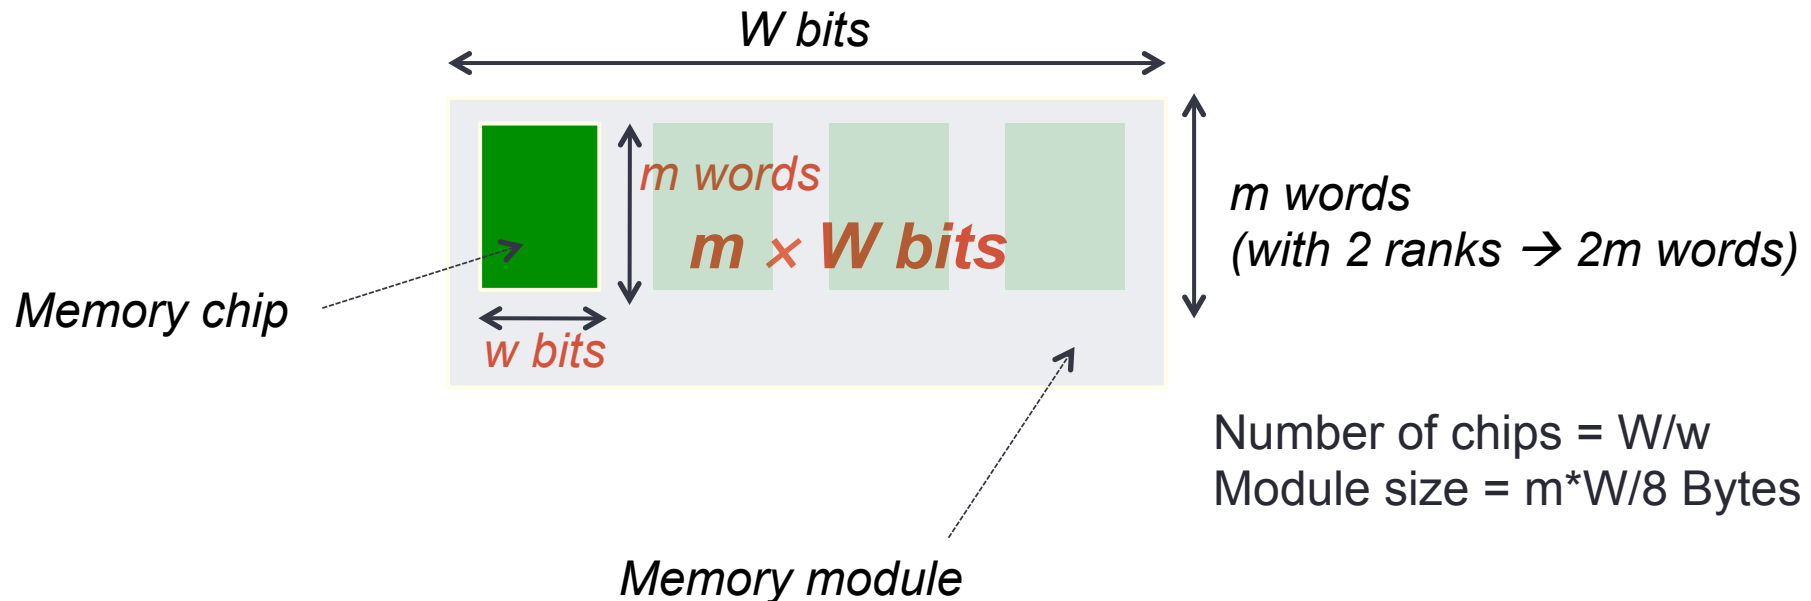
Basic concepts

Address bits
(logical or effective address)

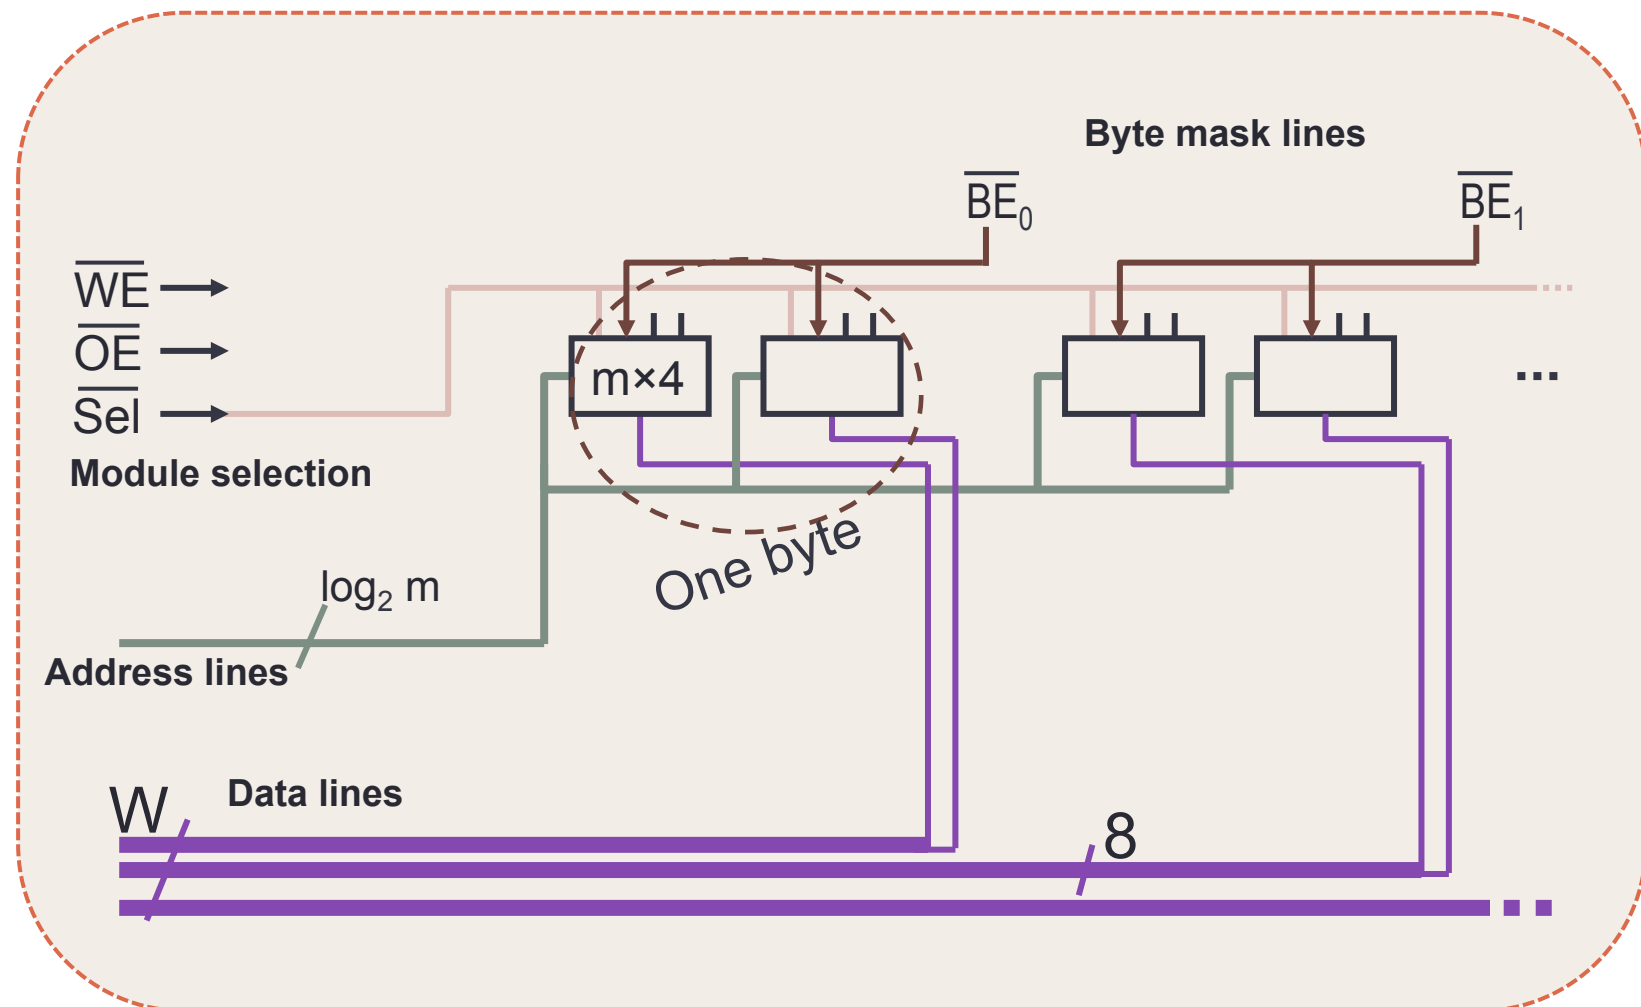


Internal organisation of modules

- For simplicity, modules are built with identical chips
- Memory chips are typically arranged as one or two rows (called *ranks*)
 - More ranks \rightarrow more words; more columns \rightarrow wider words
 - It is convenient to express capacities as *words* \times *word_size*




Internal organisation of modules



Internal organisation of modules

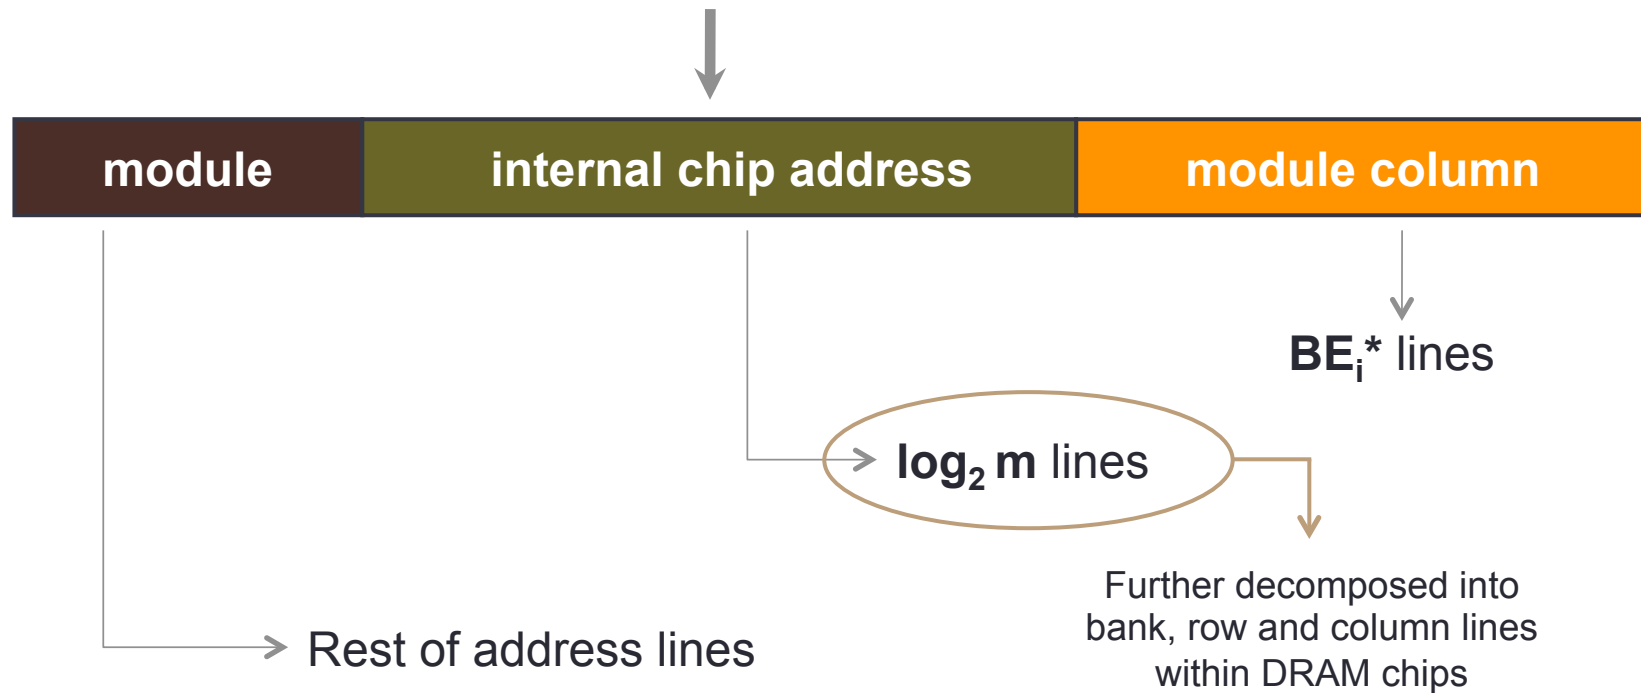
- Connections
 - CK, selection, command and address lines are common to all chips
 - Data mask lines enable selection of individual bytes
 - Total bandwidth equals the sum of chips' bandwidth
- Example: module of $m \times 32$ b based on $m \times 4$ chips



Internal organisation of modules



- Address structure with modules

Physical address bus (e.g., 64-bit data bus)



DRAM Modules

- Modules are inserted into motherboard's sockets
- Some modules include one additional parity bit per byte for error detection
- Example: $m \times 64$ b





Standard DRAM Modules

- A standard definition facilitates module interchange
 - Modules easily fit standard slots in the motherboard
 - Optionally, modules include one additional parity bit per byte to enable error detection
 - Examples




Format	Word size	Chip type	Capacity
DIMM 168 contacts	64/72 bits	EDO, SDRAM	4/.../512 MB
SODIMM 144 c.	64 bits	SDRAM	8/.../2 GB

Note: SODIMM = *Small Outline* DIMM (used for portable computers)



DIMM Modules

- DIMM (Dual Inline Memory Module)
 - 168/184/240/288 contacts, 64(72) data bits, 5¼" long (13.335 cm)



SPD: Serial Presence Detect

- EEPROM memory chip
- Stores a few bytes of relevant information about the module
 - Timing
 - Capacity
 - Manufacturer
 - Serial number
- Facilitates automatic configuration of the main memory system
- Some programs may access the information contained in the SPD
 - CPU-Z, SiSoftware Sandra, ...



SPD: View from CPU-Z

The screenshot shows the CPU-Z application window with the SPD tab selected. The window title is "CPU-Z". The tabs are CPU, Cache, Mainboard, Memory, SPD, and About. The SPD tab displays the following information:

Memory Slot Selection

Slot #1 (dropdown) | DDR

Module Size	512 MBytes	Correction	None
Max Bandwidth	PC2700 (166 MHz)	Registered	no
Manufacturer	Samsung	Buffered	no
Part Number		SPD Ext.	
Serial Number	FFFFFFFF	Week/Year	255 / 255

Timings Table

Frequency	133 MHz	166 MHz		
CAS# Latency	2.0	2.5		
RAS# to CAS#	3	3		
RAS# Precharge	3	3		
tRAS	6	7		
tRC				
Command Rate				
Voltage	2.5 V	2.5 V		

Version 1.47

CPU-Z OK

Comparative and denominations

- Packaging

- SDRAM: DIMM 168 c
- DDR SDRAM: DIMM 184 c
- DDR2 SDRAM: DIMM 240 c
- DDR3 SDRAM: DIMM 240 c

- Bandwidth and naming

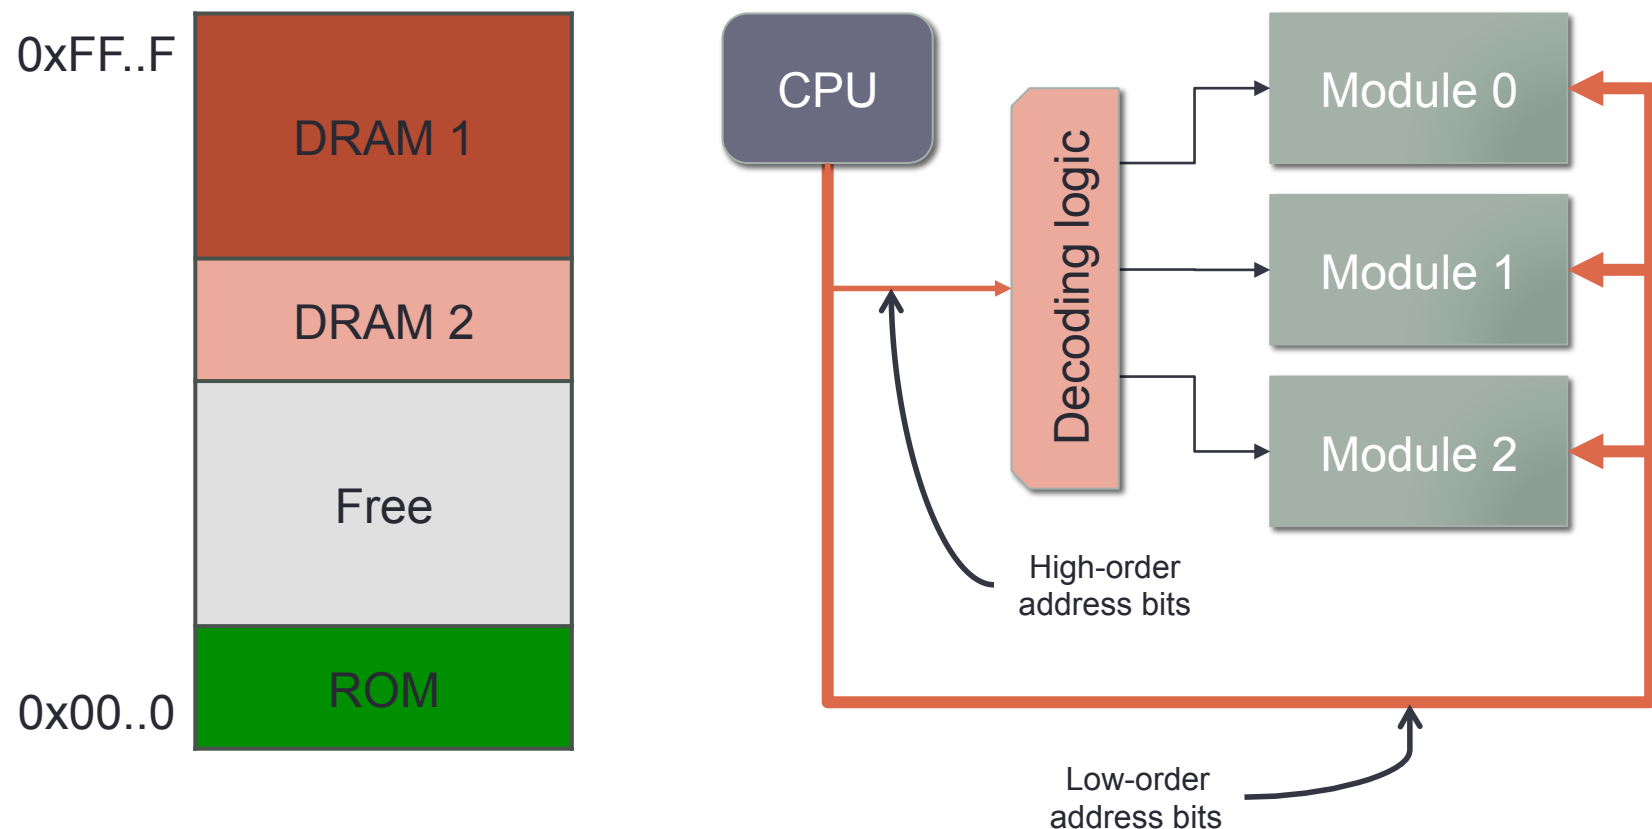
- SDRAM @ 100 MHz $\rightarrow 100 \text{ MHz} \times 8 \text{ B} = 800 \text{ MB/s}$
 - Denomination: PC100
- DDR SDRAM @ 100 MHz $\rightarrow 100 \text{ MHz} \times 8 \text{ B} \times 2 = 1600 \text{ MB/s}$
 - Denomination: PC 1600
- DDR2 SDRAM @ 200 MHz $\rightarrow 200 \text{ MHz} \times 8 \text{ B} \times 2 = 3200 \text{ MB/s}$
 - Denomination: PC2-3200
- DDR3 SDRAM @ 400 MHz $\rightarrow 400 \text{ MHz} \times 8 \text{ B} \times 2 = 6400 \text{ MB/s}$
 - Denomination: PC3-6400

6. Memory Map

- Concept
- Map decoding

Concept of memory map

- The memory map is the distribution of modules across the addressing space



Map decoding

- To implement the map decoding logic we need to know:
 - CPU addressing space, i.e., nr. of lines in the physical address bus
 - In fact, we only need to know what's the highest address bit
 - Modules' initial addresses
 - Modules' sizes (in bytes)
 - The internal organization of the module doesn't matter
 - Logic (positive or negative) of module select functions

A simple example


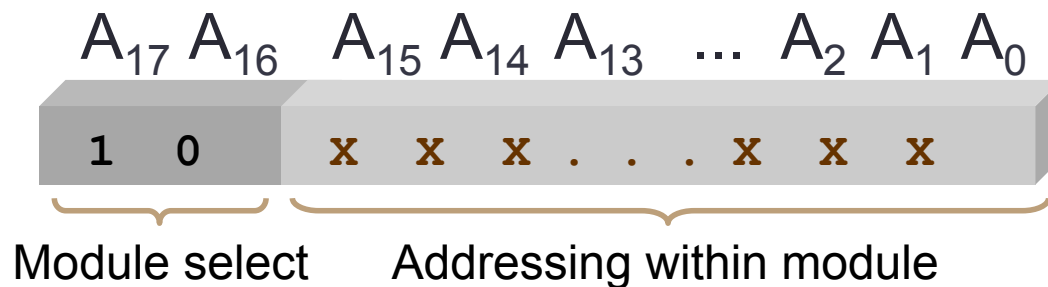
- CPU with addressing space of 256 KB, $W = 32$ b
- 64 KB DRAM module at address 0x20000
- We are asked to design the logic to implement this map
- First, see what addresses are used by the module

A_{17}	A_{16}	A_{15}	A_{14}	A_{13}	...			A_2	A_1	A_0	
1	0	0	0	0	.	.	.	0	0	0	0x20000
1	0	0	0	0	.	.	.	0	0	1	0x20001
1	0	0	0	0	.	.	.	0	1	0	0x20002
					.	.	.				
1	0	1	1	1	.	.	.	1	1	0	0x2FFFE
1	0	1	1	1	.	.	.	1	1	1	0x2FFFF

- All addresses in the module have in common $A_{17}=1$ and $A_{16}=0$

A simple example

- CPU with addressing space of 256 KB, $W = 32$ b
- 64 KB DRAM module at address 0x20000
- We are asked to design the logic to implement this map
- Second, use the regularity in module's addresses to determine the SEL function
 - All addresses in the module have the form 0x2XXXX




- Third, the functions

$$\begin{aligned} \text{SEL} &= A_{17} \cdot A_{16}^* \\ \text{SEL}^* &= A_{17}^* + A_{16} \end{aligned}$$

A simple example

- CPU with addressing space of 256 KB, $W = 32$ b
 - 64 KB DRAM module at address 0x20000
 - We are asked to design the logic to implement this map
-
- Finally, consider the physical lines




Example with two modules

- MIPS R2000: 4 GB addressing (2^{32} B)
 - Module RAM1: 1 GB starting at 0x00000000
 - Module RAM2: 2 GB starting at 0x80000000



Example with two modules: RAM1

	A_{31}	A_{30}	A_{29}	A_{28}	A_{27}	...	A_4	A_3	A_2	A_1	A_0
0x00000000	0	0	0	0	0	. . .	0	0	0	0	0
0x00000001	0	0	0	0	0	. . .	0	0	0	0	1
0x00000002	0	0	0	0	0	. . .	0	0	0	1	0
...											
0x3FFFFFFE	0	0	1	1	1	. . .	1	1	1	1	0
0x3FFFFFFF	0	0	1	1	1	. . .	1	1	1	1	1





A_{31} and A_{30} select the module

Module addressing

$$SEL_{RAM1}^* = A_{31} + A_{30}$$

Example with two modules: RAM2

	A_{31}	A_{30}	A_{29}	A_{28}	A_{27}	...	A_4	A_3	A_2	A_1	A_0
0x80000000	1	0	0	0	0	.	0	0	0	0	0
0x80000001	1	0	0	0	0	.	0	0	0	0	1
0x80000002	1	0	0	0	0	.	0	0	0	1	0
...											
0xFFFFFFFFE	1	1	1	1	1	.	1	1	1	1	0
0xFFFFFFFFF	1	1	1	1	1	.	1	1	1	1	1



A_{31} selects module

Module addressing

$$SEL_{RAM2}^* = A_{31}^*$$

7. The Memory Controller


- Description and functions
- Initialization of the memory system

Description and functions


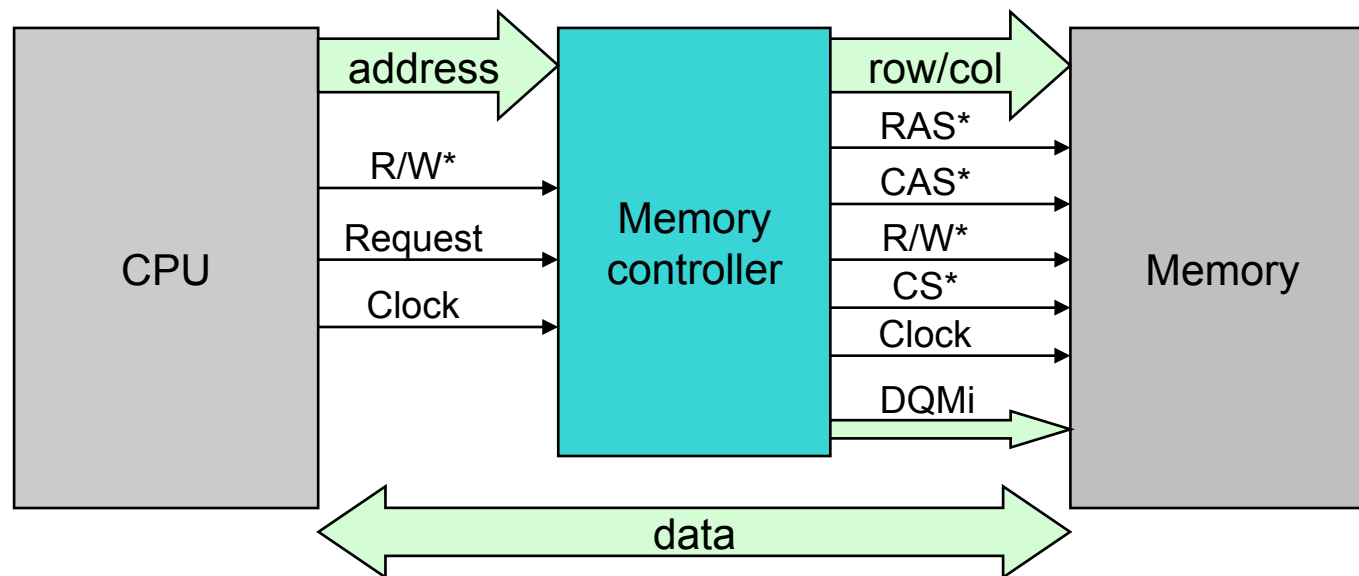
- The memory controller is the circuit in charge of handling the data flow to/from main memory
- It can be a separate chip (for flexibility) or integrated in a microprocessor (for minimum memory latency)
- Its main functions are
 - To identify and properly **initialize** the memory modules
 - To **timely** supply addresses and commands
 - To **translate** CPU addresses to modules' row, col and bank lines
 - To transfer BE* lines to **DQM*** inputs
 - To ensure DRAM is constantly **refreshed** within the refresh period
 - To support interfacing to multiple memory modules (**map**)
 - Ideally, to **maximize the throughput** of memory transfers

Description and functions

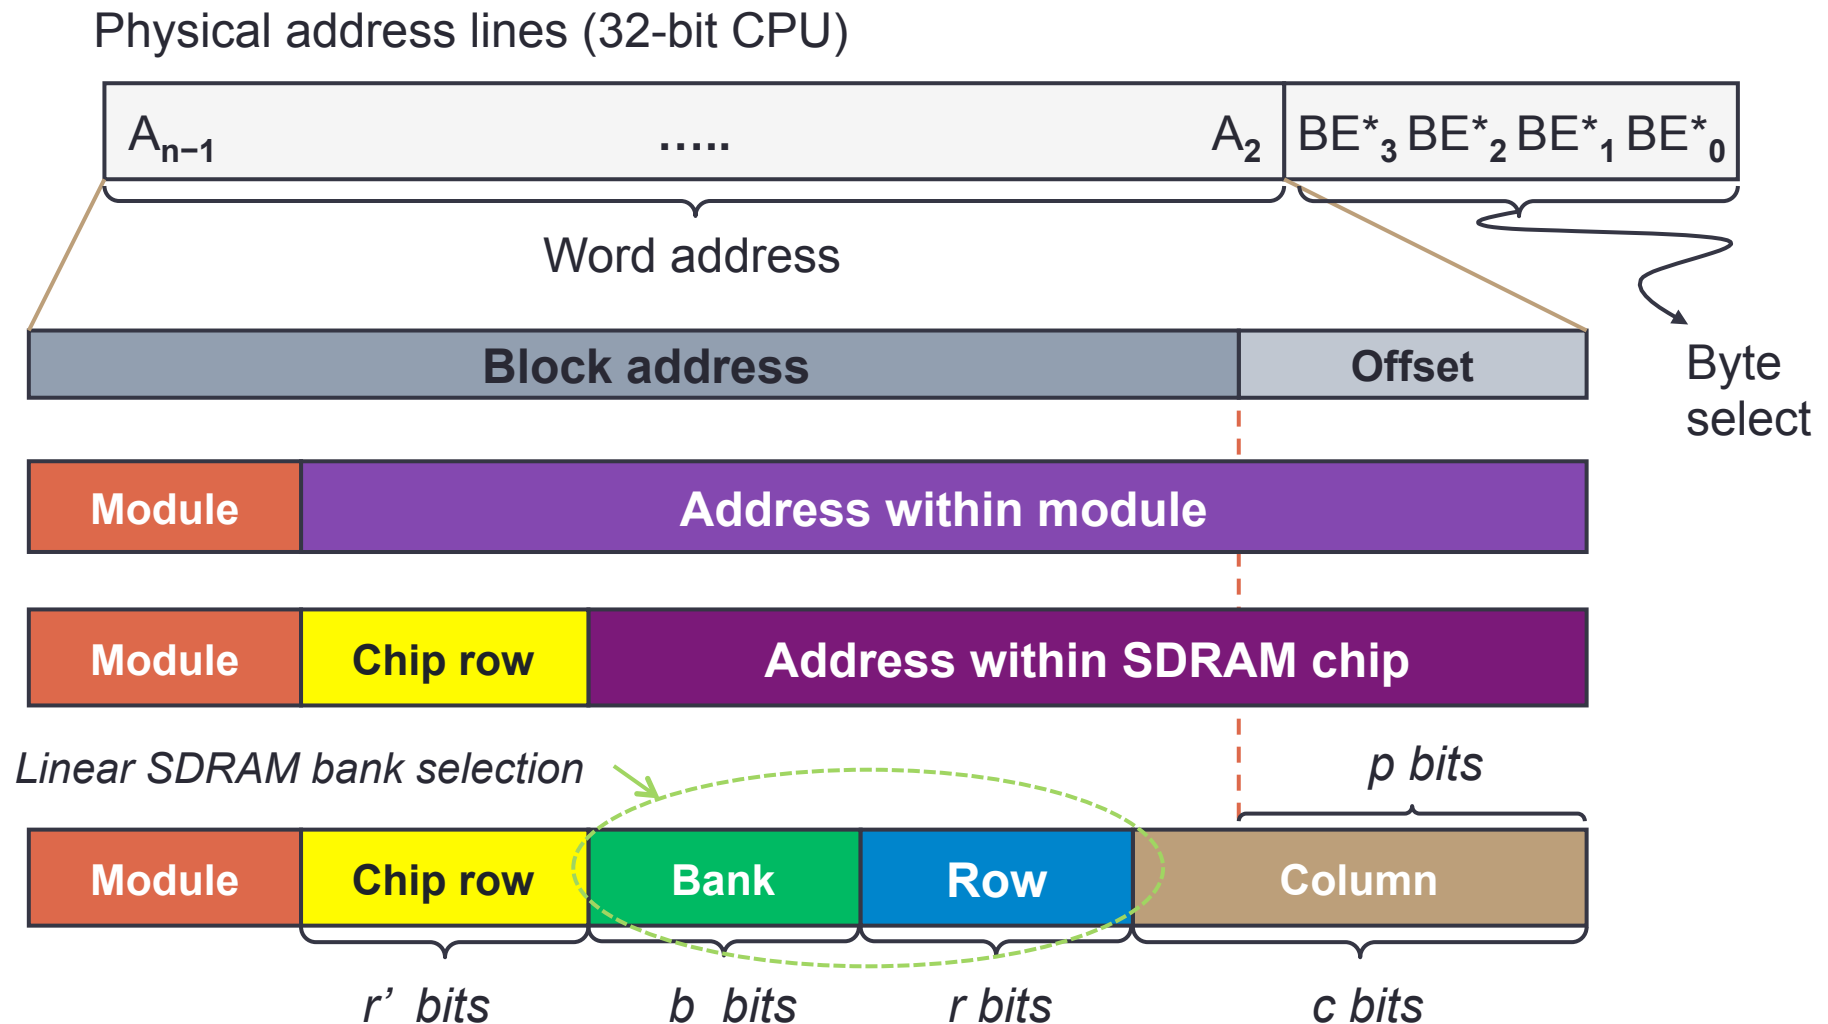
- Maximizing bandwidth
 - Support for one or more independent **channels**
 - Total bandwidth is the sum of each channel's bandwidth – when all channels can work in parallel
- The controller may impose restrictions on the distribution of modules across channels
 - You need to know these restrictions before extending the main memory



Description and functions

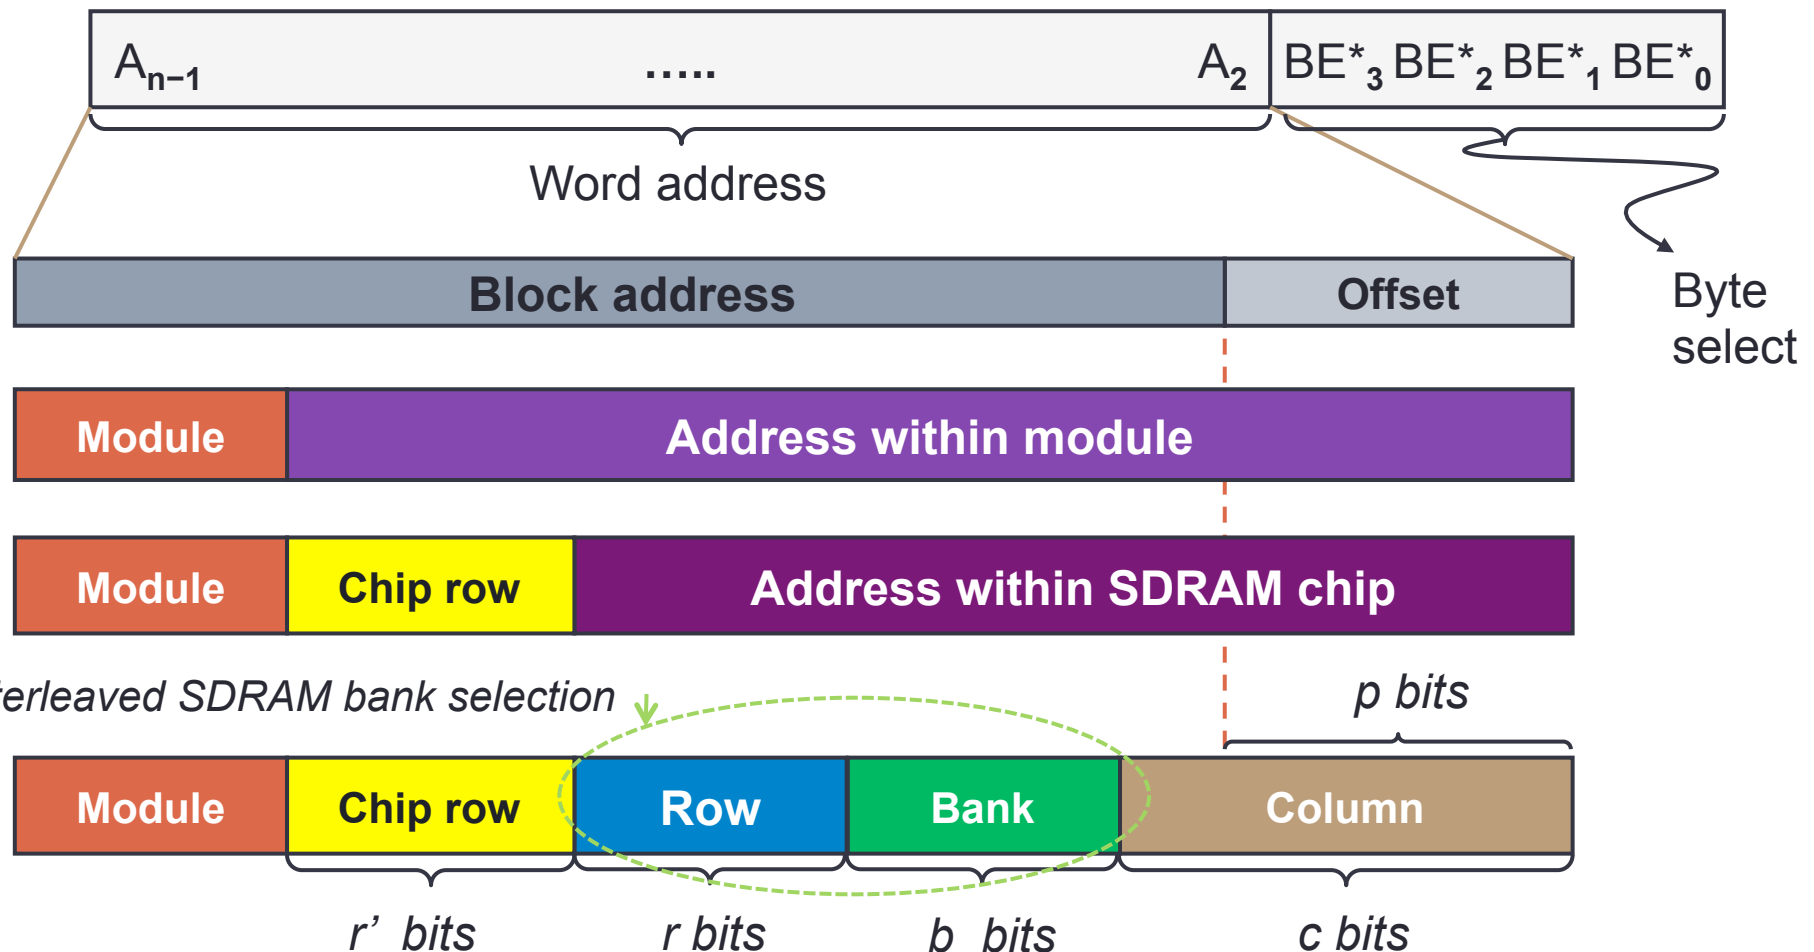


Memory controller: address handling



Memory controller: address handling

Physical address lines (32-bit CPU)



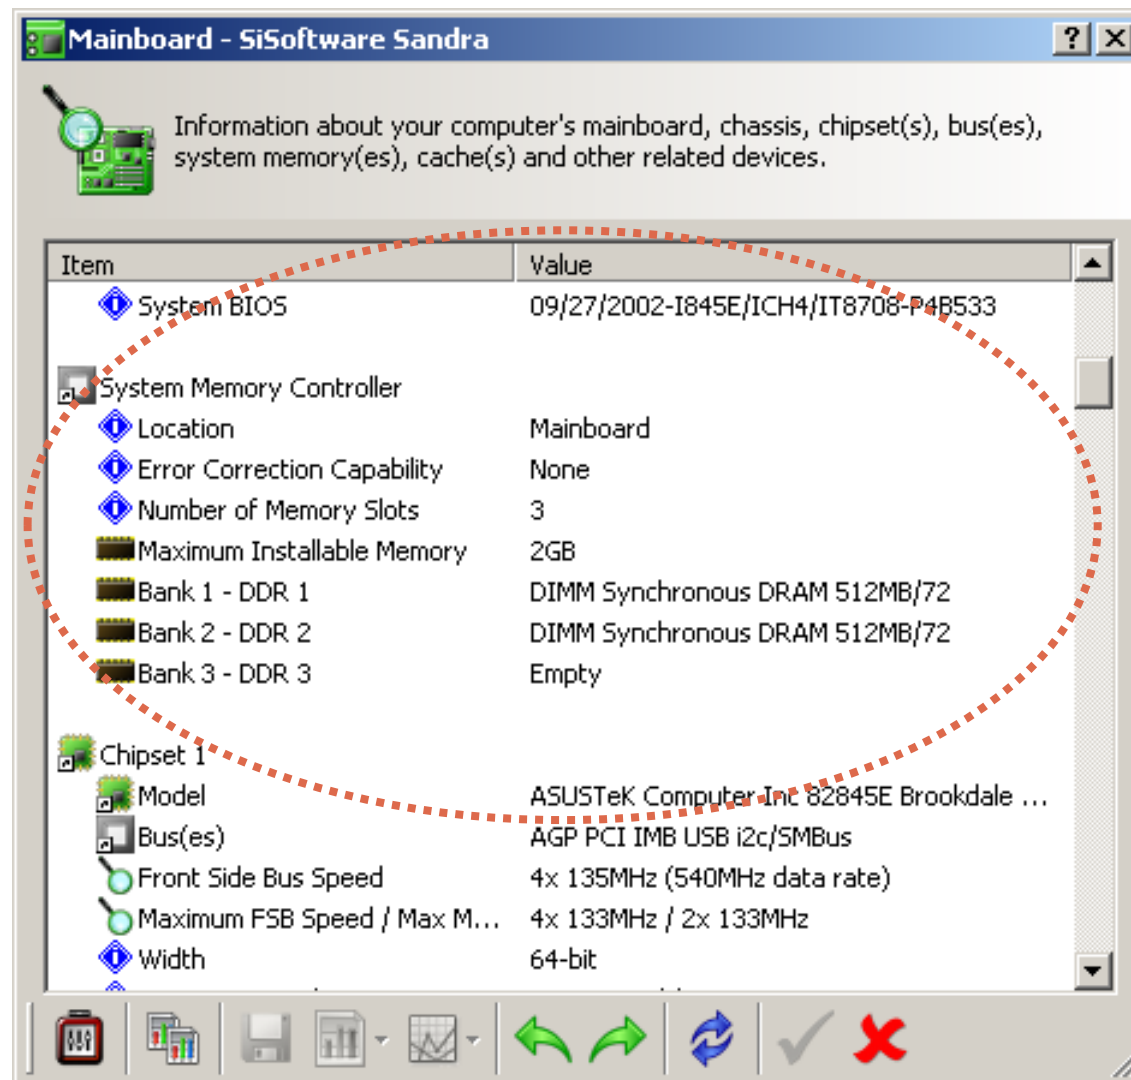
Memory controller: address handling

- Address structure:
 - n : number of address bits
 - Depends on the particular CPU
 - b : bank address bits
 - $b = \log_2 (\text{banks})$
 - r : row address bits
 - $r = \log_2 (\text{rows_per_chip})$
 - c : column address
 - $c = \log_2 (\text{columns_per_chip})$
 - r' : address of row of chips in the module
 - $r' = \log_2 (\text{chip_rows_per_module})$
 - p : number of word within accessed block
 - Depends on the pre-configured block (burst) size

Initialization of the memory system

- The memory controller is part of the system controller (chipset)
 - The chipset is roughly as complex as a CPU and includes many critical functions
- At startup, the controller
 - checks all the connected modules
 - obtains their characteristics (via the SPD chip)
 - capacity, matrix geometry, timing restrictions, ...
 - Initializes the chips in the modules
- The controller also configures
 - the memory map
 - clock frequency
 - proper timing for accessing the modules
 - the frequency of AUTOREFRESH commands

Memory controller info



- Three memory sockets
- Two of them used by modules of identical capacity