



# Unit 3:

## Database Management Systems (DBMS)

# Unit 3. Database Management Systems

---

## **1. The ANSI/SPARC Architecture**

- 1.1. Schemas
- 1.2. DBMS fundamentals
- 1.3. Data independence

## 2. Transactions, Integrity, and Concurrency

- 2.1. Transactions
- 2.2. Semantic integrity
- 2.3. Concurrent access control

## 3. Recovery and Security

- 3.1. DB Recovery
- 3.2. Security

## 1.1. Schemas

# Original proposal

---

Proposal of a DBMS architecture by the working group ANSI/SPARC (1977). They proposed the database definition with 3 levels of abstraction:

- **Internal level** → Internal schema  
Description of the DB in terms of its physical representation
- **Conceptual level** → Conceptual schema  
Description of the DB independently of the DBMS. It is usually a graphical representation.
- **External level** → External schemas  
Description of the different users' partial views

# Refined proposal

---

Since there was no generalized conceptual model for different kinds of DBMS (it is difficult to obtain the physical data structures from a conceptual graphical representation) the “logical” level was added:

- **Internal (physical) level** → Internal schema (*not in this course*)  
DB description in terms of its physical representation. Describe how the database is store in secondary memory.
- **Logical level** → Logical schema (Unit 1)  
DB description in terms of the DBMS data model. It does not include any details of the physical representation.
- **Conceptual level** → Conceptual schema (Unit 4)  
Description of the information system from the organizational point of view. It is independent of the DBMS. We will use a UML class diagram.
- **External level** → External schemas (*authorizations and views*)  
Description of the partial views which the different users have on the DB.

# Example: Logical schema

**Department** (cod\_dep: char(4), nombre: char(50), teléfono: char(8), director: char(9))

**PK**:{cod\_dep}

**NNV**:{nombre}

**FK**:{director} -> Lecturer(dni) On delete set nulls. On update cascade

**Subject** (cod\_asg: char(5), nombre: char(50), semestre: char(2), cod\_dep: char(4),  
teoría: real, prácticas: real)

**PK**:{cod\_asg}

**NNV**:{nombre, semestre, cod\_dep, teoría, prácticas}

**Uni**:{nombre}

**FK**:{cod\_dep} -> Department(cod\_dep)

On delete restrict. On update cascade

**IC**<sub>1</sub>:(teoría <= prácticas)

**IC**<sub>2</sub>:(semestre IN {'1A','1B','2A','2B','3A','3B','4A','4B'})

**Lecturer** (dni: char(9), nombre: char(80), teléfono: char(8), cod\_dep: char(4),  
provincia: char(25), edad: entero)

**PK** :{dni}

**NNV** :{nombre, cod\_dep}

**FK** :{cod\_dep} -> Department(cod\_dep)

On delete restrict. On update cascade

**Teaching** (dni: char(9), cod\_asg: char(5), gteo: entero, gpra: entero)

**PK** :{dni,cod\_asg}

**NNV** :{gteo,gpra}

**FK** :{dni} -> Lecturer(dni) On delete cascade. On update cascade

**FK** :{cod\_asg} -> Subject(cod\_asg) On delete restrict. On update cascade

**General constraint:** **GC1**: "All teacher must lecture at least one subject".

# Example: Internal schema

---

Depends on the DBMS.

***Subject:***

Hash file by cod\_dep

B+ index over (semestre + cod\_dep)

***Lecturer:***

Hash file by nombre

***Department:***

Hash file by cod\_dep

B+ index over nombre

***Teaching:***

Disordered file

# Example: External schema

---

External schema for the **Maths Department**

```
CREATE VIEW Maths-Lecturer AS
    SELECT dni, nombre, teléfono, categoría, edad
    FROM Lecturer
    WHERE dptocod_dep = 'DMA';

CREATE VIEW Maths-Subject AS
    SELECT cod_asg, nombre, semestre, teoría, prácticas
    FROM Subject
    WHERE cod_dep = 'DMA';

CREATE VIEW Maths-Teaching AS
    SELECT T.dni, T.cod_asg, T.gteo, T.gpra
    FROM Lecturer L, Teaching T, Subject S
    WHERE L.cod_dep = 'DMA'
        AND S.cod_dep = 'DMA'
        AND L.dni = T.dni
        AND S.cod_asg = T.cod_asg;
```



---

A DBMS that supports the 3-level architecture must:

- Allow the **definition** of the different **schemas** for the database (except the conceptual schema),
- Establish the **correspondence** between schemas,
- **Isolate the schemas**: changes in one schema should not affect neither the schemas at upper levels nor the application programs.




**DATA INDEPENDENCE**

## 1.2. DBMS Fundamentals

# Functions of a DBMS

---

**DBMS:** Software which allows the creation and manipulation of databases (DB).



A DBMS must maintain the independence, integrity and security of data.

# Functions of a DBMS

---

## Objectives of DB techniques

- Unified and independent **data description**
- **Application** independence
- Partial **view** definition

## DBMS Functions

**Data definition** at several **levels**

- Logical schema
- Internal schema
- External schema

## DMBS Components

Schema definition languages and their associated translators

# Functions of a DBMS

---

## Objectives of DB techniques

- Information **management**

## DBMS Functions

### Data **manipulation**

- Query
- Update

Management and **administration** of the database

## DMBS Components

Manipulation languages and their associated translators

Tools for:

- Restructuring
- Simulation
- Statistics
- Printing and reporting

# Functions of a DBMS

---

## Objectives of DB techniques

- Data **integrity** and **security**

## DBMS Functions

Control of:

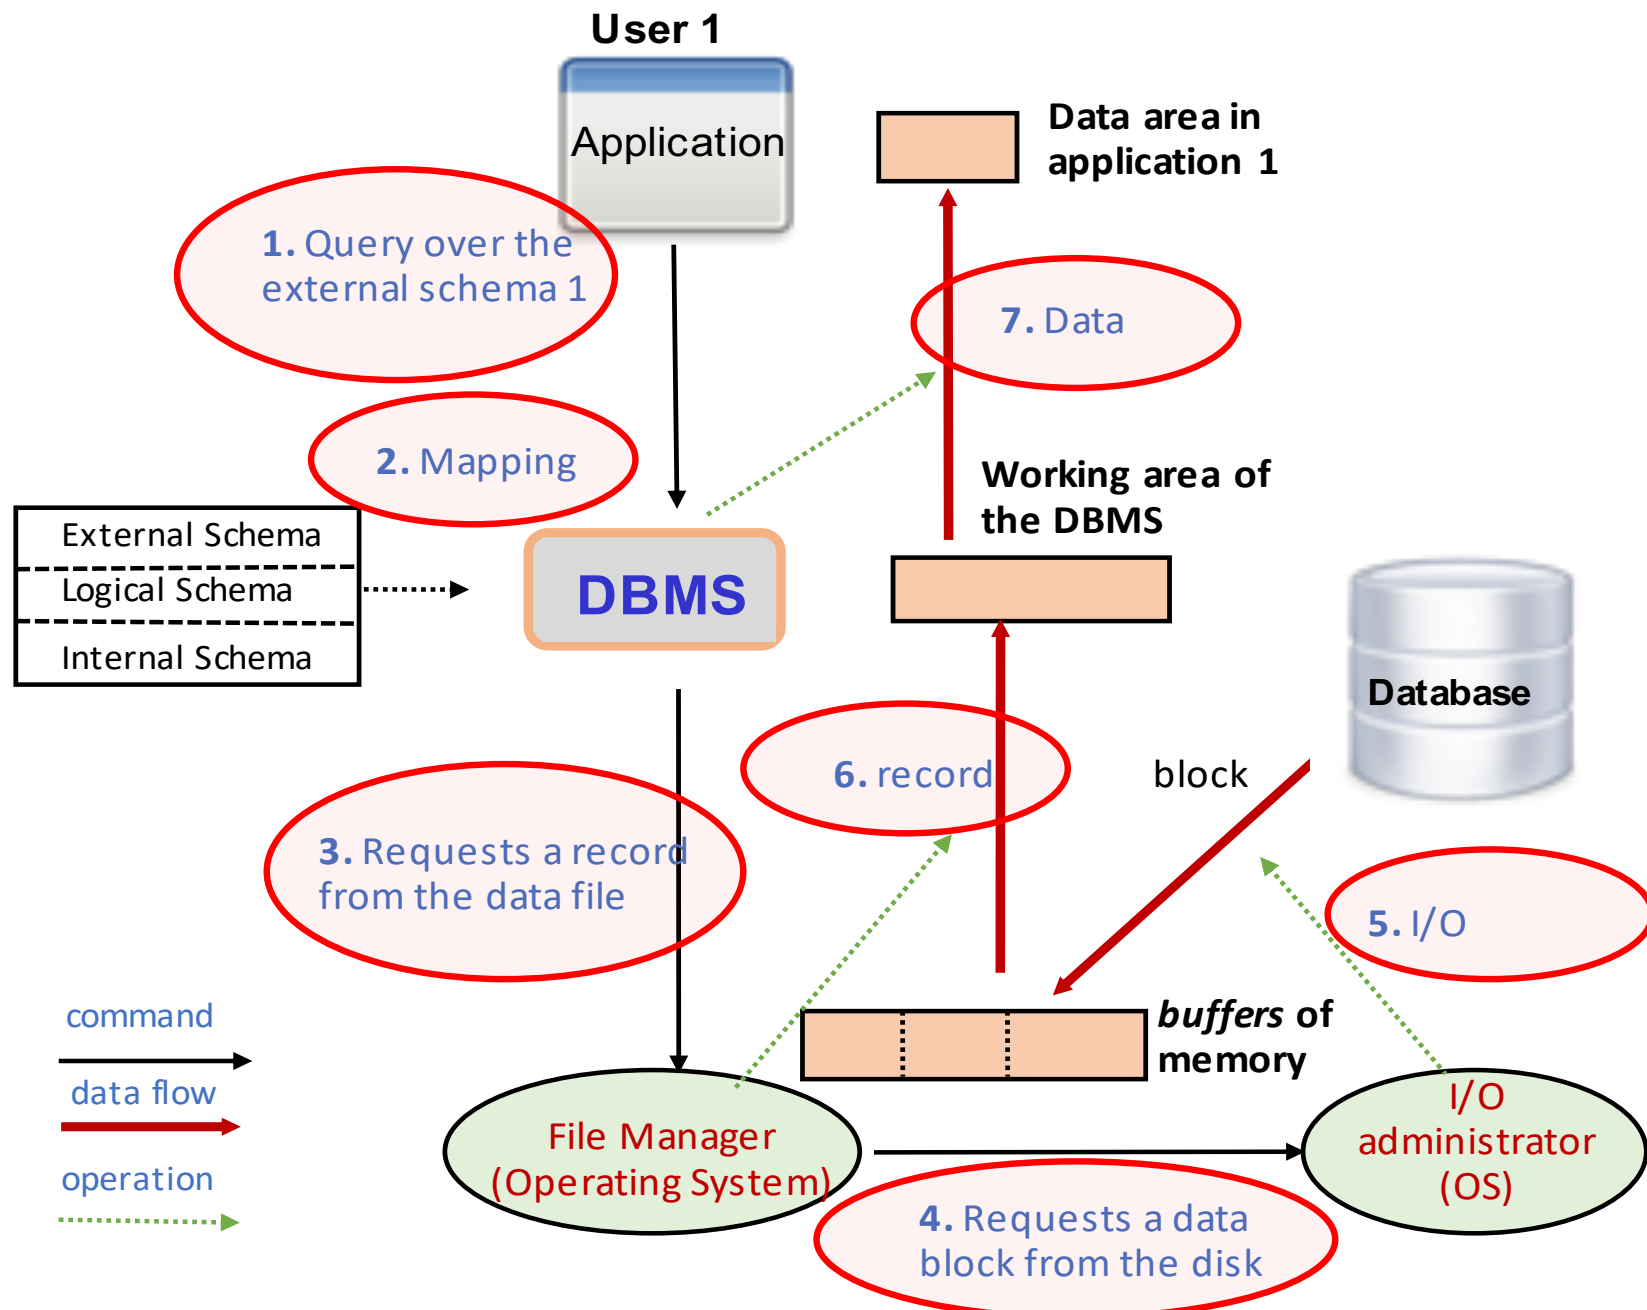
- Semantic integrity
- Concurrent access
- Recovery in case of failure
- Security (privacy)

## DMBS Components

Tools for:

- Integrity control
- Reconstruction
- Security control

# Accessing the data



## 1.3. Data Independence



# Data independence

---

Property which ensures that the application programs are independent of

- the changes which are performed on data which they do not used
- or
- the physical representation details of the accessed data

# Logical independence

---

*Logical independence* between the logical schema and the external schemas:

The external schemas and the application programs cannot be affected by the modifications in the logical schema of data which are not used by these programs

## EXAMPLE:

If we add new attributes to the “*Department*” table, such as the date in which the department was created, the building,... the external schema of the “*Maths-department*” does not need to be modified

# Physical independence

---

*Physical independence* between the internal schema and the logical schema:

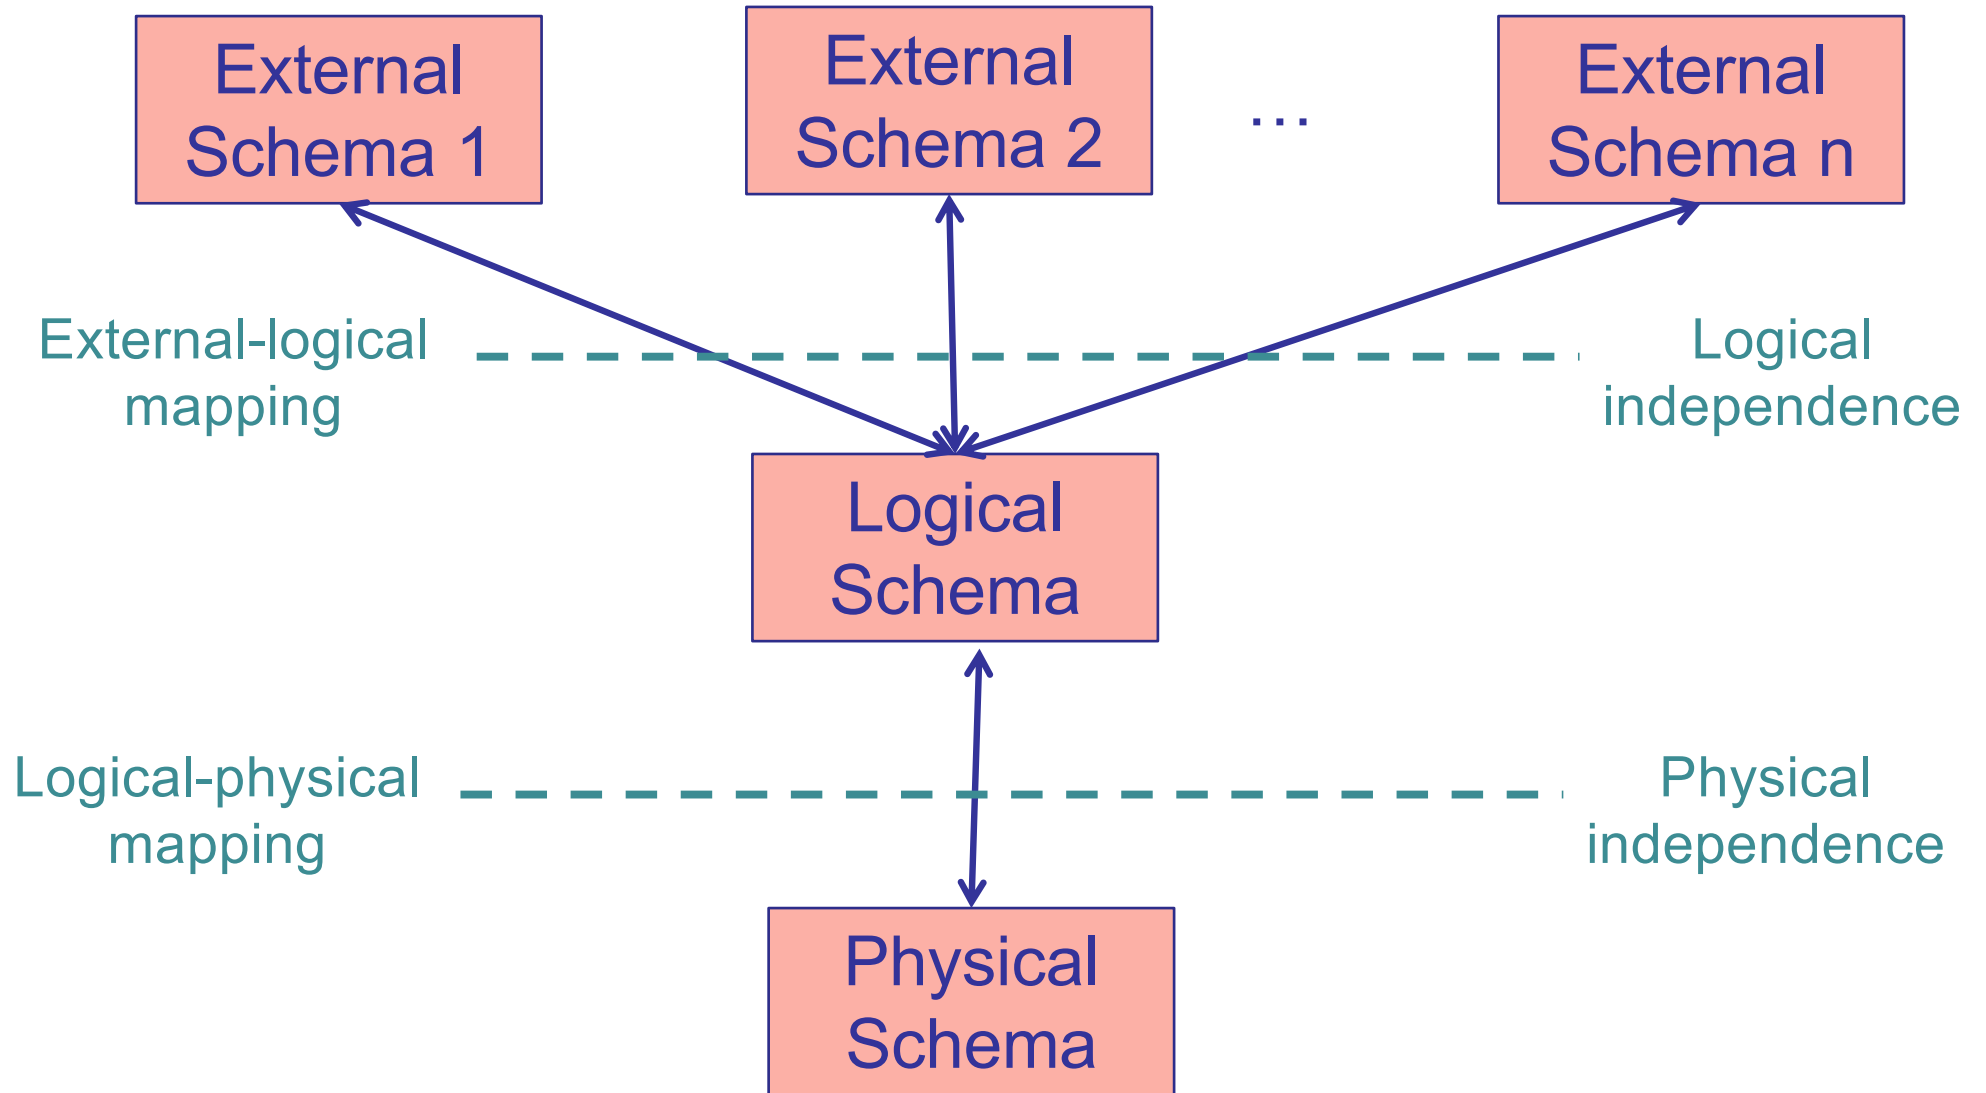
The logical schema cannot be affected by changes in the internal schema which refer to the implementation of the data structures, access modes, page size, search path, etc.

## EXAMPLE:

If the data structures used in the implementation of the “*Subject*” table are changed, the logical schema does not need to be modified

# Data independence

---

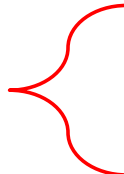


# Binding

---

## Binding (ligadura):

Transformation of the external schema into the internal schema.

**Types** 

- Logical binding (steps 2 and 7).
- Physical binding (steps 3 and 6).

When the binding is performed, independence disappears



It is important to determine the binding moment

Currently, most DBMS do the binding for each query.

*Physical independence is found in most DBMS, while logical independence is more difficult to find.*

# Unit 3. Database Management Systems

---

## 1. The ANSI/SPARC Architecture

1.1. Schemas

1.2. DBMS fundamentals

1.3. Data independence

## **2. Transactions, Integrity, and Concurrency**

2.1. Transactions

2.2. Semantic integrity

2.3. Concurrent access control

## 3. Recovery and Security

3.1. DB Recovery

3.2. Security

## 2. Transactions, Integrity, and Concurrency

---

Objective of DB technology



**Information quality**

*“Data must be structured in such a way to adequately **reflect** the **objects**, **relations**, and **constraints** which exist in the part of the real world modeled by the database model.”*

When reality changes → User **updates** the database

The information contained in the DB must preserve the schema definition.

## 2. Transactions, Integrity, and Concurrency

---

**Information quality** (integrity perspective) means that:

- The DBMS must ensure that the data are **correctly stored**
- The DBMS must ensure that **user updates** over the DB are correctly executed and become **permanent**.



## 2. Transactions, Integrity, and Concurrency

---

DBMS Tools oriented towards integrity:

- Check (when an update is performed) the **integrity constraints** defined in the schema.
- Control the correct execution of the **updates** in a concurrent environment.
- Recover (**reconstruct**) the DB in case of loss or accident

## 2.1. Transactions

## 2.1. Transactions

---

The operations in a DB are organized in transactions.

Transaction:

Sequence of access operations to the DB which constitute a **logical execution unit**.

# Example

---

**Emp** (id, name, address, dept)

PK: {id}

FK: {dept} → Dep(code)

**Dep** (code, name, location)

PK: {code}

IC<sub>1</sub>: All departments have at least one employee

Insert a new department:

<“d2”, “Human Resources”, “2nd floor”>

whose first employee is the id 20

# Example

---

1st  
Idea

1) Insert in *Dep*:

<d2, "Human Resources", "2nd floor">

**ERROR: IC<sub>1</sub> is violated**

2) Modification of *Emp* on the tuple with *id* 20

2nd  
Idea

1) Modification of *Emp* on the tuple with *id* 20

**ERROR: the FK over dept in Emp is violated**

2) Insertion in *Dep*:

<d2, "Human Resources", "2nd floor">

# Defining transactions

---

Actions which change transactions states:

**Begin:**

Indicates the **beginning** of the execution of the transaction

**Cancellation (user rollback):**

The user aborts the transaction.

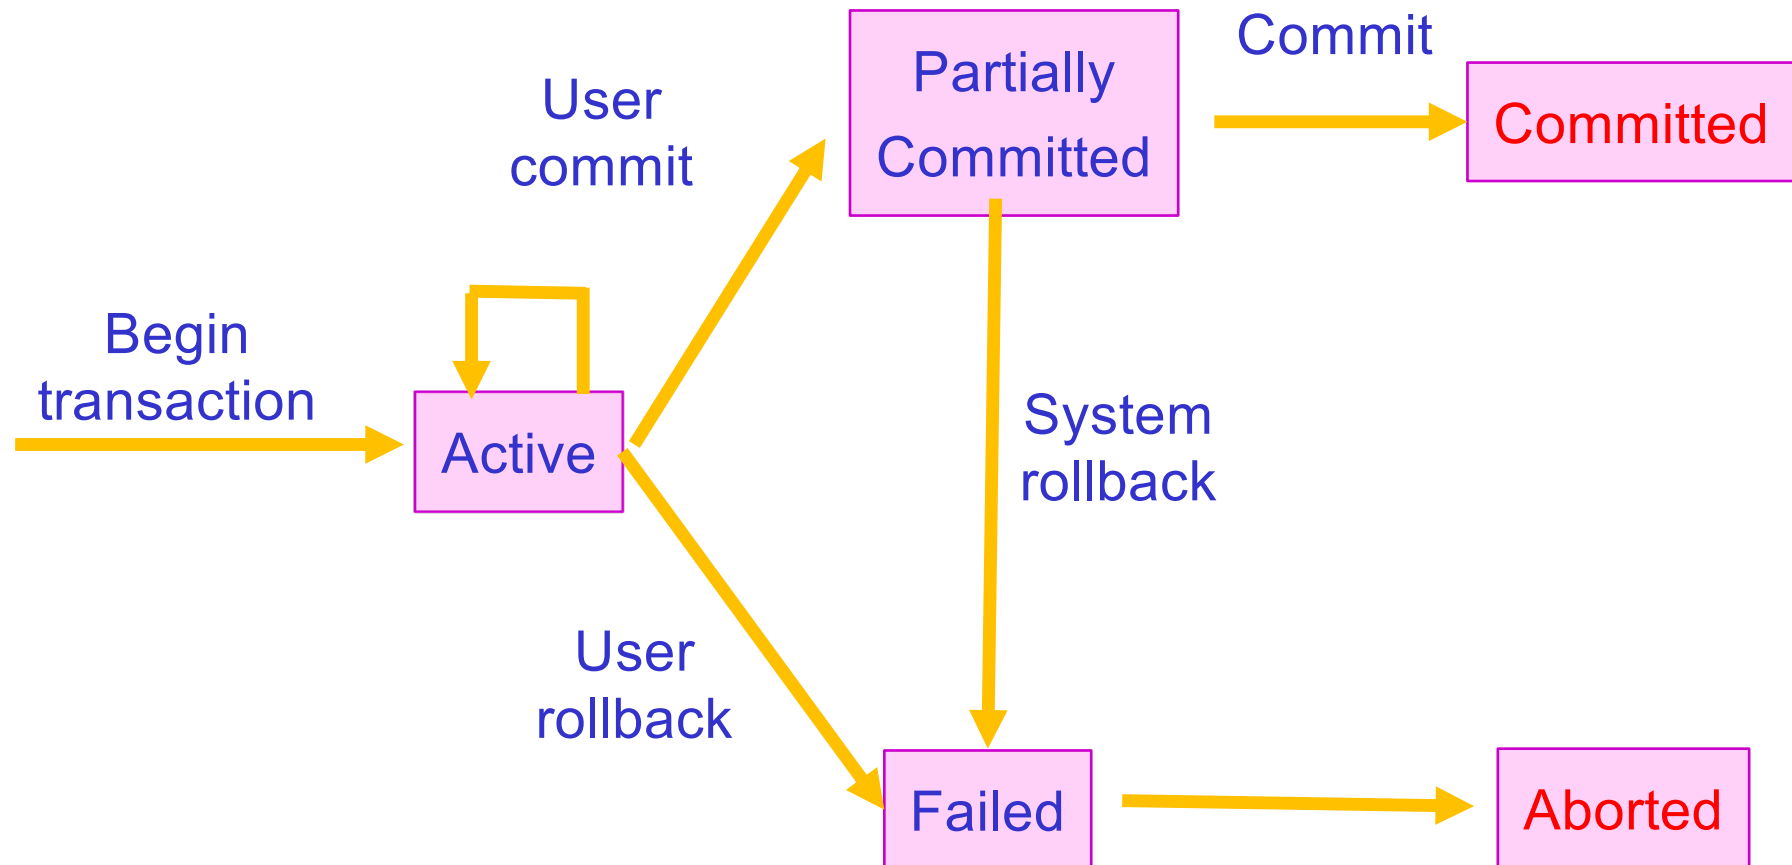
**Confirmation (*user commit*):**

The user considers the transaction as ended. Then the DBMS performs some checking to determine how the transaction will end:

- **Success (system commit):**  
Indicates the **success** of the transaction, making the DBMS store the changes performed on the DB.
- **Failure (system rollback):**  
Indicates the **failure** of the transaction, or that the transaction hasn't passed the checking. The DBMS undoes all the possible changes performed by the transaction.

# States of a transaction

---



# Properties of Transactions (ACID)

---



## Atomicity:

A transaction is an indivisible unit that is either **performed** in its **entirety** or is not performed at all (“All or nothing”).

## Consistency:

The transaction must transform the DB from one consistent state **to another consistent state** (all integrity constraints must be met)

## Isolation:

Concurrent transactions execute independently: All the partial effects of **incomplete transactions should not be visible** to other transactions

## Durability:

The effects of a **successfully** completed (committed) transaction are **permanently recorded** in the DB and must not be lost because of a subsequent system or other transaction failure



## 2.2. Semantic Integrity

## 2.2. Semantic Integrity

---

### Integrity constraint:

*Property of the real world which is modelled by the DB*

- Constraints are defined in the **logical schema** and the DBMS must ensure that they are met.
- Checking is **performed** whenever the **DB changes** (when any updating operation is executed)
- Constraints **not included in the DB** schema must be maintained by the application programs

*This situation is, in general, **inappropriate** if the constraints are common to more than one application, since the responsibility to check them is dispersed*

# Types of integrity constraints

---

**Static:** They must be met in each state of the DB (they can be represented by logical expressions)

**Example:**

The *lab credits* in a *subject* cannot be greater than the *lectures credits*.

**Dynamic (Transition):** They must be met regarding two consecutive states.

They are not usually implemented by commercial DBMS

**Example:**

The *credits* of a *subject* cannot decrease.

# Expressing static constraints in SQL

---

- Constraints over possible **data values**. e.g. Domains
- Constraints over **attributes**. e.g. NNV.
- Constraints over **relations**. e.g. PK, FK
- **General constraints** over the DB. e.g: “*All subject must be lectured by at least one teacher.*”

**When** are checked:

- After every command (**IMMEDIATE**)
- At the end of the transaction(**DEFERRED**)

# Expressing transition constraints in SQL

---

## Triggers (*“disparadores”*)

- Using triggers, the designer can program the system response when some events are produced.
- This allows to incorporate complex constraints into the DB.

### A trigger includes:

1. **Events**: Operations over the DB which trigger it
2. **Conditions** to determine if the actions must be executed or not.
3. **Actions** to be executed when an event happen and the conditions are met. They are usually written in a data-oriented high level programming language, that can include SQL commands.

# Expressing transition constraints in SQL

---

## Example:

A trigger to implement the integrity constraint: *“A lecturer can only teach a subject assigned to his/her department”*

1. **Events:** INSERTion of a tuple in the “Subject” table
2. **Conditions:** The lecturer and the subject are not in the same department.
3. **Actions:** Reject the operation (the insertion).

## **2.3. Concurrent access control**

## 2.3. Concurrent access control

---

In order to keep the integrity of the database, the DBMS must control concurrent access to the database:

To avoid that the results of the execution of **several processes** (users or programs) lead to **incorrect, incoherent** or **lost results** because of the simultaneous execution of other program accessing the same data



# Basic operations

---

Basic operations in a transaction which are relevant to the DBMS:

**read(X):**

Reading or access to a piece of data X in the DB over the program variable with the same name.

**write(X):**

Update (insertion, deletion, or modification) of a piece of data X in the DB by using the program variable with the same name

# Reading steps

---

read(X):

1. Seek the address of the block which contains the datum X
2. Copy the block to a buffer into main memory
3. Copy the datum X from the buffer to the program variable X

# Writing steps

---

write(X):

- If not read before {
1. Find the **address of the block** containing the datum X
  2. **Copy** the block into a database **buffer** in main memory
  3. **Copy** the **datum** X from the program variable into the database **buffer**
  4. **Write the updated block** from the database buffer to the disk

# Possible problems

---

The DBMS must control the concurrent access by the applications.

**Problems** due to interference of concurrent accesses:

- a) **Loss of updates**. An apparently successfully completed update operation by one user can be overridden by another user.
- b) **Inconsistent** information corresponding to several valid database states. One transaction reads several values but a second transaction updates some of them during the execution of the first.
- c) Access to updated data (but **still not confirmed**) that can still be cancelled. One transaction is allowed to use the intermediate results of another transaction before it has committed (*dirty read*).

# A. The lost update problem

Time	Program 1	Program 2
t1	read(11548, teoría) teoría=4,5	
t2		read(11548, teoría) teoría=4,5
t3	teoría←teoría+1,5 teoría=6	
t4		teoría←teoría+2 teoría=6,5
t5	write(11548, teoría )	
t6		write(11548, teoría )

Two programs reading and updating the theory credits of subject 11548:

- P1 adds 1,5 credits
- P2 adds 2 credits

Subject

cod_asg	nombre	semestre	cod_dep	teoría	prácticas
11545	Análisis Matemático	1A	DMA	4,5	1,5
11546	Álgebra	1B	DMA	4,5	1,5
11547	Matemática Discreta	1A	DMA	4,5	1,5
11548	Bases de Datos y Sistemas de Información	3A	DSIC	6,5	1,5

$$4,5 + 1,5 + 2 = 7$$

## B. The inconsistent analysis problem

*Program 1: List for each lecturer his/her amount of credits*

Teaching				Time	Program 1	Program 2
dni	cod_asg	gteo	gpra	t1	Calculate credits for lecturer 111 Credits 111= 9 (1×4,5 + 3×1,5)	Change a theory group of subject 11545 from lecturer 564 to lecturer 111
111	11545	2	3	t2	Calculate credits for lecturer 123 Credits 123= 9 (0×4,5 + 2×1,5 + 1×4,5 + 1× 1,5)	
123	11545	0	2	t3		
123	11547	1	1	t4	Calculate credits for lecturer 453 Credits 453= 0	
564	11545	1	2	t5	Calculate credits for lecturer 564 Credits 564= 7,5 (1×4,5 + 2×1,5)	
Subject				Result:		
cod_asg	...	teoría	prácticas			
11545	...	4,5	1,5			
11546	...	4,5	1,5			
11547	...	4,5	1,5			
11548	...	4,5	1,5			

DNI	Credits
111	9 credits
123	9 credits
453	0 credits
564	7,5 credits

## C. The uncommitted dependency problem

Subject

cod_asg	...	teoría	prácticas
11545	...	4,5	1,5
11546	...	4,5	1,5
11547	...	4,5	1,5
11548	...	6	1,5

Time	Program 1	Program 2
t1	read(11548, teoría)	
t2	teoría ← teoría + 1,5	
t3	write(11548, teoría)	
t4		read(11548, teoría)
		teoría = 6
t5		<i>Use this value (teoría = 6) in its instructions</i>
t6		confirmation
t7	cancellation	

# Techniques

---

## Reserving some data occurrences (locks)

- In examples **a)** and **c)** P1 must lock a record.
- In examples **b)** all the table must be locked.
- Need for controlling deadlocks

**Other solutions** (for the example c): Cascade cancellation or transaction isolation



# Unit 3. Database Management Systems

---

## 1. The ANSI/SPARC Architecture

- 1.1. Schemas
- 1.2. DBMS fundamentals
- 1.3. Data independence

## 2. Transactions, Integrity, and Concurrency

- 2.1. Transactions
- 2.2. Semantic integrity
- 2.3. Concurrent access control

## **3. Recovery and Security**

- 3.1. DB Recovery
- 3.2. Security

# Recovery and Security

---

A database must guarantee:

- **Recovery:**

*(Part of integrity, but not from the point of view of consistency.  
Recovery is focus on durability and persistence)*

A database must always be recovered from any type of failure.

- **Security:**

A database cannot allow non-authorized access.

## 3.1. DB Recovery

# DB recovery

---


The transaction properties of **atomicity** and **durability** force a DBMS to ensure that:

- If **confirmed**, the changes performed are recorded in the DB to make them persistent.
- If **cancelled**, the changes performed over the DB are **undone**.

# Database recovery: Example

---

## Backups are not sufficient



20-11-2019: Update on accounts

- Transaction #1: Update ....
- Transaction #2: Update ....
- Transaction #3: Update ....
- ...
- Transaction #51: **System failure !**

## Recovery Procedure:

- Replace the file “Accounts” with its backup

## Negative effect:

- The updates of 50 transactions are lost

# DB Recovery

---

**Backups** alone are **not** the **solution** to the recovery problem.

- The increase of the backup frequency is not a feasible solution.

DB technology provides much more efficient and robust techniques for DB recovery.

Lost of confirmed data is inadmissible  
with current technology.

# DB Recovery

---

**Backups** alone are **not** the **solution** to the recovery problem.

- The increase of the backup frequency is not a feasible solution.

DB technology provides much more efficient and robust techniques for DB recovery.

Lost of confirmed data is inadmissible  
with current technology.

# Causes of transaction failure

---

## 1. **Local to the transaction** (normal system operation)

- **Transaction error** (incorrect DB access, failed calculations, etc.)
- **Exceptions** (integrity violation, security problems, etc.)
- **Concurrency control** (locked state between two transactions)
- **Human decision** (inside a program or explicitly).

## 2. **Extern to the transaction** (system error)

- A. System failures with **loss of main memory**.
- B. Failures in the storage system with **loss of secondary memory** (disk failure, human errors, virus infection,...)



# A. Failures of system main memory

---

- The changes performed by a transaction are located in memory buffers (main memory).
- When the transaction is **confirmed** its changes must be recorded in secondary memory.
- If a failure with loss of main memory occurs between the **transaction confirmation** and **flushing the buffers** to secondary memory, the blocks in the buffers will be lost.

## B. Failures of secondary memory

---

- The changes performed by a **confirmed** transaction are recorded into the DB
- If there is a **failure in secondary memory**, the changes will be lost.

### **3.1.1 Recovery from failures of system main memory**

# Recovery from failures of system main memory

---

## What to do:

- Recover confirmed transactions which have not been recorded.
- **Cancel** transactions which have **failed**.

## Who:

Recovery module

## How:

Most used technique: Use a *journal file* (or *log*, “*fichero diario*”).

# Transaction implementation

---

Two ways of implementing transactions:

- Transactions with Immediate Update  
Updates have an immediate effect on secondary memory.  
In case of cancellation, they have to be undone.
- Transactions with Deferred Update  
Updates only have immediate effect on main memory.  
The updates will be transferred to secondary memory when confirmed.
- Both need a log file

# Log (journal) file

---

Activities and events are recorded in the log file:

- **Record** the update operations performed by existing transactions during a period of time (e.g. one day)
- The log (journal) file is stored on **disk** to avoid loss after a system failure.
- It is **dumped periodically** into a massive storage unit (magnetic tape, optical disk,...).

# Types of entries in a log file

---

[**start**, T]: A transaction with identifier T has been started.

[**write**, T, X, value\_before, value\_after]: The T transaction has performed an update instruction on data X.

[**read**, T, X]: The T transaction has read data X.

[**confirm**, T]: The T transaction has been confirmed.

[**cancel**, T]: The T transaction has been cancelled.

---

## Problems:

- Size of log file can increase very quickly.
- Recovery in case of failure is very expensive (many instructions have to be redone).

## Solution:

**checkpoints** (*“puntos de control o verificación”*)



---

Checkpoints → They are recorded in the log file periodically

### How it works ?

- **Suspend** the execution of transactions temporally.
- **Record** a checkpoint in the log file.
- **Record** all updates performed by confirmed transactions (copy all main memory **buffers to disk**).
- **Resume** the execution of the suspended transactions..

## A. IMMEDIATE DATABASE UPDATES

---

Updates have an **immediate effect on secondary memory**. In case of cancellation, they have to be undone

Failure of a transaction  $T \rightarrow$  **Undo** changes performed by  $T$

Update the data which has been modified by  $T$  with its original value (*value\_before*): Search for the entries in the logfile [**write**,  $T$ ,  $X$ , *value\_before*, *value\_after*]

---

## Unconfirmed transactions

$[start, T]$  in the log file without  $[confirm, T]$

→ Undo changes performed by  $T$  (previous process)

## Confirmed transactions

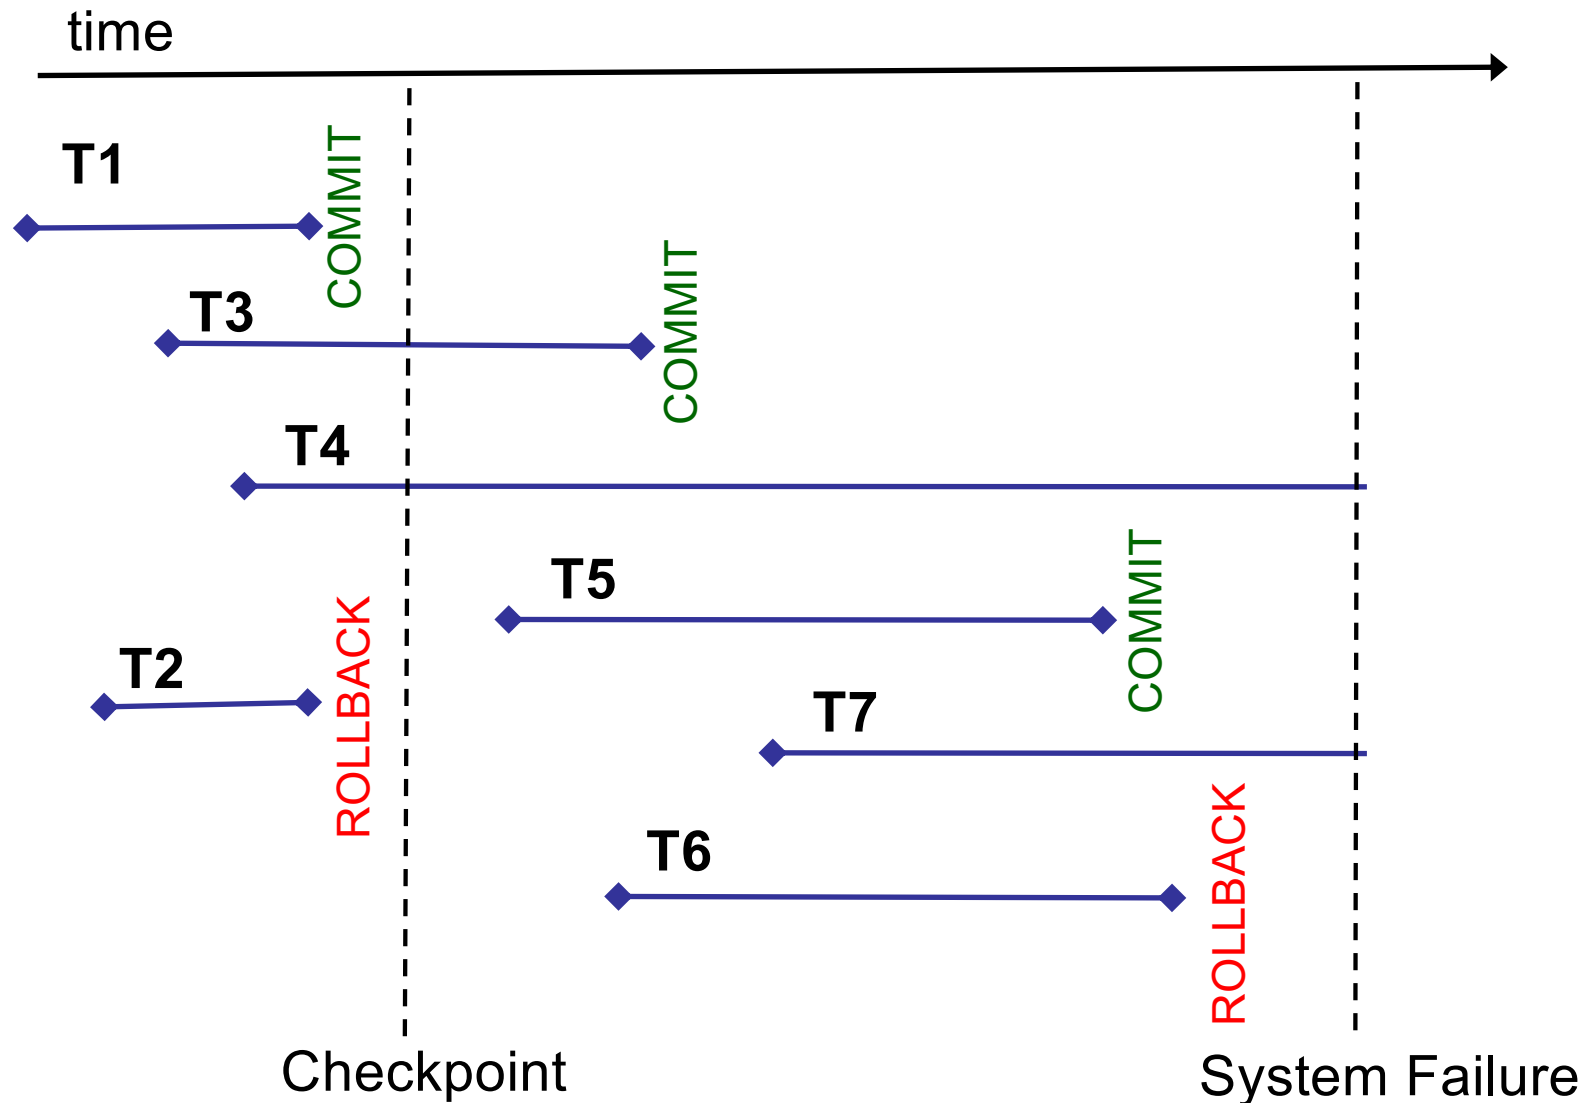
$[confirm, T]$

→ Execute (redo) them again:

$[write, T, X, value\_before, value\_after]$

# DB recovery with **immediate** Updates

Recovery the DB from the last checkpoint



# DB recovery with immediate Update

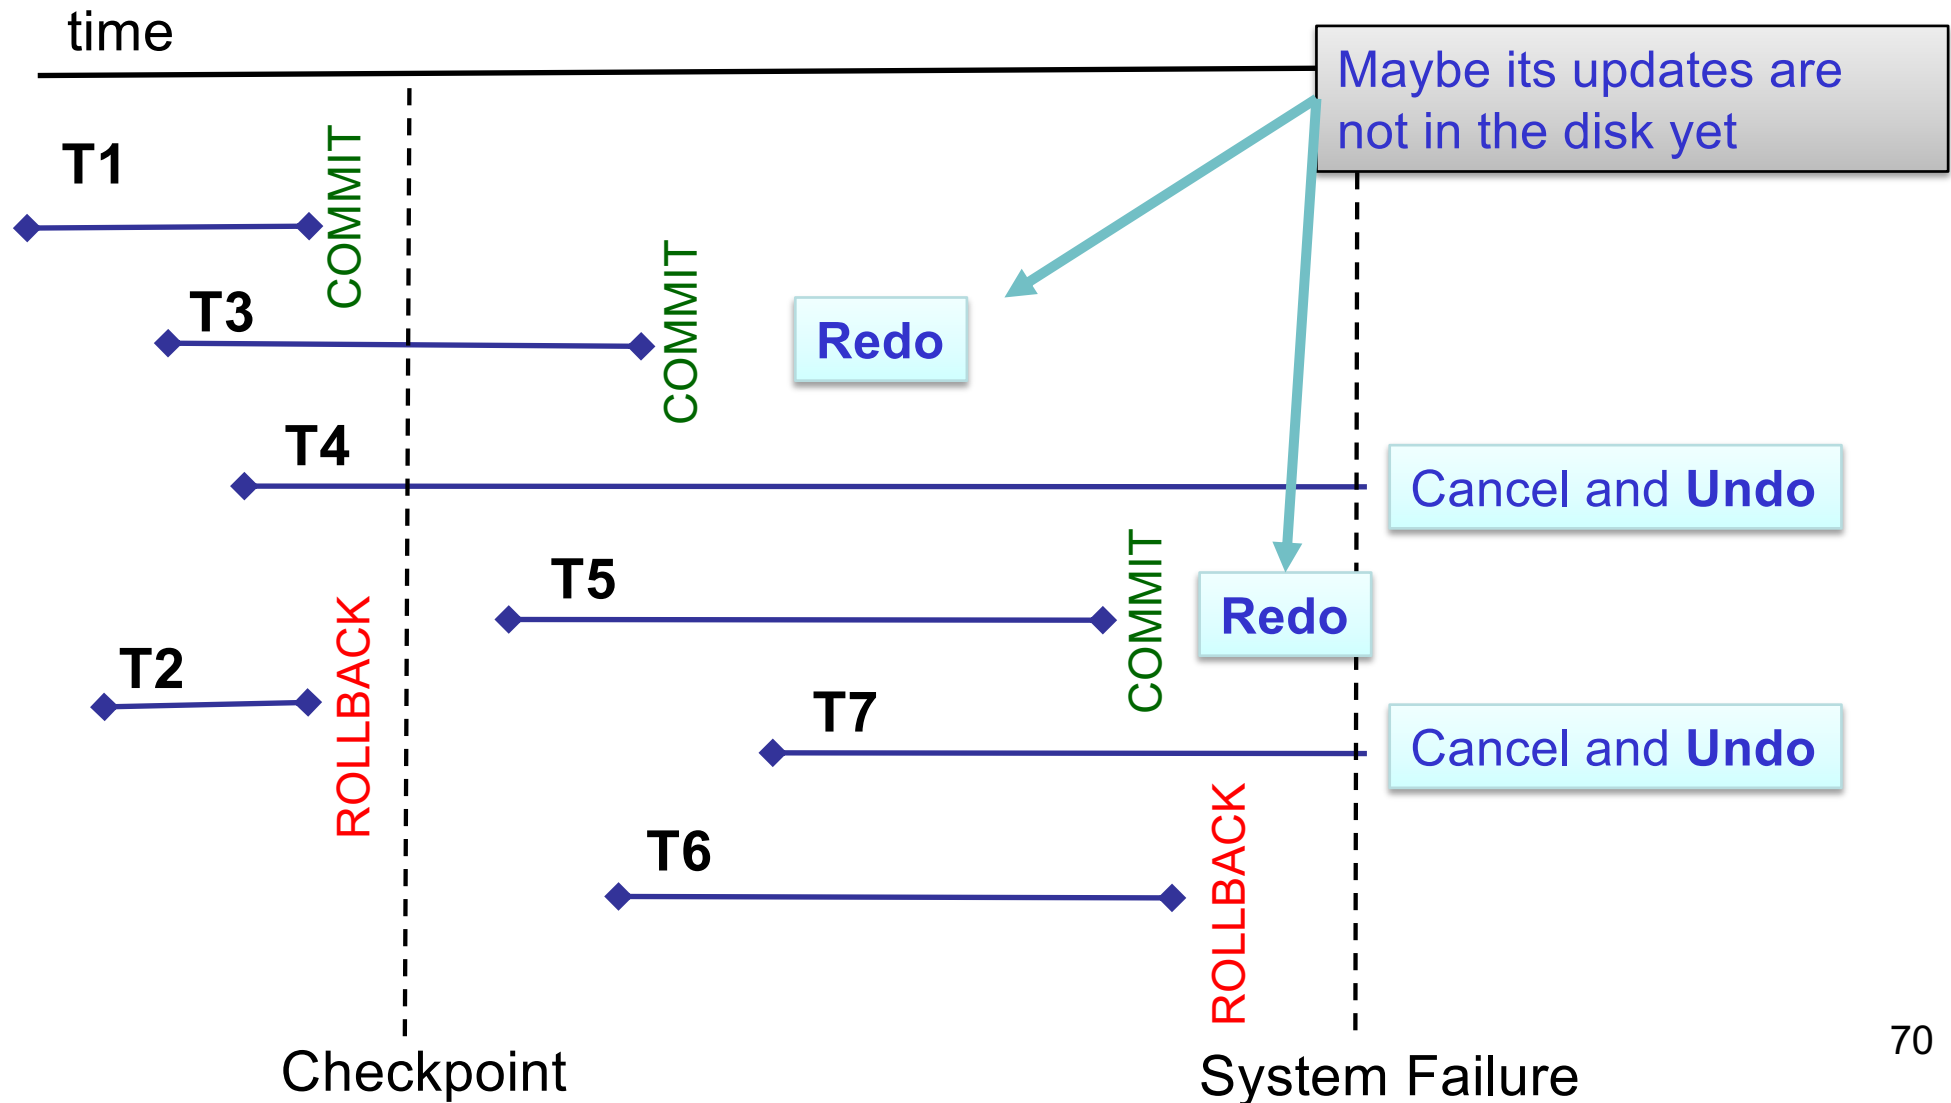
---

## Basic considerations:

- Updates performed by confirmed transaction (*a transaction commit appears in the log file*) could have not been transferred to disk because the buffer block where they are, has not been recorded yet: Redo
- Updates performed by non-confirmed transactions (there is no transaction commit in the log file) could be in disk because their main memory blocks were transferred to disk: Undo (only used in immediate updates)
- When a checkpoint is recorded, the DBMS records all the updates performed by the confirmed transaction.

# DB recovery with immediate Update

Recovery the DB from the last checkpoint



## B. DEFERRED DATABASE UPDATES

---

Updates only have immediate effect on main memory. The updates are only written to the secondary memory **until confirmed** (after the commit).

- **Unconfirmed** transactions

$[start, T]$  in the log file without  $[confirm, T]$

→ **do nothing** they are not in secondary memory

- **Confirmed** transactions

$[confirm, T]$

→ **Execute (redo)** them again:

$[write, T, X, value\_before, value\_after]$

# DB recovery with deferred Update

---

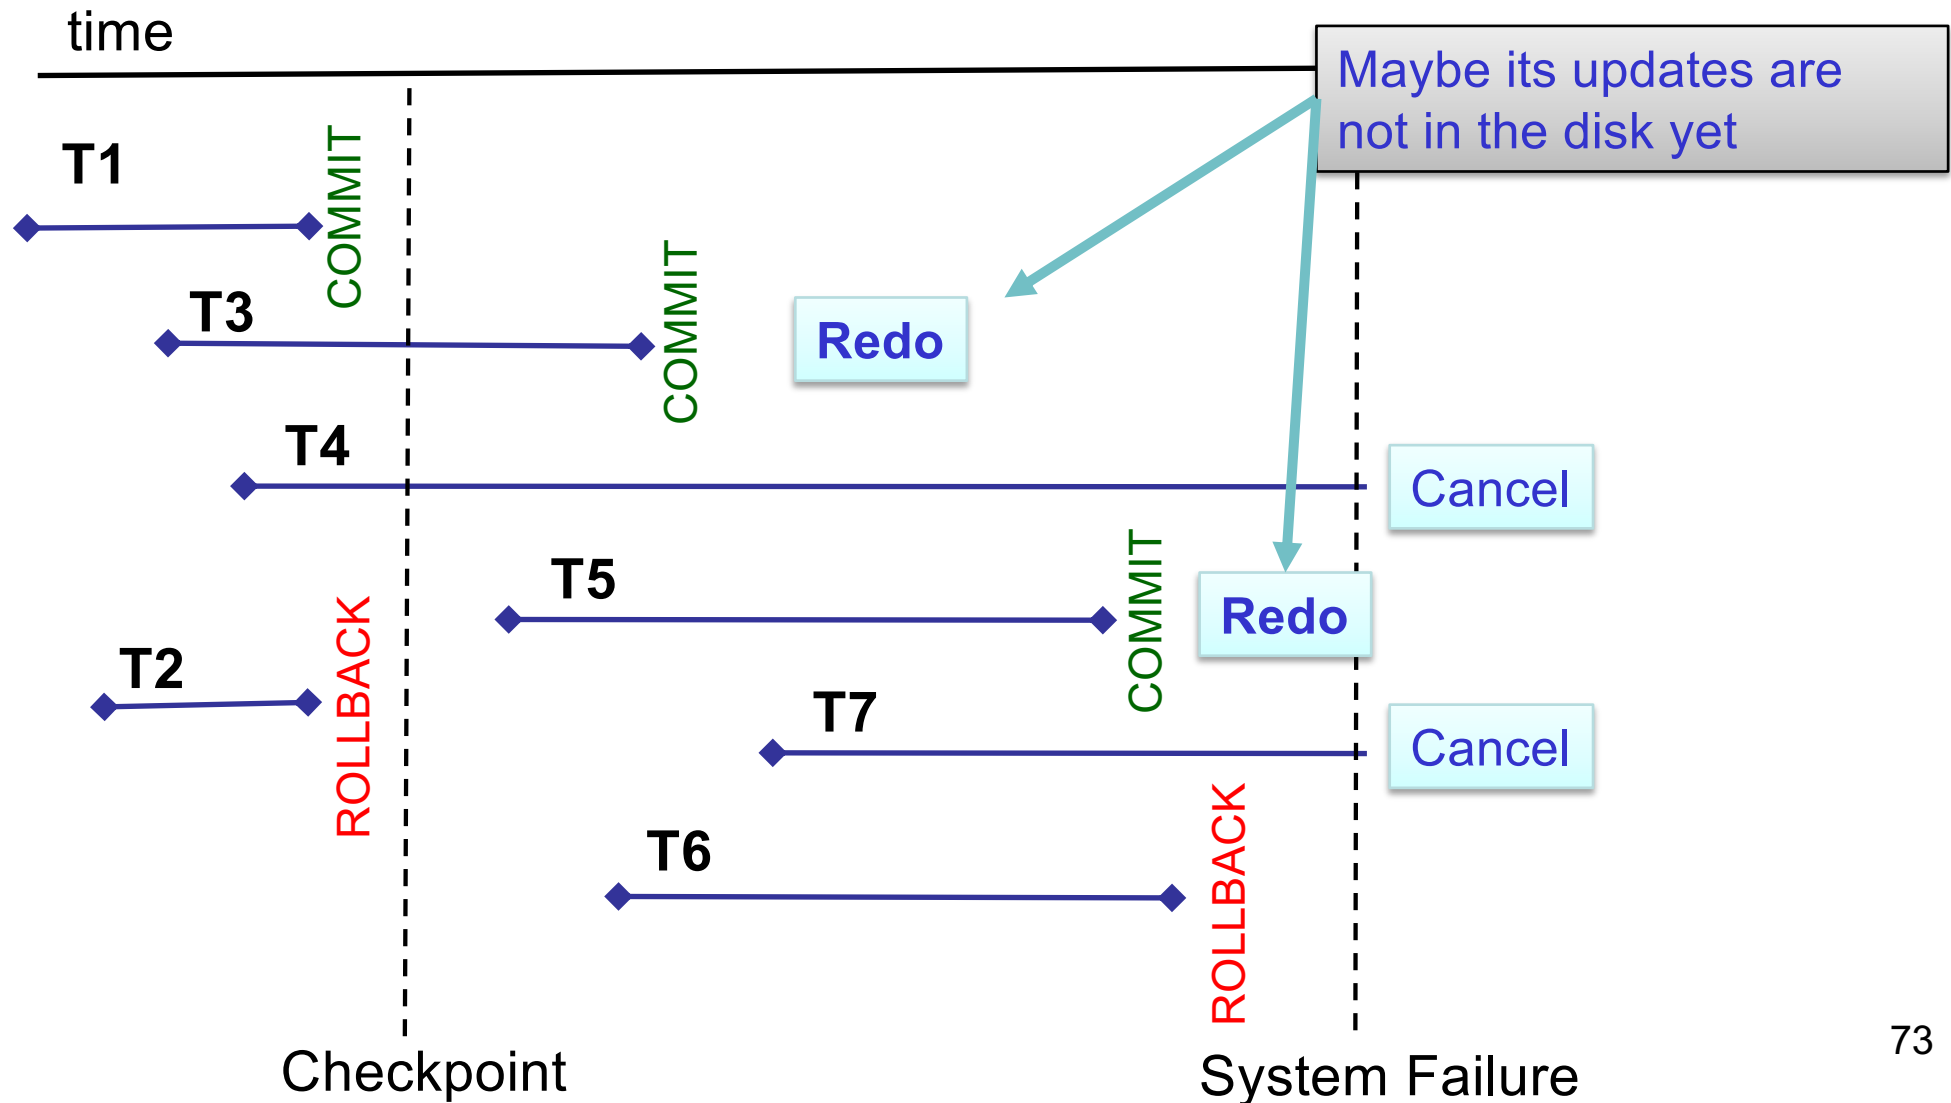
## Basic considerations:

- Updates performed by confirmed transaction (*a transaction commit appears in the log file*) could have not been transferred to disk because the buffer block where they are, has not been recorded yet: Redo
- Updates performed by non-confirmed transactions (there is no transaction commit in the log file) are not in disk : **Do nothing**
- When a checkpoint is recorded, the DBMS records all the updates performed by the confirmed transaction.



# DB recovery with deferred Update

Recovery the DB from the last checkpoint



## **3.1.2 Recovery from failures of secondary memory**

# Recovery from failures of secondary memory

---

When a failure of the **storage system** occurs, the database might be **damaged** totally or partially.

## Technique:

### **Reconstruction** of the database:

- Using the most recent **backup**
- From the backup instant, the system uses the **log file** to redo all the instructions performed by the confirmed transactions.

## 3.2. DB Security

## 3.2. Security

---

### Objective:

Information can only be accessed by the people and processes that are authorized and in the authorized way

# Techniques

---

- User identification
- Establishment of allowed accesses:


## Modes:

- **Authorization list** associated to each user containing the allowed objects and operations. (*GRANT*)
  - **Level of authorization** (less flexible). There several users groups with different authorization.
- 
- Management of transferrable authorizations:  
Handover of authorizations from one user to another. (*WITH GRANT OPTION*)

# Management of transferrable authorizations


It is necessary to know the access **authorizations** of each **user** (some authorizations will be transferable to other users).

- One authorization can be transferred to other user in **mode transferable or not**




When one authorization is **revoked**:

- If the authorization was transferable, it is necessary to revoke all the transferred authorizations



if one user receives more than one authorization, each of the authorizations can be independently revoked.



B maintains the authorization a

## **3.2.1. Privacy and Security**



# Ethical and legal implications

---

## Extreme care on:

- Protection against the access or spreading of **personal data** to **non-authorized users**.
- Control the **flow** to third parties of **information** that can contain personal data or **information** that is apparently **aggregated** (parameterized queries) but that might **reveal particular** information for some parameters.
- **Custody of security backups**, retired or malfunctioning disks, etc.
- **Small devices** (USB disks or sticks, smart phones, tablets, etc.): **lost** or **stolen** very easily.

## Personal Data:

- Personal data is “any information relating to an identified or identifiable natural person.”
- In the European Union: The EU General Data Protection Regulation (GDPR) became enforceable across Europe on 25 May 2018
- Spain “***Agencia de Protección de Datos***”.

# Exercises

---

1.- The physical schema of a database is modified due to a change of disks. If an application uses a view that in turn uses tables that are stored in any of these disks, what happens to the application?

Justify the answer according to the ANSI/SPARC architecture and the concept of independence.

2.- Consider two transactions T1 and T2, which are running concurrently, both working on a piece of data X. Indicate whether one or more properties of transactions are not satisfied here. Briefly justify the answer.

	T1	X in T1	T2	X in T2
Time	read(X)	5		
	X=X+1	6		
	write(X)	6		
		6	read(X)	6
	confirm	6	...	6
	...			

---

3. How can be recovered a DB from a main memory failure using a logfile and checkpoints ? (Assume immediate updates)

4. Consider the following time diagram, and assuming a DBMS with immediate update. If a system failure occurs, as illustrated in the following figure

- What should the DBMS do if the system failure is a main memory loss?
- What should the DBMS do if the failure is a secondary memory loss ?

