



Unit 4. Service Deployment. Docker



Network Information System Technologies



Index

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



I. Goals

- ▶ Describe aspects to have into account when deploying distributed applications.
- ▶ Show the problems derived from dependencies, and how they can be addressed.
- ▶ Discuss deployment and its problems in a generic system.
- ▶ Know some deployment tools.
- ▶ Manage a specific deployment example.



Index

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



2. Concept of deployment

- ▶ **Deployment:** Activities that allow a software system to be ready for its use.
- ➡ Activities related with the installation, activation, upgrade and removal of components or the whole system.



2.1. Deployment of a distributed application

- ▶ A distributed application is a collection of heterogeneous components spread over a computer network
 - ▶ Many components, built by different developers
 - ▶ Those components may change in a fast and independent way (e.g., new versions, etc.)
- ▶ Application components should cooperate → there are dependences among them
- ▶ Those nodes may be heterogeneous (different hardware, operating systems, etc.), but each component has its own execution requirements
- ▶ There may be additional security requirements (privacy, authentication, etc.)



2.2. Deployment example

Let us assume the broker architecture developed in the second lab project

- ▶ It is composed of 3 autonomous agents(client, broker, worker)
 - ▶ They may have been developed in different programming languages, by different programmers
- ▶ Variable amount of instances in each component (e.g., several clients and several workers). Each instance...
 - ▶ ...may be started/stopped/restarted independently of the rest of instances
 - ▶ ...may fail on its own, independently of the others
 - ▶ ...has its own location (i.e., is placed in a given node)



2.2. Deployment example

- ▶ Dependences and requirements
 - ▶ Each client has to...
 - ▶ Know the broker location (IP address and frontend port)
 - ▶ Have a unique identity
 - ▶ Each worker has to...
 - ▶ Know the broker location (IP address and backend port)
 - ▶ Have a unique identity
 - ▶ All components have a given set of development requirements
 - ▶ JavaScript as their programming language
 - ▶ NodeJS as the JavaScript interpreter
 - ▶ ZeroMQ as the communication library



2.2. Deployment example: Manual deployment

- ▶ Copy the source code of each component in those computers where each instance should be run
 - ▶ We should guarantee that in those nodes the software base (NodeJS, ZeroMQ...) is correctly installed in the proper version
- ▶ Launch every instance of each component in the appropriate order (broker → workers → clients)
 - ▶ In the command line needed for starting each component, the appropriate arguments should be stated (e.g., IP address and port number to refer to the broker, agent identity, etc.)



Index

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



3. Service deployment


- ▶ We develop distributed applications in order to provide **services** (functionality) to remote clients
 - ▶ Application + Deployment = Service
- ▶ Every service sets a SLA (Service Level Agreement)
 - ▶ Functional definition (what it does)
 - ▶ Throughput (maximum req/sec, response time...)
 - ▶ Availability (percentage of time that the service remains active)
 - ▶ Although there are **ephemeral** services (i.e., short-lived, available for a short time interval), we focus on continuously available (i.e., **persistent**) services. For instance: gmail, Dropbox...
- ▶ **Service deployment** = installation, activation, upgrade and adaptation of the service.




3. Service deployment

- ▶ Installation and activation.- Software execution
 - ▶ Software dependence resolution (e.g., needed libraries)
 - ▶ Software configuration (compatible versions)
 - ▶ Determine the amount of instances of each component and their distribution in different nodes
 - ▶ Agent dependence resolution (e.g., port numbers)
 - ▶ Set the appropriate component start order
- ▶ Deactivation.- Stop the software system in an ordered way
- ▶ Upgrade.- Replace components (newer versions)
- ▶ Adaptation.- (With service continuity) after:
 - ▶ Failure/recovery of an agent
 - ▶ Agent configuration changes
 - ▶ Scaling (reaction to workload changes)

2. Services



Distributed Application



Service



Index

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



4. Deployment automation

A large scale deployment cannot be manually managed → An automation tool is needed

- ▶ Configuration of every component
 - ▶ File with a list of configuration parameters and dependency descriptions
 - ▶ The tool creates a specific configuration por every component instance
- ▶ Global configuration plan
 - ▶ Inter-component connection plan (list of public endpoints, list of dependences)
 - ▶ The location of each instance is decidec
 - ▶ Dependency resolution or binding (endpoints of interacting agents are connected, including dependences on external services)

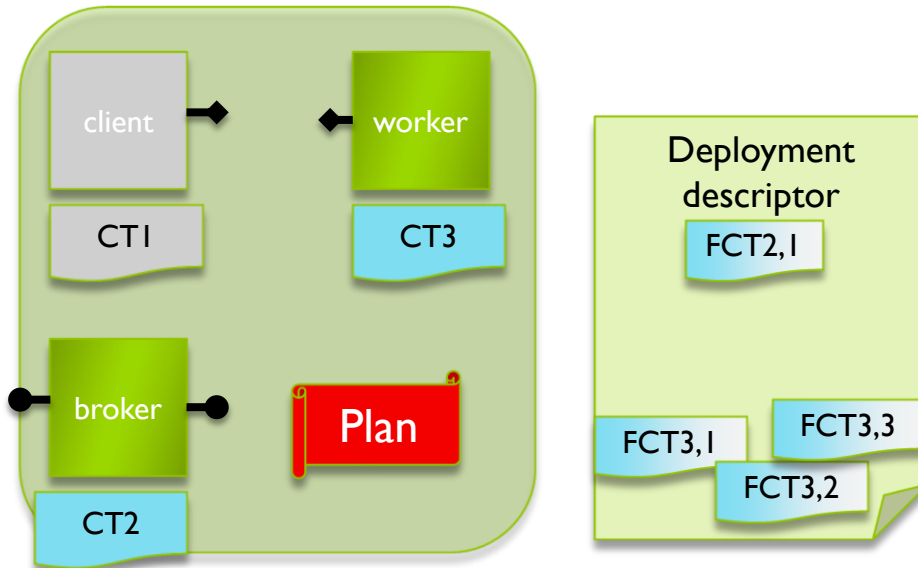


4. Deployment automation: Example

- ▶ Several client instances
 - ▶ Arguments: frontendURL, id
 - ▶ Dependency on a *broker*
- ▶ One broker instance
 - ▶ Arguments: frontendPort, backendPort
- ▶ Several worker instances
 - ▶ Arguments: backendURL, id
 - ▶ Dependency on a *broker*
- ▶ Plan
 - ▶ Their start order is: *broker, workers, clients*
 - ▶ The broker endpoints are: frontend (external) and backend (internal)
 - ▶ There are no dependencies on other external services

4. Deployment automation: Example

Distributed application



—◆— Dependency

●— Endpoint

—◆— Bound dependency

●— Bound endpoint


●— Exposed service endpoint

CT_i Configuration template for component *i*


FCT_{i,j} Filled configuration template for instance *j* of component *i*

4. Deployment automation: Example


Distributed application



Deployment



Deployed service



◆ Dependency

● Endpoint

◆ Bound dependency

● Bound endpoint

● Exposed service endpoint




Configuration template for component i




Filled configuration template for instance j of component i

4. Deployment automation: Example


Distributed application



Deployment



Deployed service



◆ Dependency

● Endpoint

◆ Bound dependency

● Bound endpoint

● Exposed service endpoint




Configuration template for component i




Filled configuration template for instance j of component i

4. Deployment automation: Example

Distributed application



Deployed service



◆ Dependency

● Endpoint

◆ Bound dependency

● Bound endpoint

● Exposed service endpoint

CT_i Configuration template for component *i*

FCT_{i,j} Filled configuration template for instance *j* of component *i*



4. Deployment automation

Dependency resolution. Options:

1. The program code defines how dependences must be resolved
 - ▶ E.g., reading data from a configuration file, or receiving data from a socket
 - ▶ Low level
2. **Dependency injection** (recommended)
 - ▶ The program code exposes local names for relevant interfaces
 - ▶ The container fills those variables with object instances
 - ▶ This generates a graph with all service component instances
 - ▶ The edges of the graph are links dependency-endpoint
 - As it has been shown in the previous slides



Index

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



5.1. Deployment in the cloud: IaaS

- ▶ Based on hardware virtualisation
 - ▶ Virtual machines in different “sizes” (i.e., hardware characteristics)
 - ▶ Flexibility in resource allocations
- ▶ However, there are some deployment limitations
 - ▶ Allocation decisions cannot be automated (low level)
 - ▶ Number of instances per component, location, VM type...
 - ▶ No choice on networking characteristics (latency, bandwidth...)
 - ▶ Inaccurate failure models
 - ▶ Not really independent failure modes
 - ▶ Limited help in recovering failed instances



5.2. Deployment in the cloud: PaaS

- ▶ SLA as the central element → SLA parameters for every component
 - ▶ Trying to achieve deployment automation
 - ▶ Deployment plans derived from the SLA
 - ▶ Upgrading and configuration plans
 - ▶ Current state
 - ▶ Still with a limited automation (initial deployment and autoscaling of simple services are automated, but service upgrading complying with its SLA is not completely supported yet)
 - ▶ Microsoft Azure is one of the best providers in this regard

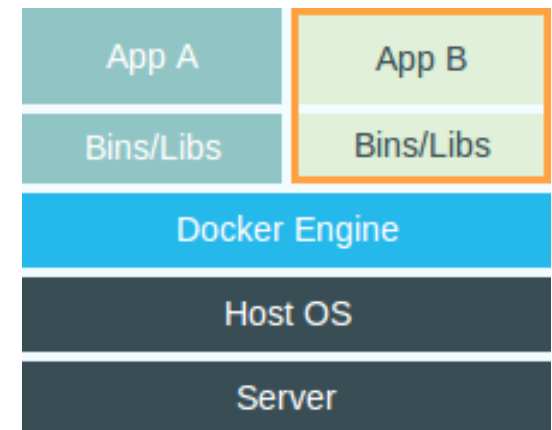
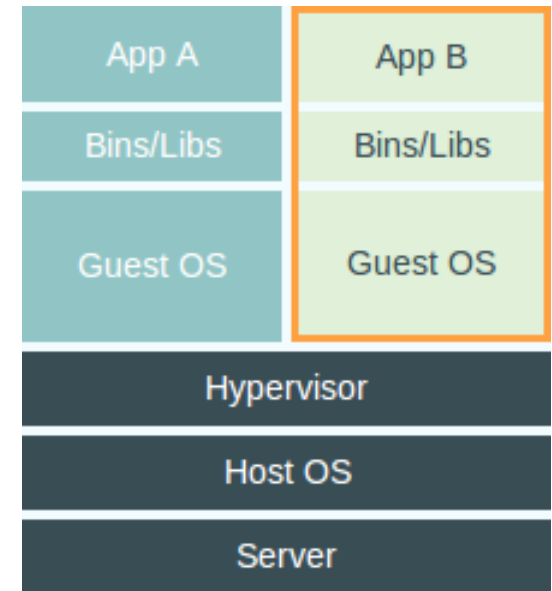


Index

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References

6. Containers

- ▶ **Provisioning** = booking the infrastructure needed by a distributed application.
- ▶ Resources needed for component intercommunication.
- ▶ Specific resources needed by each instance (CPU+ main memory + secondary storage). Alternatives:
 - ▶ Instances on **VM** → OS + Libraries
 - ▶ Instances on **containers** (light version of a VM) → Libraries
 - The host OS is used





6. Containers

Let us assume the usage of containers instead of VMs

- ▶ Lower flexibility
 - ▶ Software in each instance must be compatible with the host OS
 - ▶ Container isolation is not perfect
- ▶ Uses fewer resources
 - ▶ E.g., let us assume 100 instances of a component that needs 900MB (OS) + 100MB (libraries + program)
 - ▶ With VMs: $100 \times (900\text{MB} + 100\text{MB}) = 100\text{GB}$
 - ▶ With containers: $900\text{MB} + 100 \times 100\text{MB} = 10.9\text{GB}$
 - ▶ Space and time (e.g., in order to install the image) are saved!
- ▶ Easier deployment (automation tools with configuration files)
- ▶ May be applied in many scenarios



6. Containers: Docker

- ▶ Its Dockerfile configuration file automates the deployment of every instance
- ▶ It supports a version control system (Git)
- ▶ Besides the native file system, it defines a read-only file system that may be shared among containers
- ▶ Allows cooperation among developers through a public repository

NOTE.- We assume that the host OS is Linux (although there is a Docker version for Windows)



Index

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



7. Docker

The Docker ecosystem consists of 3 components:

1. **Images** (builder component) Read-only templates taken as a basis for container instantiation.
 - ▶ E.g., we may build an image that provides Linux+NodeJS+zmq, called **centos-zmq**
2. **Containers** (executor component) Created from images. They maintain all items needed for executing an instance.
 - ▶ E.g., in order to run our second lab project, each instance may be run in a container created from the **centos-zmq** image
3. **Repository** (distributor component) Global store where images can be pushed and pulled (hub.docker.com)
 - ▶ E.g., we may push the **centos-zmq** image there, and everyone may pull and use it for generating containers



7.1. Docker: Image

- ▶ In the repository, there are predefined images for different Linux distributions (e.g., image **Centos:7.4.1708**)
- ▶ New image = baseImage + instructions. Examples:
 - ▶ `Centos:7.4.1708` + instr. to install node → `centos-nodejs`
 - ▶ `centos-nodejs` + instr. to install zmq → `centos-zmq`
 - ▶ `centos-zmq` + ... → build images for each component (*client, broker, worker*)
- ▶ Docker uses console commands
- ▶ General structure: `docker` action options arguments
 - ▶ Information on local images: `docker images`
 - ▶ Information on a given image: `docker history imageName`



7.2. Docker. Container

- ▶ Create and start a container from an image

```
docker run options image initialProgram
```

- ▶ E.g., `docker run -i -t centos bash` downloads the image `centos`, creates a container, allocates a file system, allocate a networking interface and an internal IP address, and runs `bash`
 - ▶ Options `-i -t` choose the interactive mode (the host console remains open and connected to the container)
- ▶ Now, we may modify the container state using command in an interactive session at the console
- ▶ Finally, we may create a new image saving the current state of the container

```
docker commit containerName imageName
```




7.2. Docker. Command groups

Group	Description
config	Configuration management
container	Operations with containers
context	Contexts for distributed deployment (k8s, ...)
image	Image management
network	Network management
service	Service (i.e. multiple containers that run the same image) management in distributed deployments (e.g. with docker-compose or docker swarm)
system	Global management
volume	Secondary storage management



7.3. Docker. Main commands

1. Life cycle management (in the container group)
 - ▶ **docker commit | run | start | stop**
2. Informative (multiple groups)
 - ▶ **docker logs | ps | info | images | history**
3. Repository access (image group)
 - ▶ **docker pull | push**
4. Other (container group)
 - ▶ **docker cp | export**



Index

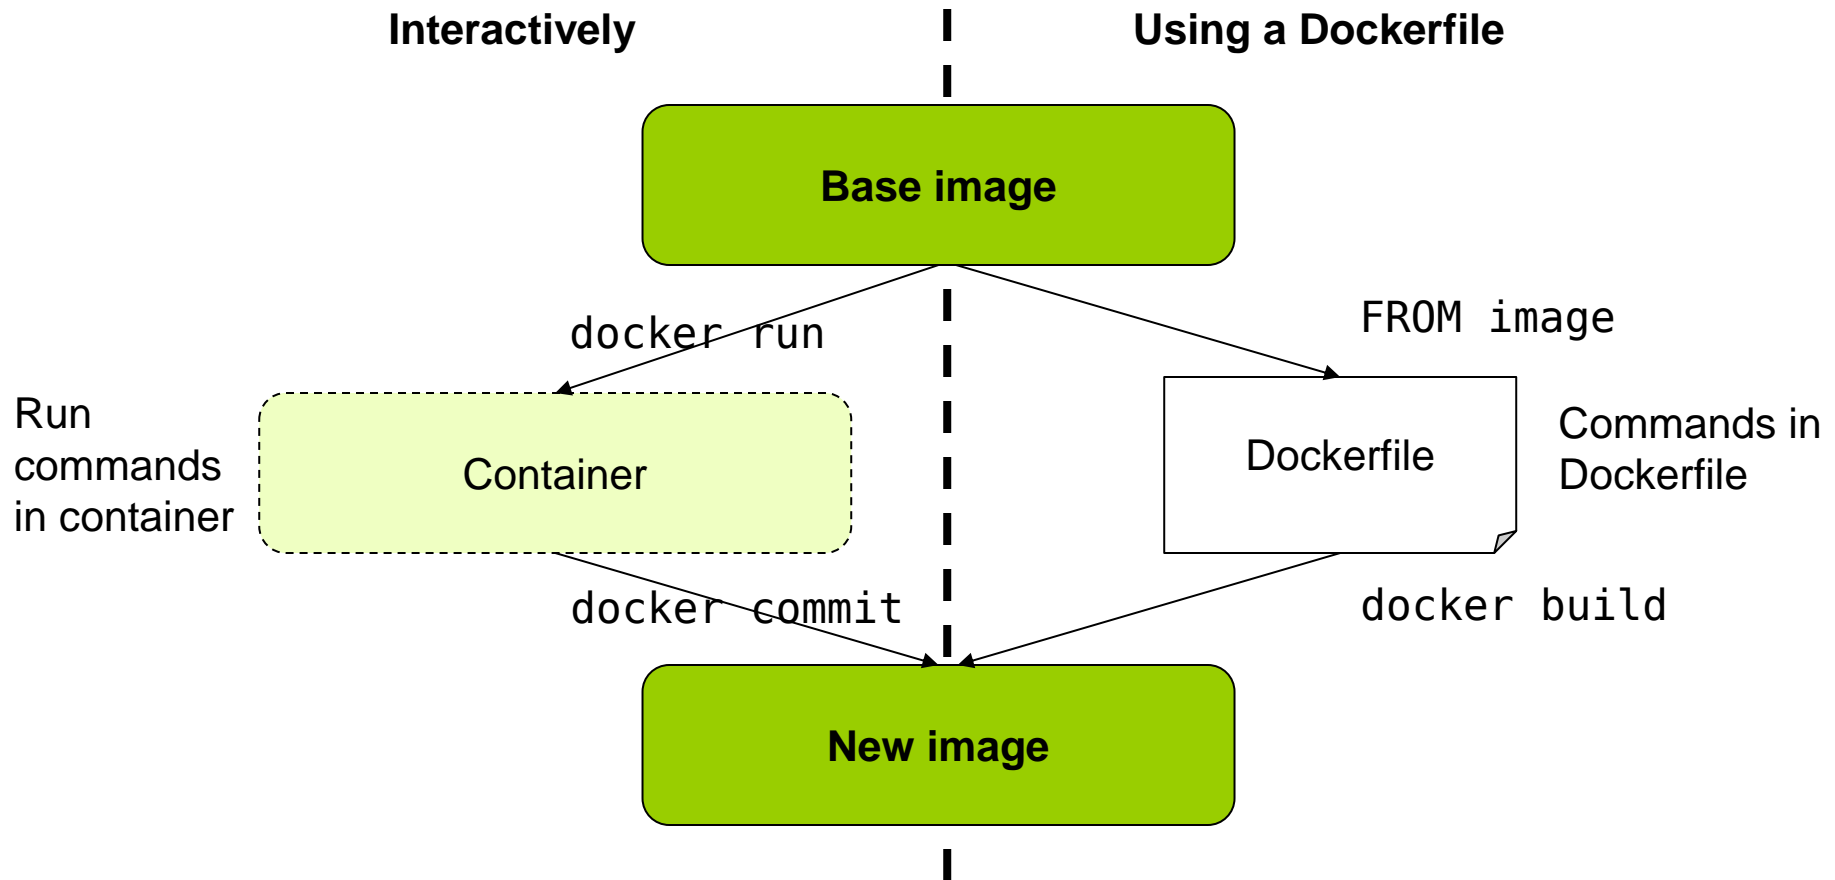
1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



8. Image creation

Interactively

Using a Dockerfile





8. Image creation

- ▶ Two alternatives for image creation:

- ▶ Interactively

- ▶ Choose a base image and start a container

- ```
docker run -i -t baseImage initialProgram
```

- ▶ Modify interactively the container state

- ▶ Save the current state as a new image:

- ```
docker commit containerID newImageName
```

- ▶ Create a new image from the instructions saved in a Dockerfile text file

- ```
docker build -t newImageName pathToDockerfile
```



## 8.1. Interactive image creation. Example

- ▶ Image that may run NodeJS programs with 'zeromq'.
  - Step 1: Launch Docker using an interactive CentOS image (running bash):

```
$ docker run -i -t centos:7.4.1708 bash
```

- Step 2: Run these commands in the container:

```
$ curl --silent --location https://rpm.nodesource.com/setup_10.x | bash -
$ yum install -y nodejs
$ yum install -y epel-release
$ yum install -y zeromq-devel
$ yum install -y make python gcc-c++
$ npm install zeromq@4
$ exit
```

- From the host console, get the name and ID of that container: `docker ps -a`
  - Create the new image:  
`docker commit containerID newImageName`



## 8.2. Image creation with Dockerfile. Example

1. This Dockerfile is written to disk:

```
FROM centos:7.4.1708
RUN curl --silent --location https://rpm.nodesource.com/setup_10.x | bash -
RUN yum install -y nodejs
RUN yum install -y epel-release
RUN yum install -y zeromq-devel make python gcc-c++
RUN npm install zeromq@4
```

2. Change the current directory to that where the Dockerfile is placed
3. Run this command:
  - ▶ `docker build -t tsr1718/centos-zmq .`



## 8.3. Docker. Dockerfile

---

- ▶ Each line starts with a command (in uppercase, by convention)
- ▶ The first instruction (in the first line) must be `FROM imageBase`
- ▶ `RUN` command runs that command in the container shell
- ▶ `ADD` source destination
  - ▶ Copy files from a source (URL, directory or file) to a path in the container
  - ▶ If source is a directory, then it copies all its contents. If it is a compressed file, then it is extracted in the destination
- ▶ `COPY` source destination is equal to `ADD`, but it does not extract compressed files
- ▶ `EXPOSE` port states the port number where the container listens to incoming requests





## 8.3. Docker. Dockerfile

---

- ▶ **WORKDIR** path specifies the container working directory for subsequent **RUN**, **CMD**, or **ENTRYPOINT** commands
- ▶ **ENV** variable value assigns value to an environment variable that may be used by the programs to be run in the container
- ▶ **CMD** command arg1 arg2 ... provides default values for the command and arguments to be run in the container
- ▶ **ENTRYPOINT** command arg1 arg2 ... runs that command when the container is started (and the container is terminated when that command ends)

There must be a single **CMD** or **ENTRYPOINT** in the Dockerfile (otherwise, only the last one is used)



# Index

---

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



## 9. Multiple components in a node. Example

- ▶ Example: Service that consists of a “broker” and a “worker” that may be replicated as many times as needed.
- ▶ Almost identical to that seen in part 6.3 of lab 2.

```
// ROUTER-ROUTER request-reply broker in NodeJS
const zmq = require('zmq')
let cli=[], req=[], workers=[]

let args = process.argv.slice(2)
let fePortNbr = args[0] || 9998
let bePortNbr = args[1] || 9999

let sc = zmq.socket('router') // frontend
let sw = zmq.socket('router') // backend

sc.bind('tcp://*:'+fePortNbr)
sw.bind('tcp://*:'+bePortNbr)
```

```
sc.on('message',(c,sep,m)=> {
 if (workers.length==0) {
 cli.push(c); req.push(m)
 } else {
 sw.send([workers.shift(),' ',c,' ',m])
 }
})

sw.on('message',(w,sep,c,sep2,r)=> {
 if (c=='') {workers.push(w); return}
 if (cli.length>0) {
 sw.send([w,' ',
 cli.shift(),' ',req.shift()])
 } else {
 workers.push(w)
 }
 sc.send([c,' ',r])
})
```



## 9. Multiple components in a node. Example

- ▶ Worker and client programs.
  - ▶ Similar to those in lab 2, with command-line arguments and default values.

*// worker in NodeJS with URL & id arguments*

```
const zmq = require('zeromq')
let req = zmq.socket('req')

let args = process.argv.slice(2)
let backendURL = args[0] ||
'tcp://localhost:9999'

let myID = args[1] ||
'WID'+parseInt(Math.random()*5000)
let replyText = args[2] || 'resp'

req.identity = myID
req.connect(backendURL)
req.on('message', (c, sep, msg) => {
 setTimeout(() => {
 req.send([c, '', replyText])
 }, 1000)
})
req.send(['', '', ''])
```

*// client in NodeJS with URL & id arguments*

```
const zmq = require('zeromq')
let req = zmq.socket('req')

let args = process.argv.slice(2)
let brokerURL = args[0] ||
'tcp://localhost:9998'

let myID = args[1] ||
'CID'+parseInt(Math.random()*10000)
let myMsg = args[2] || 'Hola'

req.identity = myID
req.connect(brokerURL)
req.on('message', (msg) => {
 console.log('resp: '+msg)
 process.exit(0);
})
req.send(myMsg)
```



## 9. Multiple components in a node. Example

---

- ▶ With multiple components, some dependencies arise:
  - ▶ Clients need to know the frontend URL (IP address and port)
  - ▶ Workers need to know the backend URL (IP address and port)
- ▶ But we do not know the IP address until the broker container is started
- ▶ Manual alternative
  - ▶ Once the broker container is started, we get its IP address
  - ▶ We manually update those values in the client and worker Dockerfiles in order to generate their correct images
  - ▶ Start then workers and clients
- ▶ Automation:
  - ▶ A deployment plan is defined = Description of components, properties and relations
  - ▶ We use a tool that completes the deployment reading that plan.



## 9.1. Manual method: Broker

- ▶ In a directory, we copy broker.js and this Dockerfile

```
FROM tsr1718/centos-zmq
COPY ./broker.js broker.js
EXPOSE 9998 9999
CMD node broker
```

- ▶ Ports to be used: frontend=9998, backend=9999
- ▶ Build the image: **docker build -t broker .**
- ▶ Start the broker container: **docker run -d broker**
- ▶ Find out its IP address
  - ▶ **docker ps -a** in order to the container ID
  - ▶ **docker inspect ID** in order to get its IP address



## 9.1. Manual method: Worker

- ▶ In another directory, we copy worker.js and this Dockerfile (assuming that the broker IP address is a.b.c.d)

```
FROM tsr1718/centos-zmq
COPY ./worker.js worker.js
CMD node worker tcp://a.b.c.d:9999
```

- ▶ Take a look at the previous slide on how to get the a.b.c.d address
- ▶ In this directory, run: **docker build -t worker .**
- ▶ Start a worker with **docker run -d worker**
- ▶ We may start as many instances as needed



## 9.1. Manual method: Local client

- ▶ In another directory, we copy client.js and write this Dockerfile:

```
FROM tsr1718/centos-zmq
COPY ./client.js client.js
CMD node client tcp://a.b.c.d:9998
```

- ▶ In that directory, we run: `docker build -t client .`
- ▶ Start a client with: `docker run -d client`
  - ▶ We may start as many instances as needed





## 9.1. Manual method: Remote client

---

- ▶ We start the broker with `docker run -p 8000:9998 -d broker`
  - ▶ Option `-p hostPort:containerPort` allows external accesses to the frontend port (through port 8000 in the host)
  - ▶ Option `-d` starts the container in background mode
- ▶ Clients may be started directly, without containers. They must connect to `tcp://hostIPAddress:8000`



## 9.2. Automated method. Example

- ▶ Create broker subdir., copy broker.js and this Dockerfile

```
FROM tsr1718/centos-zmq
COPY ./broker.js broker.js
EXPOSE 9998 9999
CMD node broker
```

- ▶ Create worker subdir., copy worker.js and this Dockerfile

```
FROM tsr1718/centos-zmq
COPY ./worker.js worker.js
CMD node worker $BROKER_URL
```

- ▶ Create client subdir., copy client.js and this Dockerfile

```
FROM tsr1718/centos-zmq
COPY ./client.js client.js
CMD node client $BROKER_URL
```

## 9.2. Automated method. Example

- ▶ Copy the text on the right to a `docker-compose.yml` file
- ▶ This command...  
`docker-compose up -d`
  - ▶ builds all images
  - ▶ starts one instance of each one in the correct order
- ▶ We may start n instances of service X with  
`docker-compose up -d --scale X=n`
- ▶ There are other commands
  - ▶ Read the guide of this unit

```
version: '2'
services:
 cli:
 image: client
 build: ./client/
 links:
 - bro
 environment:
 - BROKER_URL=tcp://bro:9998
 bro:
 image: broker
 build: ./broker/
 expose:
 - "9998"
 - "9999"
 wor:
 image: worker
 build: ./worker/
 links:
 - bro
 environment:
 - BROKER_URL=tcp://bro:9999
```



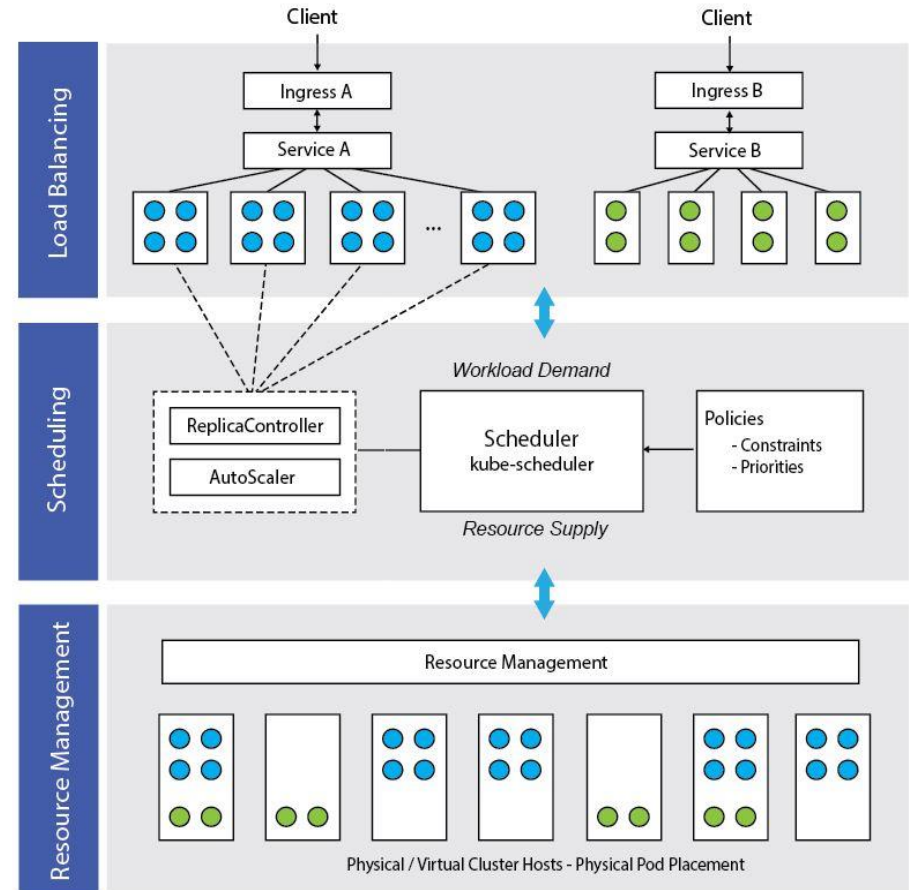
# Index

---

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References

## 10. Multiple components in different nodes

- ▶ Docker-compose only manages component in a single host
- ▶ If we need multiple hosts → the best option is **kubernetes**
  - ▶ It is a container orchestrator, but it does not depend on Docker
  - ▶ The guide for this unit provides a general description of its elements





## 10. Multiple components in different nodes

---

Kubernetes. Main elements:

- ▶ **Cluster** and node (real or virtual)
- ▶ **Pod**: smallest deployment unit
  - ▶ It includes containers that share a *namespace* and storage volumes
- ▶ **Replication controllers**: manage the life cycle of a group of pods,
  - ▶ ensuring that a specified number of instances is running,
    - scaling, replicating and recovering pods
- ▶ **Deployment controllers**: upgrade the distributed application
- ▶ **Service**: defines a set of pods and their public access
- ▶ **Secrets** (credentials management)
- ▶ **Volumes** (persistence)



# Index

---

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



## II. Learning Results

---

- ▶ After this unit, the student must be able to
  - ▶ Know in some detail the salient aspects to be considered in the deployment of distributed applications
  - ▶ Understand the problems posed by dependencies, and the ways to address them
  - ▶ Use Docker in order to deploy distributed services





# Index

---

1. Goals
2. Concept of deployment
3. Service deployment
4. Deployment automation
5. Deployment in the cloud
6. Containers
7. Docker
8. Image creation
9. Multiple components in a node
10. Multiple components in different nodes
11. Learning Results
12. References



## 12. References

---

- ▶ Inversion of Control/Dependency injection
  - ▶ <http://martinfowler.com/articles/injection.html>
  - ▶ <http://www.springsource.org/>
- ▶ [www.docker.com](http://www.docker.com) (Official Docker website)
  - ▶ [docs.docker.com/userguide/](http://docs.docker.com/userguide/) (**Official documentation**)
  - ▶ [docs.docker.com/compose/](http://docs.docker.com/compose/) (Compose)
- ▶ [github.com/wsargent/docker-cheat-sheet](https://github.com/wsargent/docker-cheat-sheet) (Summary of Docker)
- ▶ <http://kubernetes.io>
- ▶ [12factor.net/](http://12factor.net/) (*“The twelve-factor app”* methodology)